

# CS2102: Database Systems

## Relational Algebra

# Announcement

---

# Announcement

Assignments  
Project

## Assignments

### Updated Assignment Deadlines

#	Components	Due Week	Day	Time
1	SQL	<i>Week 6.5 (recess week)</i>	Saturday	12:00 ( <i>noon</i> )
2	Refinement	<i>Week 14 (reading week)</i>	Monday	12:00 ( <i>noon</i> )

### Note

- Assignment 2 deadline is on reading week because the original week 12 is well being day
  - We keep it "free of work" for week 12
- L00 Overview online slide has been updated
- We have also updated the [Course Syllabus](#) page

# Announcement

Assignments  
**Project**

## Project

### Project Group Formation

- Opens now on Canvas > People > Groups > [Project](#)
  - Self sign-up
  - Can be from different tutorial group
  - You may advertise on [Canvas discussion](#) or telegram
- Close on Week 3, Saturday, 28 January 2023 at 12:00 noon
  - Afterwards, we will randomly allocate people to groups and/or merge groups to form groups of at least 3 people

# Recap

---

# Recap

Relational Model  
Integrity

## Relational Model

### Basic Concept: Relations

- Unified representation of all data using tables with rows and tables
  - Definitions
    - **Domains**
      - NULL
    - **Tuples**
    - **Relations**
- Possible values of an attribute*  
*Unknown/invalid values*
- A single row on the table*  
*Set of tuples (all rows)*

### Structural Integrity Constraints

- Condition that restricts what constitutes valid data
  - **Domain**      *(what type can be inserted?)*
  - **Key**            *(which column(s) identifies others?)*
  - **Foreign Key**    *(does it exist somewhere else?)*

X	Y
1	2
100	777
:	:

$$\Rightarrow X < Y$$

Table "Movies"

<u>id</u>	title	genre	opened
101	Aliens	Action	1986
102	Logan	Drama	2017
:	:	:	:

Table "Cast"

movie_id	actor_id	role
101	20	Ellen Ripley
102	21	Logan
:	:	:

Table "Actors"

<u>id</u>	name	dob
20	Sigourney Weaver	08-10-1949
21	Hugh Jackman	12-10-1968
:	:	:

# Recap

Relational Model  
**Integrity**  
- Misconception  
- Foreign Key

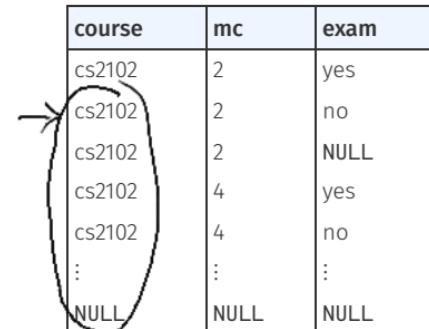
## (Structural) Integrity Constraints

### Possible Misconception

- (*Foreign*) key constraints are not an intrinsic property of a relation
  - Constraints are specified by the database designer to define what constitutes valid data

### Example

- Without any key constraints, the relation on the right is perfectly valid (*DBMS does not complain*)
  - However, the meaning is problematic from the application perspective
    - (e.g., *CS2102 gives different credits, with and without exam?*)
- **Goal:** Avoid different values for "mc" and "exam" for the same course
  - (i.e., *"course" must uniquely identify both "mc" and "exam"*)
    - How? Pick {course} as the primary key



course	mc	exam
cs2102	2	yes
cs2102	2	no
cs2102	2	NULL
cs2102	4	yes
cs2102	4	no
:	:	:
NULL	NULL	NULL

# Recap

## Relational Model Integrity

- Misconception
- Foreign Key

3	1	2

(3, 3, NULL)

## (Structural) Integrity Constraints

### Foreign Key Condition

Each foreign key in  $R_1$  must satisfy one of the following:

- Appear as **primary key** in  $R_2$
- Be a NULL value (or a tuple containing at least one NULL value)



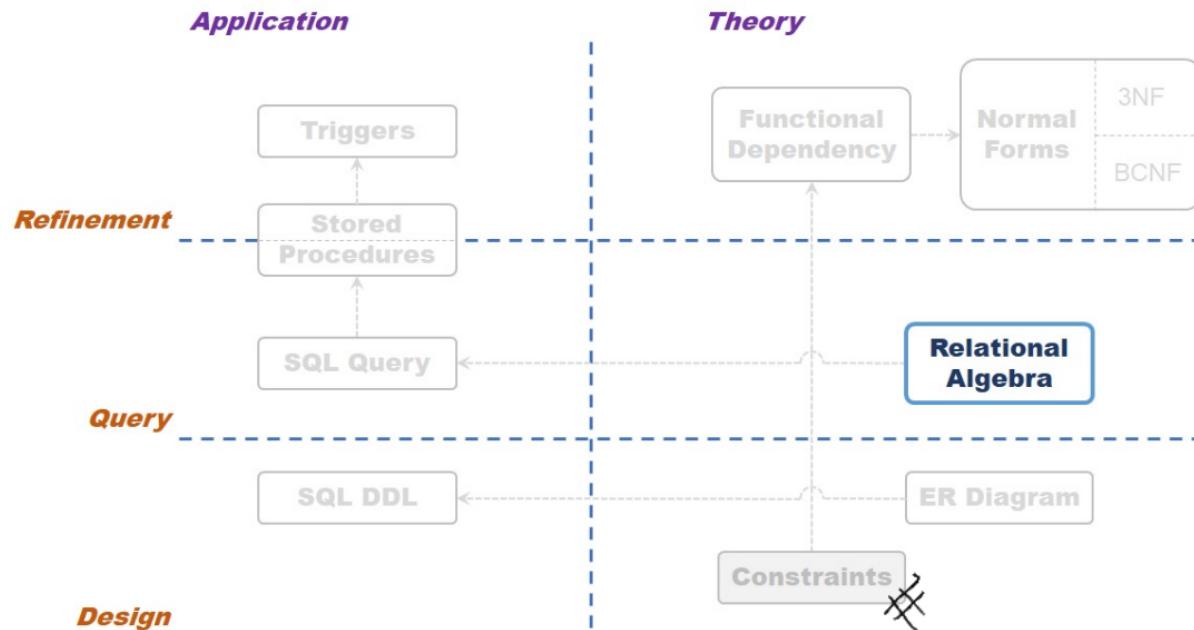
### Specification

A referential constraint is satisfied if one of the following conditions is true, depending on the <match option> specified in the <referential constraint definition>:

- If no <match type> was specified then, for each row R1 of the referencing table, either at least one of the values of the referencing columns in R1 shall be a null value, or the value of each referencing column in R1 shall be equal to the value of the corresponding referenced column in some row of the referenced table.
- If MATCH FULL was specified then, for each row R1 of the referencing table, either the value of every referencing column in R1 shall be a null value, or the value of every referencing column in R1 shall not be null and there shall be some row R2 of the referenced table such that the value of each referencing column in R1 is equal to the value of the corresponding referenced column in R2.

- at least one of the values of the referencing columns in  $R_1$  shall be a NULL value, or
- the value of each referencing column in  $R_1$  shall be equal to the value of the corresponding referenced column in some row of the referenced table

# Roadmap



# Roadmap

## Overview

## Overview

### *Relational Algebra*

- Preliminary
- Closure

### *Three-Valued Logic*

- Logical
- Relational
- Arithmetic

### *Basic Operators*

- Unary Operators
- Set Operators

### *Join Operators*

- Inner Joins
- Outer Joins

### *Complex Expressions*

# Relational Algebra

---

# Relational Algebra

l + 2      + x -

## Preliminary

- Algebra

- Why?

Closure

Example Data

## Preliminary

### Algebra

An **algebra** is a mathematical system consisting of:

- **Operands** Variables or values from which new values can be constructed
- **Operators** Symbols denoting procedures that construct new values from the given values

### Relational Algebra

- **Operands** Relations (or variables representing relations)
- **Operators** Transformation from one or more input relations into one output relation

One arg Operators

- Unary
- Binary

two args

selection ( $\sigma$ ), projection ( $\pi$ ), renaming ( $\rho$ )  
cross product ( $\times$ ), union ( $u$ ), intersection ( $\cap$ ),  
difference ( $-$ )

### Note

All other operators\* can be expressed using these basic operators\*\*.

\*\*In fact, intersection can be constructed using union and difference.

\*Except for "aggregate" and "sorting" operations.

# Relational Algebra

## Preliminary

- *Algebra*

- *Why?*

Closure

Example Data

## Preliminary

### Why Algebra?

- Allows for query optimization\*
- Rewrite a query such that:
  - It is *equivalent* to the original query
  - It can be evaluated faster than the original query

## Example

### Question

Which of the following operation is "faster"?

- A.  $(a \times b) + (a \times c)$
- B.  $a \times (b + c)$

### Quiz #1

Is it equivalent?

\*While we do not penalize for inefficient queries, it is useful to know.

# Relational Algebra

Preliminary

Closure

- Definition

- Theorem

- Implication

Example Data

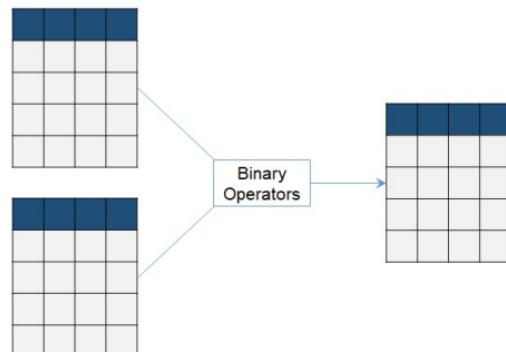
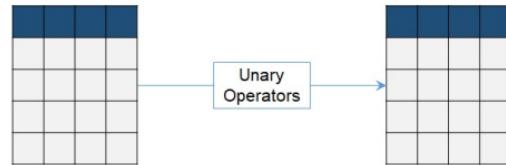
## Closure Property

### Definition

We say that a set of values is **closed** under the set of operators if any combination of the operators produces only values in the given set.

### Example

- $\mathbb{Z}^+$  is closed under  $\{+\}$ 
  - but not under  $\{+, -\}$
- $\mathbb{Z}$  is closed under  $\{+, -, \times\}$ 
  - but not under  $\{+, -, \times, /\}$
- $\mathbb{R}$  is closed under  $\{+, -, \times, /\}^*$ 
  - but not under  $\{+, -, \times, /, \sqrt{\cdot}\}$
- $\mathbb{C}$  is algebraically closed



\*Division operation (~~1/0~~) should not be a division by zero.

# Relational Algebra

Preliminary

Closure

- Definition

- Theorem

- Implication

Example Data

## Closure Property

### Definition

We say that a set of values is **closed** under the set of operators if any combination of the operators produces only values in the given set.

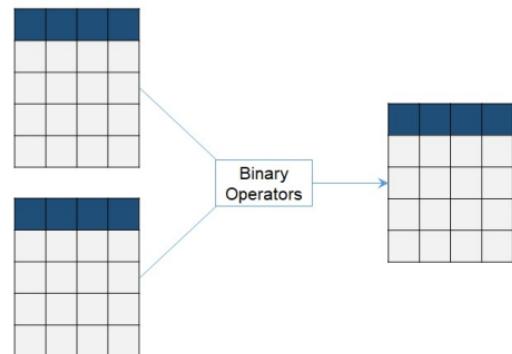
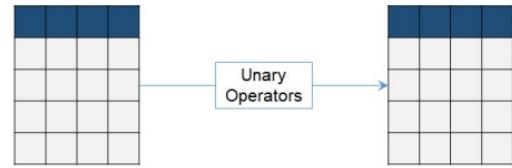
### Theorem

Relations are *closed* under Relational Algebra.

#### Proof

- Inputs are relations
- Outputs are relations

⇒ Output of one relation → input to another relation  
⇒ No other output results are possible



# Relational Algebra

Preliminary

Closure

- Definition

- Theorem

- Implication

Example Data

$$(\underline{a \times b}) \oplus (\underline{a \times c})$$

## Closure Property

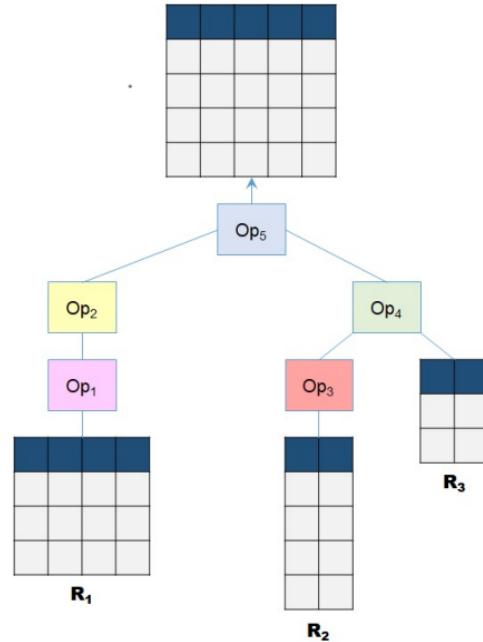
Implication

We can chain operations because all inputs and outputs are **relations** (*output of one operators can be input of another operators*).

### Example

Using the image on the right we get

```
Op5 (  
    Op2 ( Op1 ( R1 ) ),  
    Op4 ( Op3 ( R2 ),  
        R3  
    )  
)
```



# Relational Algebra

Preliminary

Closure

**Example Data**

- Schema

- Instance

## Example Data

### Schema

A simplified company schema

*Employees(name: TEXT, age: INT, role: TEXT)*

*Managers(name: TEXT, office: TEXT)*

*Teams(ename: TEXT, pname: TEXT, hours: INT)*

*Projects(name: TEXT, manager: TEXT, start\_year: INT, end\_year: INT)*

### Foreign Key Constraints

- (Manager.name)  $\rightsquigarrow$  (Employees.name)
- (Teams.ename)  $\rightsquigarrow$  (Employees.name)
- (Teams.pname)  $\rightsquigarrow$  (Projects.name)
- (Projects.manager)  $\rightsquigarrow$  (Manager.name)

# Relational Algebra

Preliminary  
Closure

**Example Data**  
- Schema  
- Instance

## Example Data

Instance

Table "Teams"

ename	pname	hours
Sarah	BigAI	10
Sam	BigAI	5
Bill	BigAI	15
Judy	GlobalDB	20
Max	GlobalDB	5
Sarah	GlobalDB	10
Emma	GlobalDB	35
Max	CoreOS	40
Bill	CoreOS	30
Sam	CoolCoin	40
Sarah	CoolCoin	25
Emma	CoolCoin	10

Table "Projects"

name	manager	start_year	end_year
BigAI	Judy	2020	2025
FastCash	Judy	2018	2025
GlobalDB	Jack	2019	2023
CoreOS	Judy	2020	2020
CoolCoin	Jack	2015	2020

Table "Employees"

name	age	role
Sarah	25	dev
Judy	35	sales
Max	52	dev
Marie	36	hr
Sam	30	sales
Bernie	19	NULL
Emma	28	dev
Jack	40	dev
Bill	45	dev

Table "Managers"

name	office
Judy	#03-20
Jack	#03-10

L02.tbl

Boolean : { True, False, NULL }

# Three-Valued Logic

---

# Three-Valued Logic

Logical  
- Truth Table  
- Alternate  
Relational &  
Arithmetic

## Logical Operations

### Truth Table

$c_1$	$c_2$	$c_1 \wedge c_2$	$c_1 \vee c_2$	$\neg c_1$
False	False	False	False	True
False	NULL	False	NULL	True
False	True	False	True	True
NULL	False	False	NULL	NULL
NULL	NULL	NULL	NULL	NULL
NULL	True	NULL	True	NULL
True	False	False	True	False
True	NULL	NULL	True	False
True	True	True	True	False

# Three-Valued Logic

Logical  
- Truth Table  
- Alternate  
Relational &  
Arithmetic

## Logical Operations

### Alternate Truth Table

#### Conjunction

$c_1 \wedge c_2$		$c_1$		
		False	NULL	True
$c_2$	False	False	False	False
	NULL	False	NULL	NULL
	True	False	NULL	True



#### Disjunction

$c_1 \vee c_2$		$c_1$		
		False	NULL	True
$c_2$	False	False	NULL	True
	NULL	NULL	NULL	True
	True	True	True	True

# Three-Valued Logic

Logical  
Relational &  
Arithmetic

$$\text{NULL} + n \longrightarrow \text{NULL}$$

## Relational & Arithmetic Operations

### Basic Relational & Arithmetic

Any relational/arithmetic operation with NULL produces NULL values

Operation	Result	Operation	Result
NULL < 2023	NULL	NULL = NULL	NULL
NULL ≥ 0	NULL	NULL <> NULL	NULL

### NULL as Values

We add operators to treat NULL as values

$\equiv$        $\neq$

$v_1$	$v_2$	$v_1 = v_2$	$v_1 \equiv v_2$
NULL	NULL	True	False
$v_1$	NULL	False	True
NULL	$v_2$	False	True
$v_1$	$v_2$	$v_1 = v_2$	$v_1 \neq v_2$

$\times$

$\neq$

$\neq$

$\equiv$

# Basic Operators

---

# Basic Operators

## Selection

- Syntax
- Semantics
- Example
- Exercise

## Projection

## Renaming

## Caution

## Selection

Selection  $\sigma[c](R)$   
greek

### Syntax

→  $\sigma[c](R)$  where c is a condition that returns a boolean value:

	Expr	Example	Name
1	(expr)	$\sigma[(start\_year=2020)](Projects)$	precedence
2	attr op const	$\sigma[start\_year=2020](Projects)$	constant selection
2	attr <sub>1</sub> op attr <sub>2</sub>	$\sigma[start\_year=end\_year](Projects)$	attribute selection
3	¬ expr	$\sigma[-(start\_year=2020)](Projects)$	negation
4	expr <sub>1</sub> ∧ expr <sub>2</sub>	$\sigma[start\_year=2020 \wedge manager='Judy'](Projects)$	conjunction
5	expr <sub>1</sub> ∨ expr <sub>2</sub>	$\sigma[start\_year=2020 \vee manager='Judy'](Projects)$	disjunction

with  $op \in \{=, <, <, \leq, \geq, >, \equiv, \not\equiv\}$ .

Operator Precedence: (), op, ¬, ∧, ∨\*

not  
and  
or

\*Without the parentheses and op, this is often called Disjunctive Normal Form.

# Basic Operators

Accept

Reject

## Selection

- Syntax

- Semantics

- Example

- Exercise

Projection

Renaming

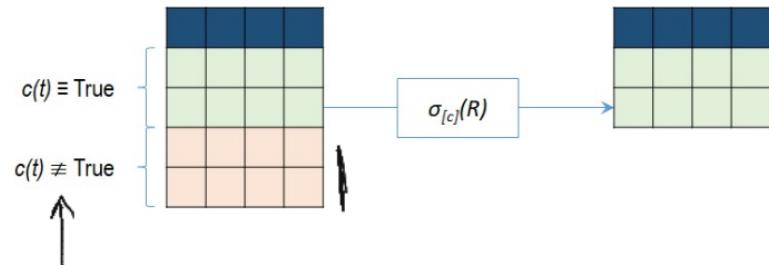
Caution

## Selection

### Semantics

$\sigma_{[c]}(R)$  selects all tuples from a relation  $R$  (i.e., rows from a table) that satisfy the selection condition  $c$ .

### Visualization



### Properties

- The result have the same schema as the input relation
- The number of rows are often smaller

<sup>#</sup>We can always "rearrange" the rows such that the visual above works because we are working with a **set** of tuples.

# Basic Operators

## Selection

- Syntax
- Semantics
- Example
- Exercise

## Projection

## Renaming

## Caution

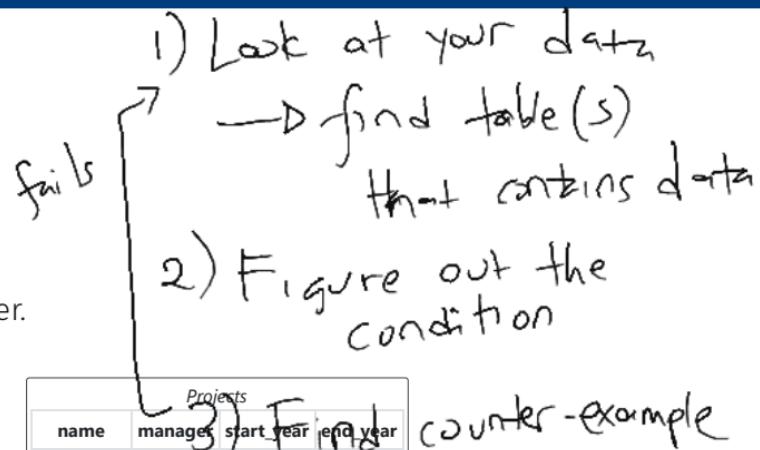
## Selection

### Example

Find all projects where Judy is the manager.

$\$T1 := \sigma_{[(\text{manager} = 'Judy')]}(\text{Projects})$			
name	manager	start_year	end_year
"BigAI"	"Judy"	2020	2025
"FastCash"	"Judy"	2018	2025
"CoreOS"	"Judy"	2020	2020

3 Rows



Projects			
name	manager	start_year	end_year
"BigAI"	"Judy"	2020	2025
"FastCash"	"Judy"	2018	2025
"GlobalDB"	"Jack"	2019	2023
"CoreOS"	"Judy"	2020	2020
"CoolCoin"	"Jack"	2015	2020

5 Rows

# Basic Operators

## Selection

- Syntax
- Semantics
- Example
- **Exercise**

Projection

Renaming

Caution

## Selection

### Exercise #1

Find all employees that (a) do not work in Sales and are younger than 30 or  
(b) work in HR.

$$\overline{\sigma}_{[(\text{role} \neq 'Sales') \wedge (\text{age} \leq 30)]}(\text{Emps}) \\ \vee \text{role} \equiv 'HR'$$

# Basic Operators

Selection

## Projection

- Syntax

- Semantics

- Example

- Exercise

Renaming

Caution

## Projection

### Syntax

$\pi_{[\ell]}(R)$  where  $\ell$  is an **ordered** list of attributes

#### Note:

- For simplicity, we do not allow the following on  $\ell$ :
  - Operations (e.g.,  $\pi[A_1 + A_2](R)$ )
  - Duplicate (e.g.,  $\pi[A_1, A_1](R)$ )
- The ordering of the attribute matters (i.e.,  $\pi[A_1, A_2](R) \neq \pi[A_2, A_1](R)$ )

# Basic Operators

Selection  
Projection

- Syntax
- Semantics
- Example
- Exercise

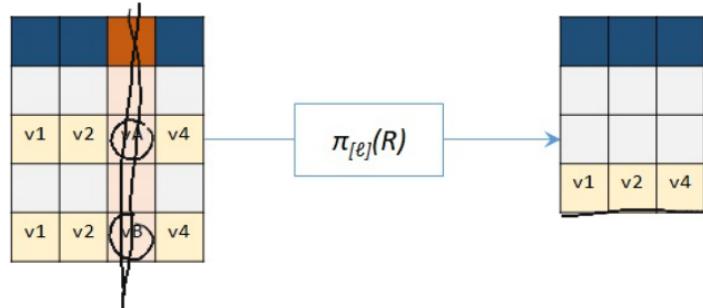
Renaming  
Caution

## Projection

### Semantics

$\pi_{[\ell]}(R)$  keeps only the column (*i.e., projects*) specified in the ordered list  $\ell$  and in the same order.

### Visualization



### Properties

- The resulting schema is specified by  $\ell$
- The number of rows may be smaller
  - Because relation is defined as a **set** of tuples

# Basic Operators

Selection  
Projection

- Syntax
- Semantics
- Example
- Exercise

Renaming  
Caution

## Projection

### Example

Find all project names that have team members.

$\$T1 := \pi_{pname}(Teams)$		
pname		
	"BigAI"	
	"GlobalDB"	
	"CoreOS"	
	"CoolCoin"	

4 Rows

Teams		
ename	pname	hours
"Sarah"	"BigAI"	10
"Sam"	"BigAI"	5
"Bill"	"BigAI"	15
"Judy"	"GlobalDB"	20
"Max"	"GlobalDB"	5
"Sarah"	"GlobalDB"	10
"Emma"	"GlobalDB"	35
"Max"	"CoreOS"	40
"Bill"	"CoreOS"	30
"Sam"	"CoolCoin"	40
"Sarah"	"CoolCoin"	25
"Emma"	"CoolCoin"	10

12 Rows

# Basic Operators

Selection  
Projection

- Syntax
- Semantics
- Example
- Exercise

Renaming  
Caution

## Projection

### Exercise #2

Which of the following relational algebra expression resulted in the output below?

name	manager	start_year	end_year
BigAI	Judy	2020	2025
FastCash	Judy	2018	2025
GlobalDB	Jack	2019	2023
CoreOS	Judy	2020	2020
CoolCoin	Jack	2015	2020

manager	name
Judy	FastCash
Jack	GlobalDB
Jack	CoolCoin

### Choice

### Comment

A	$\sigma_{[start\_year \leq 2019]}(\pi_{[manager, name]}(Projects))$	?
B	$\pi_{[manager, name]}(\sigma_{[start\_year < 2020]}(Projects))$	?
C	$\pi_{[manager, name]}(\sigma_{[manager='Jack']}(Projects))$	?
D	$\pi_{[name, manager]}(\sigma_{[start\_year \leq 2019]}(Projects))$	?



# Basic Operators

Selection  
Projection  
**Renaming**  
- Syntax  
- Semantics  
- Example  
Caution

## Renaming

### Syntax

$\rho_{[\mathcal{R}]}(R)$  where  $\mathcal{R}$  is a collection of attribute renaming which is either:

1.  $\mathcal{R}$  is an **unordered** collection of  $B_i \leftarrow A_i$ 
  - **Example:**  $\rho[B_1 \leftarrow A_1, B_2 \leftarrow A_2](R)$
  - The order of attributes does not matter
2.  $\mathcal{R}$  is an **ordered** list of attributes
  - **Example:**  $\rho(B_1, B_2, \dots, B_n)(R)$
  - The order of attributes matters
  - The list must specify **all** attributes in the original schema
    - If not renamed, set  $B_i = A_i$

new  $\leftarrow$  old

We will follow **option (1)** (*unordered collection of arrows (new  $\leftarrow$  old)*).

#### Note:

- Renaming will be relevant for set and join operations

# Basic Operators

Selection  
Projection  
**Renaming**  
- Syntax  
- Semantics  
- Example  
Caution

## Renaming

### Semantics

$\rho_{[\mathfrak{R}]}(R)$  renames the attributes listed in  $\mathfrak{R}$ .

### Visualization

A <sub>1</sub>	A <sub>2</sub>	A <sub>3</sub>	A <sub>4</sub>

$$\rho_{[B_1 \leftarrow A_1]}(R)$$

B <sub>1</sub>	A <sub>2</sub>	A <sub>3</sub>	A <sub>4</sub>

### Properties

- The resulting schema is modified by  $\mathfrak{R}$
- The order columns is unchanged (*modulo renaming*)
- The number of rows remains the same

# Basic Operators

Selection  
Projection  
**Renaming**  
- Syntax  
- Semantics  
- Example  
Caution

## Renaming

### Example

Rename the teams such that the project name becomes simple "project".

$\$T1 := \rho_{pname \leftarrow project}(Teams)$		
ename	project	hours
"Sarah"	"BigAI"	10
"Sam"	"BigAI"	5
"Bill"	"BigAI"	15
"Judy"	"GlobalDB"	20
"Max"	"GlobalDB"	5
"Sarah"	"GlobalDB"	10
"Emma"	"GlobalDB"	35
"Max"	"CoreOS"	40
"Bill"	"CoreOS"	30
"Sam"	"CoolCoin"	40
"Sarah"	"CoolCoin"	25
"Emma"	"CoolCoin"	10
12 Rows		

Teams		
ename	pname	hours
"Sarah"	"BigAI"	10
"Sam"	"BigAI"	5
"Bill"	"BigAI"	15
"Judy"	"GlobalDB"	20
"Max"	"GlobalDB"	5
"Sarah"	"GlobalDB"	10
"Emma"	"GlobalDB"	35
"Max"	"CoreOS"	40
"Bill"	"CoreOS"	30
"Sam"	"CoolCoin"	40
"Sarah"	"CoolCoin"	25
"Emma"	"CoolCoin"	10
12 Rows		

# Basic Operators

$$B \leftarrow \underline{A}, c \leftarrow B$$

Selection  
Projection  
Renaming  
**Caution**

## Caution

### Selection

$$\sigma_{[c]}(R)$$

The condition  $c$  must specify only attributes in  $R$

### Projection

$$\pi_{[\ell]}(R)$$

The ordered list of attributes  $\ell$  must specify only attributes in  $R$

### Renaming

$$\rho_{[\mathfrak{R}]}(R)$$

No two different attributes may be renamed to the same name

No attributes may be renamed more than once in a single operation

# Set Operators

---

# Set Operators

$$\{(1, 2), (3, 1)\} \cup \{('A', 2022-01-01, 3)\}$$



## Operations

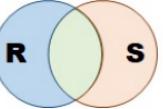
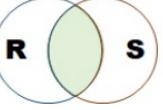
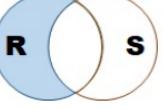
### - Operators

- Compatibility

- Example

Product

## Operations

Operation	Operator	Visualization	Description
Union	$R \cup S$		A relation containing all tuples that are in $R$ or $S$
Intersection	$R \cap S$		A relation containing all tuples that are in $R$ and $S$
Difference	$R - S$		A relation containing all tuples that are in $R$ but not in $S$

### Problem

What should be the result of union between  $\{1, 2\}$  and  $\{'A', 'B', 'C'\}$ ?



# Set Operators

## Operations

- Operators
  - Compatibility
  - Example
- Product

## Operations

### Union Compatibility

#### Definition

Two relations R and S are **union-compatible** if it satisfies both of the following:

- R and S have the same number of attributes
- The corresponding attributes have the same or compatible domains

#### Note:

- Just because two relations are union-compatible, does not mean the set operation is meaningful

### Example

#### Union-Incompatible

- Employees(name: TEXT, age: INT, role: TEXT )
- Teams(ename: TEXT, pname: TEXT , hours: INT )



#### Union-Compatible

- Employees(name: TEXT, role: TEXT, age: INT)
- Teams(ename: TEXT, pname: TEXT, hours: INT)



# Set Operators

## Operations

- Operators

- Compatibility

- Example

Product

## Operations

### Example

Find all projects that Bill is working on but Sarah is not working on.

$\$T5 := \$T2 - \$T4$	
	pname
	"CoreOS"
1 Rows	

$\$T2 := \pi_{pname}(\$T1)$	
	pname
	"BigAI"
	"CoreOS"
2 Rows	

$\$T1 := \sigma_{(ename = 'Bill')}(Teams)$		
ename	pname	hours
"Bill"	"BigAI"	15
"Bill"	"CoreOS"	30
2 Rows		

$\$T4 := \pi_{pname}(\$T3)$	
	pname
	"BigAI"
	"GlobalDB"
	"CoolCoin"
3 Rows	

$\$T3 := \sigma_{(ename = 'Sarah')}(Teams)$		
ename	pname	hours
"Sarah"	"BigAI"	10
"Sarah"	"GlobalDB"	10
"Sarah"	"CoolCoin"	25
3 Rows		



# Set Operators

Operations

Product

- Definition

- Example

- Exercise

- Discussion

$$A \times (B \times C)$$

$\equiv$

$$(A \times B) \times C$$

## Cross Product

### Definition

The **cross product** of two relations ( $R \times S$ ) is a relation formed by combining all pairs of tuples (i.e., rows) from the two input relations.

More formally, let  $R(A_1, A_2, \dots, A_n)$  and  $S(B_1, B_2, \dots, B_n)$ , then

- $R \times S$  produce a schema  
 $(R \times S)(A_1, A_2, \dots, A_n, B_1, B_2, \dots, B_n)$
- $R \times S = \{(a_1, a_2, \dots, a_n, b_1, b_2, \dots, b_n) \mid (a_1, a_2, \dots, a_n) \in R \wedge (b_1, b_2, \dots, b_n) \in S\}$

### Note:

- The set of attributes in  $R$  and  $S$  must be **disjoint**  
 $(Attr(R) \cap Attr(S) = \emptyset)$ , where
  - $Attr(R)$  is the set of attributes in  $R$
  - $Attr(S)$  is the set of attributes in  $S$

$$\{1, 2\} \times \{a, b\} = \{(1, a), (1, b)\}$$

$\cdots \} \times$

Example  $((1, a), \chi)$

R		
A	B	C
1	5	"m"
2	3	"f"

2 Rows

ST1 := R × S				
	A	B	C	X
1	5	"m"	"a"	30
1	5	"m"	"b"	10
1	5	"m"	"c"	20
2	3	"f"	"a"	30
2	3	"f"	"b"	10
2	3	"f"	"c"	20

6 Rows

S	
X	Y
"a"	30
"b"	10
"c"	20

3 Rows

# Set Operators

Operations

Product

- Definition

- Example

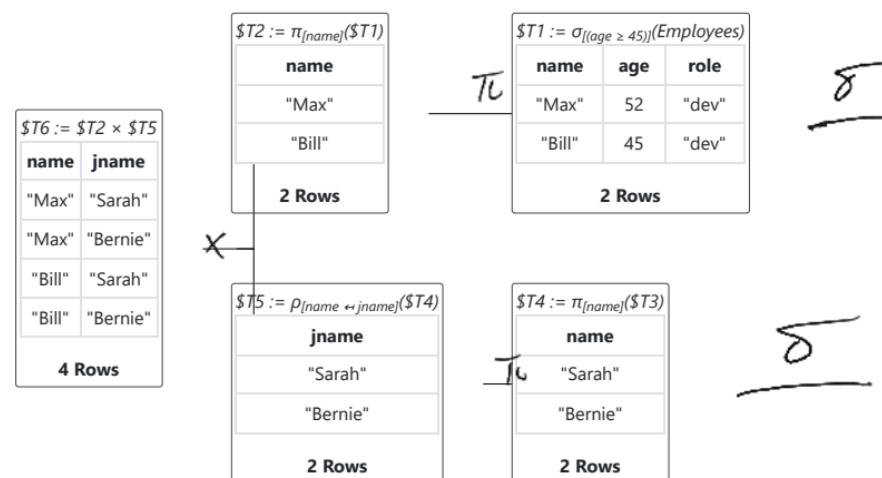
- Exercise

- Discussion

## Cross Product

### Example

Find all pairs of senior employee names (i.e.,  $age \geq 45$ ) and junior employee names (i.e.,  $age \leq 25$ ).



# Set Operators

Operations

**Product**

- *Definition*

- *Example*

- **Exercise**

- *Discussion*

## Cross Product

### Exercise #3

For all the project names, find the offices of the managers.

# Set Operators

Join

Operations

Product

- Definition

- Example

- Exercise

- Discussion

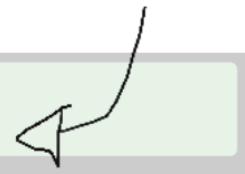
## Cross Product

Discussion

Observation

- Given two relations  $R$  and  $S$ , the size of the cross product is  $|R| \times |S|$
- In practice, many to most queries requiring a cross product also require
  - Selection operation to remove unnecessary rows
  - Projection to remove unnecessary/duplicate columns (*optional*)

$$\pi_{[name, office]}(\sigma_{[manager = mname]}(Projects \times \rho_{[mname \leftarrow name]}(Managers)))$$



## Join Operators

- Simplify relational algebra expression
  - Combine cross product, selection, and (*optionally*) projection
- Avoid generating all  $|R| \times |S|$  intermediate tuples

$$R \bowtie[\theta] S$$

# Join Operators

---

# Join Operators

Preliminary  
Inner Joins  
Outer Joins

## Preliminary

### Basic Idea

- Combine cross product, selection, and (*optionally*) projection into a single operator
- Typically results in simpler relational algebra expressions when formulating queries

### Base Types

**Join** (*for each element in R, combine with each element in S*)

#### Inner Join

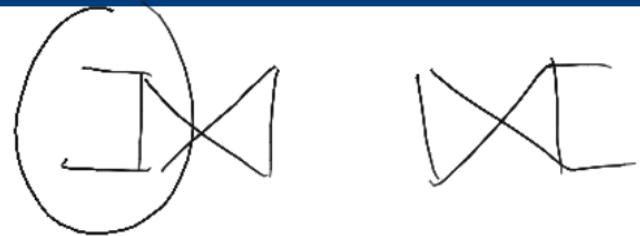
 Includes **only** tuples that satisfy the condition

- $\bowtie_{[\theta]}$  **θ-join**
- $\bowtie_{=}$  **equi join**
- $\bowtie$  **natural join**

#### Outer Join

 Also include tuples that do **not** satisfy the condition

-  **left outer join**
-  **right outer join**
-  **full outer join**



# Join Operators

Preliminary

Inner Joins

-  $\theta$ -Join

- Equi Join

- Natural Join

Outer Joins

## Inner Joins

$\theta$ -Join

### Definition

The  $\theta$ -join ( $R \bowtie_{[\theta]} S$ ) of two relations  $R$  and  $S$  is defined as:  $R \bowtie_{[\theta]} S = \sigma_{[\theta]}(R \times S)$

#### Note:

- The condition  $\theta$  can use attributes that appears in  $R$  or  $S$

### Example

For all the projects, find the offices of the managers.

Projects  $\bowtie_{[\text{manager}=\text{mname}]}$   
( $\rho_{[\text{mname}=\text{name}]}(\text{Managers})$ )

$\$T2 := \text{Projects } \bowtie_{[\text{manager} = \text{mname}]} \$T1$					
name	manager	start_year	end_year	mname	office
"BigAI"	"Judy"	2020	2025	"Judy"	"#03-20"
"FastCash"	"Judy"	2018	2025	"Judy"	"#03-20"
"GlobalDB"	"Jack"	2019	2023	"Jack"	"#03-10"
"CoreOS"	"Judy"	2020	2020	"Judy"	"#03-20"
"CoolCoin"	"Jack"	2015	2020	"Jack"	"#03-10"

5 Rows

# Join Operators

Preliminary

Inner Joins

-  $\theta$ -Join

- Equi Join

- Natural Join

Outer Joins

## Inner Joins

Equi Join

### Definition

Equi join ( $R \bowtie_= S$ ) is a special  $\theta$ -join where the only relational operator that can be used is equality (e.g., = or  $\equiv$ )

#### Note:

- $\theta$ -Join can use arbitrary relational operator (e.g., =,  $<$ ,  $>$ ,  $\leq$ ,  $\geq$ ,  $\equiv$ )
- Dedicated operation because attribute selection using equality operator are the most common
  - Especially when performing inner join following foreign key constraints
  - In many cases, hash-tables can be used to increase speed

# Join Operators

Preliminary

## Inner Joins

-  $\theta$ -Join

- Equi Join

- Natural Join

## Outer Joins

## Inner Joins

### Natural Join

$R \bowtie S$

#### Definition

Let  $R(A_1, A_2, \dots, A_i, B_1, B_2, \dots, B_j, \dots, B_l)$  and  $S(B_1, B_2, \dots, B_j, C_1, C_2, \dots, C_k)$ . We define **natural join** ( $R \bowtie S$ ):

$$R \bowtie S = \{(a_1, a_2, \dots, a_i, b_1, b_2, \dots, b_j, c_1, c_2, \dots, c_k) \mid (a_1, a_2, \dots, a_i, b_1, b_2, \dots, b_j) \in R \wedge (b_1, b_2, \dots, b_j, c_1, c_2, \dots, c_k) \in S\}$$

In other words:

- The join is performed over all common attributes between  $R$  and  $S$  (i.e.,  $(B_1, B_2, \dots, B_j)$ )
  - We only combine if the value of the common attributes are equal
  - The output relations contains the common attribute of  $R$  and  $S$  only once (via projection)

Alternatively,  $R \bowtie S = \pi_{[\ell]}(R \bowtie_{[\theta]} S)$ , where:

- $\ell = \text{Attr}(R) \cup (\text{Attr}(S) - \text{Attr}(R))^*$
- $\theta = \forall A_i \in \text{Attr}(R) \cap \text{Attr}(S): R_{A_i} = S_{A_i}$

\*Equivalent to  $\text{Attr}(R) \cup \text{Attr}(S)$  except that we want to show it more explicitly all attribute in  $S$  appears only after  $R$ .

# Join Operators

Preliminary

Inner Joins

Outer Joins

- Motivation

- Informal Steps

- Definition

- Natural

- Exercise

## Outer Joins

### Motivation

- Inner joins *eliminate* all tuples that do not satisfy the condition  
(i.e., *selection operation*)
- Sometimes the tuple in R or S that do **NOT** match are also of interest  
(i.e., *dangling tuple*)

### Example

Find all employees that are not assigned to any project.

### Note

Inner join can only find all employees that are assigned to at least one project.

Table "Teams"

ename	pname	hours
Sarah	BigAI	10
Sam	BigAI	5
Bill	BigAI	15
Judy	GlobalDB	20
Max	GlobalDB	5
Sarah	GlobalDB	10
Emma	GlobalDB	35
Max	CoreOS	40
Bill	CoreOS	30
Sam	CoolCoin	40
Sarah	CoolCoin	25
Emma	CoolCoin	10

Table "Employees"

name	age	role
Sarah	25	dev
Judy	35	sales
Max	52	dev
Marie	36	hr
Sam	30	sales
Bernie	19	NULL
Emma	28	dev
Jack	40	dev
Bill	45	dev

# Join Operators

Preliminary

Inner Joins

**Outer Joins**

- Motivation

- **Informal Steps**

- Definition

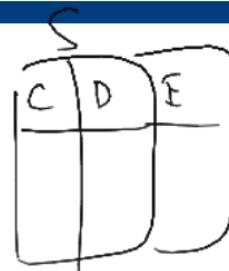
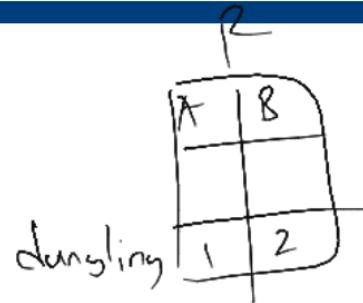
- Natural

- Exercise

## Outer Joins

Informal Steps

1. Perform inner join  $M = R \bowtie_{\theta} S$
2. Add dangling tuples to M from:
  - R in case of a **left outer join**
  - S in case of a **right outer join**
  - R and S in case of a **full outer join**
3. "Pad" missing attribute values of dangling tuples with NULL



# Join Operators

Preliminary

Inner Joins

**Outer Joins**

- Motivation

- Informal Steps

- **Definition**

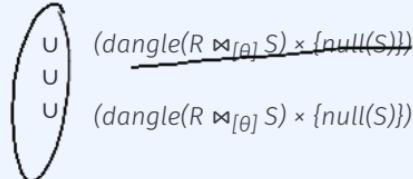
- Natural

- Exercise

## Outer Joins

### Definition

- $R \bowtie_{[\theta]} S = R \bowtie_{[\theta]} S$
- $R \bowtie_{[\theta]} S = R \bowtie_{[\theta]} S$
- $R \bowtie_{[\theta]} S = R \bowtie_{[\theta]} S$



where

- $dangle(R \bowtie S)$  = set of dangling tuples in  $R$  with respect to  $R \bowtie S$
- $null(R)$  =  $n$ -component tuple of NULL values where  $n$  is the number of attributes of  $R$

### Example

Find all employees that are not assigned to any project.

### Possible Solution

$\sigma_{[ename = \text{NULL}]}(Employees \bowtie_{[name=ename]} Teams)$

### Note

There are other solutions involving set operations as well.

# Join Operators

Preliminary

Inner Joins

## Outer Joins

- Motivation

- Informal Steps

- Definition

- Natural

- Exercise

## Outer Joins

### Natural Outer Joins

Analog to natural (*inner*) join

- The condition  $\theta$  is not specified
- Equality operation is performed over all attributes that  $R$  and  $S$  have in common

### Kinds

- **Natural left outer join**  $R \bowtie S$
- **Natural right outer join**  $R \bowtie\! S$
- **Natural full outer join**  $R \bowtie\!\! S$

# Join Operators

Preliminary

Inner Joins

Outer Joins

- Motivation

- Informal Steps

- Definition

- Natural

- Exercise

## Outer Joins

### Exercise #4

Recap the relations "Teams" and "Managers" on the right. How many **row** and **columns** are in the result of the relational algebra expression below?

$$\sigma[\text{ename} \equiv \text{NULL}] (\text{Managers} \bowtie [\text{name} = \text{ename}] \text{Teams})$$

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

Table "Teams"

ename	pname	hours
Sarah	BigAI	10
Sam	BigAI	5
Bill	BigAI	15
Judy	GlobalDB	20
Max	GlobalDB	5
Sarah	GlobalDB	10
Emma	GlobalDB	35
Max	CoreOS	40
Bill	CoreOS	30
Sam	CoolCoin	40
Sarah	CoolCoin	25
Emma	CoolCoin	10

Table "Managers"

name	office
Judy	#03-20
Jack	#03-10

dangling

### Choice

### Comment

A	1 row, 5 columns	?
B	3 rows, 5 columns	?
C	1 row, 3 columns	?
D	3 rows, 3 columns	?



# Complex Expressions

---

# Complex Expressions

## Motivation

- Question

- Tree

Equivalence

Alternatives

## Motivation

### Question

Find all managers (*with their offices*) of projects that started in 2020 or later, where at least one member of the project team has to work 30 hours or more on that project per week.

# Complex Expressions

## Motivation

- Question

- Tree

Equivalence

Alternatives

## Motivation

### Question

Find all managers (*with their offices*) of projects that started in 2020 or later, where at least one member of the project team has to work 30 hours or more on that project per week.

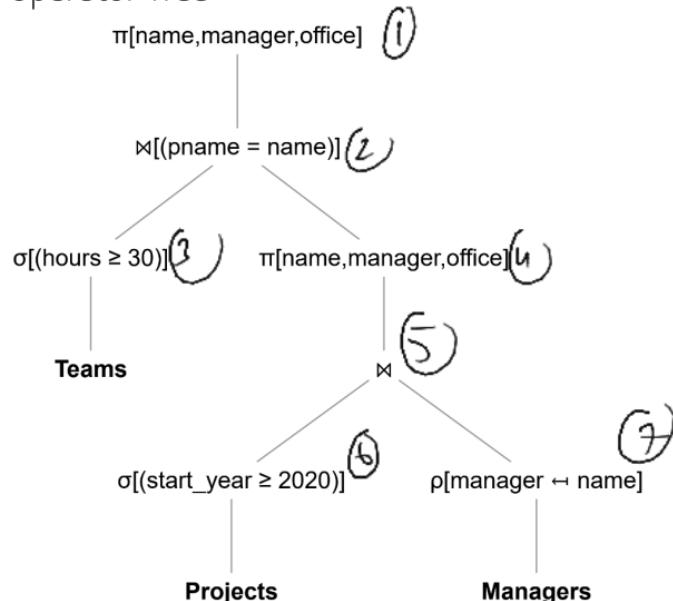
### Possible Solution #1

```
M := ρ[manager ← name](Managers)
P := σ[start_year ≥ 2020](Projects)
T := σ[hours ≥ 30](Teams)
Q1 := π[name, manager, office](P ⋈ M)
Q2 := T ⋈ [pname = name] Q1
Q := π[name, manager, office](Q2)
```



```
π[name, manager, office](
  σ[hours ≥ 30](Teams)
  ⋈ [pname = name]
  π[name, manager, office](
    σ[start_year ≥ 2020](Projects)
    ⋈
    ρ[manager ← name](Managers)))
```

### Operator Tree



# Complex Expressions

## Motivation

- Question

- Tree

Equivalence

Alternatives

## Motivation

### Question

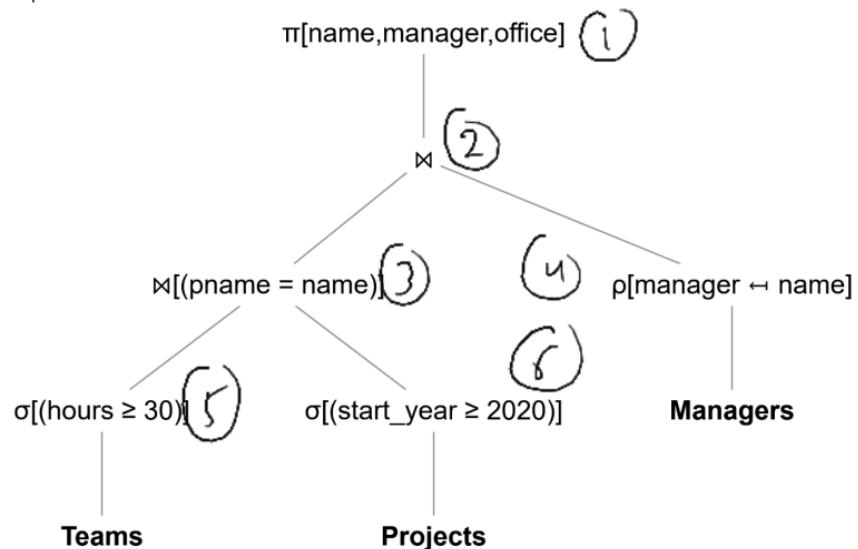
Find all managers (*with their offices*) of projects that started in 2020 or later, where at least one member of the project team has to work 30 hours or more on that project per week.

### Possible Solution #2

```
M := ρ[manager ← name](Managers)
P := σ[start_year ≥ 2020](Projects)
T := σ[hours ≥ 30](Teams)
Q1 := T ⋈ [pname = name] P
Q2 := Q1 ⋈ M
Q := π[name, manager, office](Q2)
```

```
π[name, manager, office](
    σ[hours ≥ 30](Teams)
        ⋈ [pname = name]
            σ[start_year ≥ 2020](Projects)
                ⋈
                ρ[manager ← name](Managers)
)
```

### Operator Tree



# Complex Expressions

Motivation  
**Equivalence**  
- Definition  
- Strong  
- Weak  
Alternatives

## Equivalence

### Definition

We say that two relational algebra expressions  $Q_1$  and  $Q_2$  are **equivalent** (*denoted by  $Q_1 \equiv Q_2$* ) if for any input relations either:

both produces error or both produces the same result

**Note:**

- This may be called **strongly** equivalent
- **Weakly** equivalent can be defined as if there is no error then both produce the same result

## Multiplicity

- In general, there are multiple ways to formulate a query to get the same result
- This includes:
  - Order in which join operations are performed
  - Order in which selection operations are performed (*e.g., before or after join operators*)
  - Inserting additional projections to minimize intermediate results
- **Query optimization:** finding the "best" relational algebra expression (*CS3223*)
  - Handled by the DBMS transparent to the user

# Complex Expressions

## Motivation Equivalence

- Definition
  - Strong
  - Weak
- Alternatives

## Equivalence

**Strong Equivalence** (*both cause errors or both return the same result*)

### Selection

- $\sigma_{[c_1]}(\sigma_{[c_2]}(R)) \equiv \sigma_{[c_2]}(\sigma_{[c_1]}(R))$
- $\sigma_{[c_1]}(\sigma_{[c_2]}(R)) \equiv \sigma_{[c_1 \wedge c_2]}(R)$

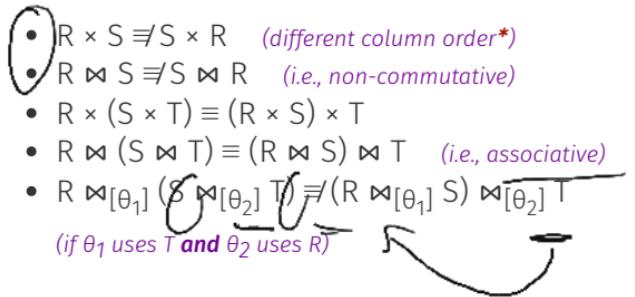
### Projection

- $\pi_{\ell_1}(\pi_{\ell_2}(R)) \not\equiv \pi_{\ell_1}(R)$     (unless  $\ell_1 \subseteq \ell_2$ )

### Combined

- $\pi_{[\ell]}(\sigma_{[\theta]}(R)) \not\equiv \sigma_{[\theta]}(\pi_{[\ell]}(R))$     (unless  $\theta$  uses only attributes in  $\ell$ )
- $\sigma_{[\theta]}(R \times S) \not\equiv \sigma_{[\theta]}(R) \times S$     (unless  $\theta$  uses only attributes in  $R$ )

### Cross Product and Joins

- $R \times S \not\equiv S \times R$     (different column order\*)
- $R \bowtie S \not\equiv S \bowtie R$     (i.e., non-commutative)
- $R \times (S \times T) \equiv (R \times S) \times T$
- $R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$     (i.e., associative)
- $R \bowtie_{[\theta_1]} (S \bowtie_{[\theta_2]} T) \not\equiv (R \bowtie_{[\theta_1]} S) \bowtie_{[\theta_2]} T$   
(if  $\theta_1$  uses  $T$  and  $\theta_2$  uses  $R$ )


\*You may find some people allow this to be equivalent but we do not.



# Complex Expressions

## Motivation Equivalence

- Definition
- Strong
- Weak

## Alternatives

## Equivalence

Weak Equivalence (if both do not cause errors then both return the same result)

### Selection

- $\sigma_{[c_1]}(\sigma_{[c_2]}(R)) \equiv \sigma_{[c_2]}(\sigma_{[c_1]}(R))$
- $\sigma_{[c_1]}(\sigma_{[c_2]}(R)) \equiv \sigma_{[c_1 \wedge c_2]}(R)$

### Projection

- $\pi_{\ell_1}(\pi_{\ell_2}(R)) \equiv \pi_{\ell_1}(R)$

—

### Combined

- $\pi_{[\ell]}(\sigma_{[\theta]}(R)) \equiv \sigma_{[\theta]}(\pi_{[\ell]}(R))$
- $\sigma_{[\theta]}(R \times S) \equiv \sigma_{[\theta]}(R) \times S$

### Cross Product and Joins

- $R \times S \not\equiv S \times R$  (*different column order\**)
- $R \bowtie S \not\equiv S \bowtie R$  (*i.e., non-commutative*)
- $R \times (S \times T) \equiv (R \times S) \times T$
- $R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$  (*i.e., associative*)
- $R \bowtie_{[\theta_1]} (S \bowtie_{[\theta_2]} T) \equiv (R \bowtie_{[\theta_1]} S) \bowtie_{[\theta_2]} T$

\*You may find some people allow this to be equivalent but we do not.

# Complex Expressions

Motivation  
Equivalence  
**Alternatives**  
- Invalid  
- Redundant

## Alternatives

Invalid Relational Expressions

Attribute no longer available after projection

$$\sigma_{\underline{\underline{role}} = 'dev'}(\pi_{name, age}(Employees))$$

Attribute no longer available after renaming

$$\sigma_{role = 'dev'}(\rho_{position \leftarrow role}(Employees))$$

Incompatible attribute types

$$\sigma_{age = role}(Employees)$$


We assume that there is no implicit type conversion (e.g., *INT* to *TEXT*)

# Complex Expressions

Motivation  
Observation  
**Alternatives**  
- Invalid  
- Redundant

## Alternatives

Valid but with Redundancy

Cross product + attribute selection  $\equiv$  join

$$\begin{aligned}\sigma_{[ \text{manager} = \text{mname} ]}(\text{Projects} \times \rho_{[\text{mname} \leftarrow \text{name}]}(\text{Managers})) \\ \equiv \text{Projects} \bowtie [\text{manager} = \text{mname}] \rho_{[\text{mname} \leftarrow \text{name}]}(\text{Managers})\end{aligned}$$

## Unnecessary operators

$$\begin{aligned}\pi_{[ \text{name} ]}(\pi_{[ \text{age} ]}(\text{Employees})) \\ \equiv \pi_{[ \text{name} ]}(\text{Employees})\end{aligned}$$

## Unoptimized query

$$\begin{aligned}\sigma_{[\text{start\_year} = 2020]}(\text{Projects} \bowtie_{[\text{manager} = \text{mname}]} \rho_{[\text{mname} \leftarrow \text{name}]}(\text{Managers})) \\ \equiv (\sigma_{[\text{start\_year} = 2020]}(\text{Projects})) \bowtie_{[\text{manager} = \text{mname}]} \rho_{[\text{mname} \leftarrow \text{name}]}(\text{Managers})\end{aligned}$$

## Quiz #2

Can you use natural join for the first and last operations?

$$\pi_{[A]}(\tau_{[B,C]}(R)) \quad \text{or}$$

$$\pi_{[A]}(R) \quad \text{no error}$$

# Summary

---

# Summary

Algebra

Unary

Set

Join

## Algebra

### What It Is

- Formal method to query relational data
- Closure property for arbitrarily complex relational expressions
- Basis for database query language such as SQL

### Operators

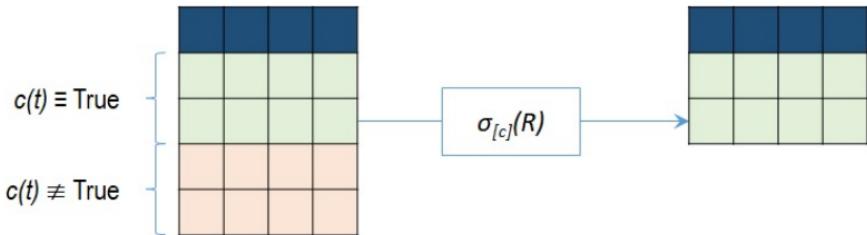
- **Unary** selection ( $\sigma$ ), projection ( $\pi$ ), renaming ( $\rho$ )
- **Binary** cross product ( $\times$ ), union ( $U$ ), intersection ( $\cap$ ), difference (-)

# Summary

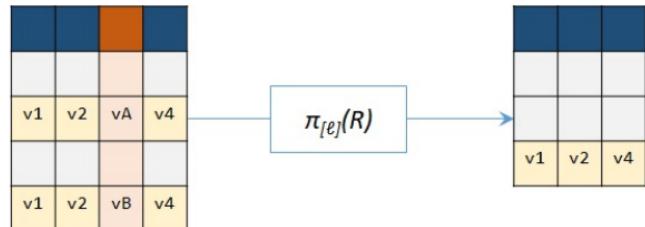
Algebra  
Unary  
Set  
Join

## Unary

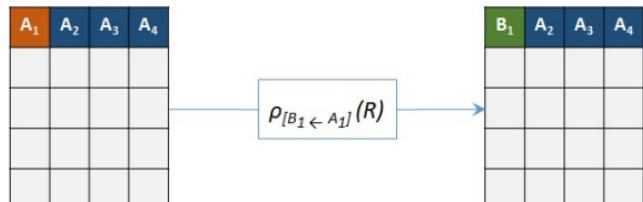
### Selection



### Projection



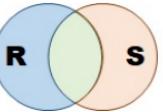
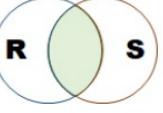
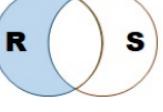
### Renaming



# Summary

Algebra  
Unary  
Set  
Join

## Set

Operation	Operator	Visualization	Description
Union	$R \cup S$		A relation containing all tuples that are in $R$ or $S$
Intersection	$R \cap S$		A relation containing all tuples that are in $R$ and $S$
Difference	$R - S$		A relation containing all tuples that are in $R$ but not in $S$

# Summary

Algebra  
Unary  
Set  
Join

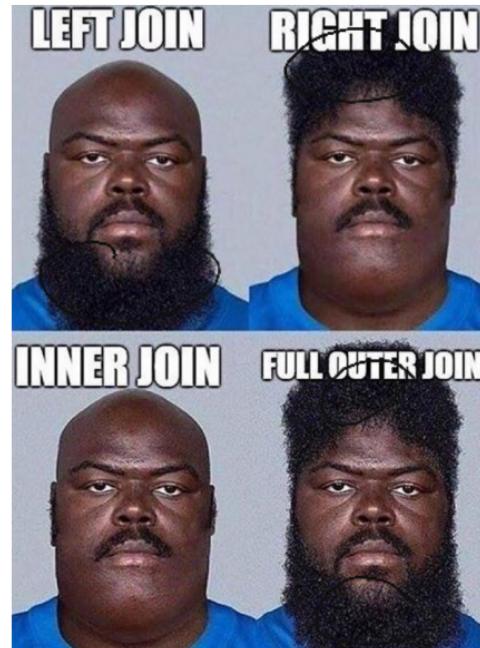
## Joins

L	Operator	R	Visualization
Keep dangling			
		Keep dangling	
Keep dangling		Keep dangling	

# Summary

Algebra  
Unary  
Set  
Join

## Joins



```
postgres=# exit
```

```
Press any key to continue . . .
```

# Solutions

## Quiz #1

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #2

## Preliminary

### Why Algebra?

- Allows for query optimization\*
- Rewrite a query such that:
  - It is *equivalent* to the original query
  - It can be evaluated faster than the original query

### Example

#### Question

Which of the following operation is "faster"?

- A.  $(a \times b) + (a \times c)$
- B.  $a \times (b + c)$

#### Quiz #1

Is it equivalent?

Depends on the type

- **YES** if
  - a:INT
  - b:INT
  - c:INT
- **NO** if
  - a:INT
  - b:TEXT
  - c:TEXT

\*While we do not penalize for inefficient queries, it is useful to know.

# Solutions

Quiz #1

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #2

## Selection

Exercise #1

Find all **employees** that **(a) do not work in Sales and are younger than 30** or **(b) work in HR**.

Solution

$$\sigma_{((role \neq 'sales') \wedge (age < 30)) \vee (role = 'hr'))}(\text{Employees})$$

$\$T1 := \sigma_{((role \neq 'sales') \wedge (age < 30)) \vee (role = 'hr'))}(\text{Employees})$

name	age	role
"Sarah"	25	"dev"
"Marie"	36	"hr"
"Emma"	28	"dev"

3 Rows

# Solutions

Quiz #1

Exercise #1

**Exercise #2**

Exercise #3

Exercise #4

Quiz #2

## Projection

### Exercise #2

Which of the following relational algebra expression resulted in the output below?

name	manager	start_year	end_year
BigAI	Judy	2020	2025
FastCash	Judy	2018	2025
GlobalDB	Jack	2019	2023
CoreOS	Judy	2020	2020
CoolCoin	Jack	2015	2020

manager	name
Judy	FastCash
Jack	GlobalDB
Jack	CoolCoin

### Choice

### Comment

A	$\sigma[\text{start\_year} \leq 2019](\pi[\text{manager}, \text{name}](\text{Projects}))$	NO: after $\pi$ , no $\text{start\_year}$	
B	$\pi[\text{manager}, \text{name}](\sigma[\text{start\_year} < 2020](\text{Projects}))$	YES: correct column order and correct condition	
C	$\pi[\text{manager}, \text{name}](\sigma[\text{manager} = 'Jack'](\text{Projects}))$	NO: missing Row 1	
D	$\pi[\text{name}, \text{manager}](\sigma[\text{start\_year} \leq 2019](\text{Projects}))$	NO: wrong column order	

# Solutions

Quiz #1

Exercise #1

Exercise #2

**Exercise #3**

Exercise #4

Quiz #2

## Cross Product

### Exercise #3

For all the project names, find the offices of the managers.

\$T4 := \pi_{[name, office]}(\$T3)	
name	office
"BigAI"	"#03-20"
"FastCash"	"#03-20"
"GlobalDB"	"#03-10"
"CoreOS"	"#03-20"
"CoolCoin"	"#03-10"

5 Rows

### Solution

$$\pi_{name, office}(\sigma_{manager=mname}(Projects \times (\rho_{mname \leftarrow name}(Managers))))$$

# Solutions

Quiz #1  
Exercise #1  
Exercise #2  
Exercise #3  
**Exercise #4**  
Quiz #2

## Outer Joins

### Exercise #4

Recap the relations "Teams" and "Managers" on the right. How many **row** and **columns** are in the result of the relational algebra expression below?

$$\sigma_{ename = \text{NULL}}(\text{Managers} \bowtie_{name=ename} \text{Teams})$$

Table "Teams"

ename	pname	hours
Sarah	BigAI	10
Sam	BigAI	5
Bill	BigAI	15
Judy	GlobalDB	20
Max	GlobalDB	5
Sarah	GlobalDB	10
Emma	GlobalDB	35
Max	CoreOS	40
Bill	CoreOS	30
Sam	CoolCoin	40
Sarah	CoolCoin	25
Emma	CoolCoin	10

Table "Managers"

name	office
Judy	#03-20
Jack	#03-10

### Choice

### Comment

A	1 row, 5 columns	YES: Only Jack has no match	<input checked="" type="checkbox"/>
B	3 rows, 5 columns	NO: Judy has no match	<input checked="" type="checkbox"/>
C	1 row, 3 columns	NO: The number of columns is the sum of both	<input checked="" type="checkbox"/>
D	3 rows, 3 columns	NO: The number of columns is the sum of both	<input checked="" type="checkbox"/>

# Solutions

Quiz #1

Exercise #1

Exercise #2

Exercise #3

Exercise #4

**Quiz #2**

## Alternatives

Valid but with Redundancy

Cross product + attribute selection  $\equiv$  join

$$\begin{aligned}\sigma_{[ \text{manager} = \text{mname} ]}(\text{Projects} \times \rho_{[\text{mname} \leftarrow \text{name}]}(\text{Managers})) \\ \equiv \text{Projects} \bowtie \rho_{[\text{manager} \leftarrow \text{name}]}(\text{Managers})\end{aligned}$$

## Unnecessary operators

$$\begin{aligned}\pi_{[ \text{name} ]}(\pi_{[ \text{name}, \text{age} ]}(\text{Employees})) \\ \equiv \pi_{[ \text{name} ]}(\text{Employees})\end{aligned}$$

## Unoptimized query

$$\begin{aligned}\sigma_{[\text{start\_year} = 2020]}(\text{Projects} \bowtie_{[\text{manager} = \text{mname}]} \rho_{[\text{mname} \leftarrow \text{name}]}(\text{Managers})) \\ \equiv (\sigma_{[\text{start\_year} = 2020]}(\text{Projects})) \bowtie \rho_{[\text{manager} \leftarrow \text{name}]}(\text{Managers})\end{aligned}$$

## Quiz #2

Can you use *natural join* for the first and last operations?

