

# CS2102: Database Systems

## SQL (Part 1)

# Announcement

---

# Announcement

## Project Group

- Deadline for group formation is this **Saturday, 28 January 2023** by 12:00 Noon
  - Groups with fewer than 4 members may be *merged* or *allocated* a random member

## Tutorial

- Starts this **Wednesday**
- Please prepare the material before hand
- Solutions will be released on Friday evening (**18:00**)

## Catch-Up Session

- **Session I:** Friday, 27 January 2023 at 18:00
  - Q&A Session
  - Will be recorded

# Recap

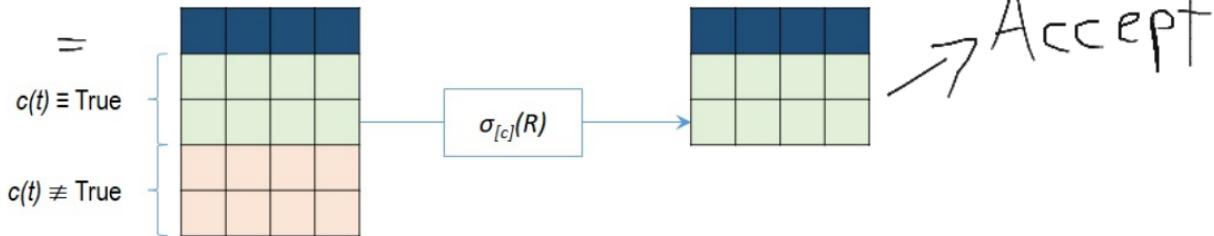
---

# Recap

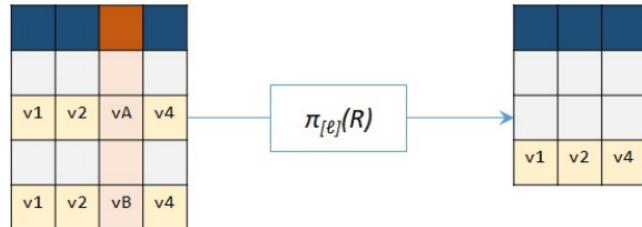
Unary  
Set  
Join

## Unary

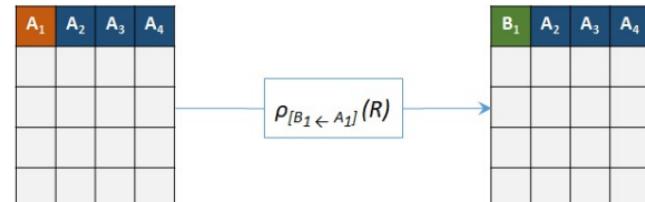
### Selection



### Projection



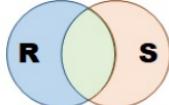
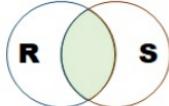
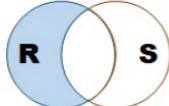
### Renaming



# Recap

Unary  
Set  
Join

## Set

Operation	Operator	Visualization	Description
Union	$R \cup S$		A relation containing all tuples that are in $R$ or $S$
Intersection	$R \cap S$		A relation containing all tuples that are in $R$ and $S$
Difference	$R - S$		A relation containing all tuples that are in $R$ but not in $S$

Union - Compatible

# Recap

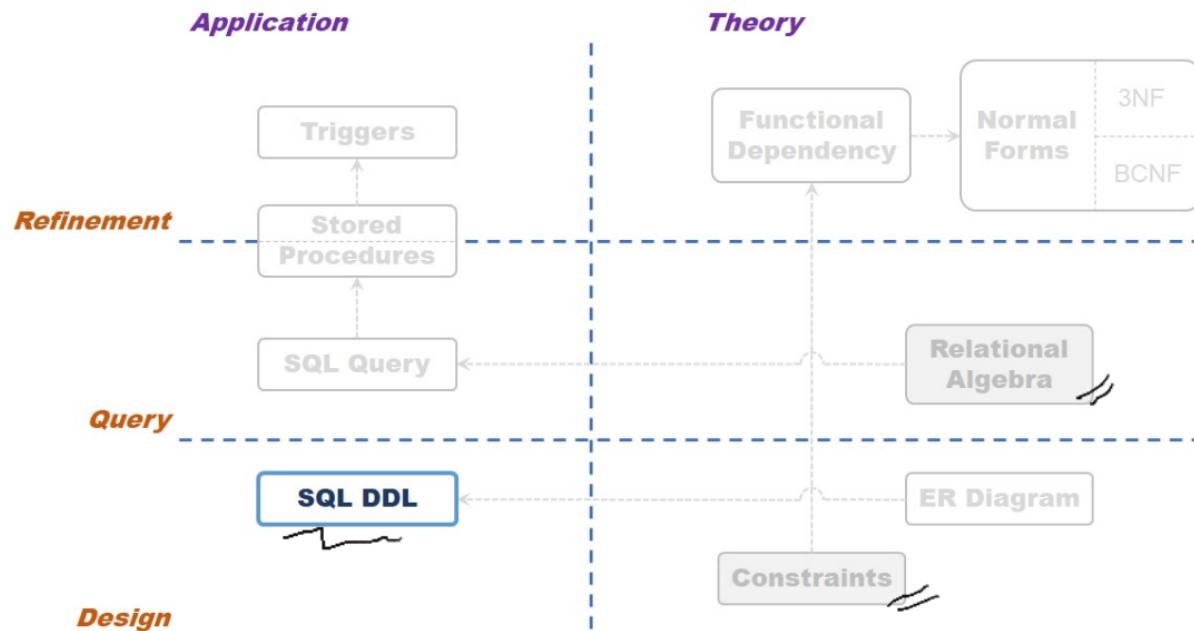
Unary  
Set  
Join

## Joins

satisfy condition

L	Operator	R	Visualization
X		X	
Keep dangling			
		Keep dangling	
Keep dangling		Keep dangling	

# Roadmap



# Roadmap

## Overview

### Overview

#### *SQL Overview*

- History and Usages

#### *Creating/Insert/Deleting*

- Basic DDL
- Basic DML

#### *Integrity Constraints*

- Unique and Keys
- General

#### *Modifying Database*

- Alter and Drop

#### *Deferrable Constraints*

# SQL Overview

---

# SQL Overview

SQL

Using SQL  
Commands

## SQL: Structured Query Language

What Is It?

- De-facto standard language to "talk" to RDBMS
  - Developed by Donald D. Chamberlin and Raymond F. Boyce (*IBM Research, 1974*)
  - Originally called *SEQUEL* (*Structured English Query Language*)
- SQL is a **domain-specific language** (*i.e., not a general-purpose like Python*)
- SQL is a **declarative language**
  - Focus on **what** to compute, instead of **how** to compute
  - Somewhat opposite to relational algebra expression (*which is imperative*)
- Based on *multi-set/bag*, (*may contain duplicate rows*) instead of set

SQL Standard



- First standard is SQL-86; the most recent standard is SQL-2019 (*new standard every ~3-5 years*)
- New standards introduce new language concepts (*e.g., support new features of RDBMS*)
- Many RDBMS add their own "flavor" to SQL



# SQL Overview

SQL

Using SQL

- Non-Interactive

- Interactive

Commands

## Using SQL

### Non-Interactive SQL

- Included in application written in a host language

#### Statement Level Interface (SLI)

- Application is a mixture of host language (e.g., C, Java, Python) and SQL statements
  - Example:** Embedded SQL, Dynamic SQL

#### Call Level Interface (CLI)

- Application is entirely written in the host language (e.g., C, Java, Python)
  - Example:** ODBC (Open DataBase Connectivity), JDBC (Java DataBase Connectivity)



# SQL Overview

SQL

Using SQL

- Non-Interactive

- Interactive

Commands

## Using SQL

Interactive SQL

| Directly writing SQL statements to an interface

Command Line

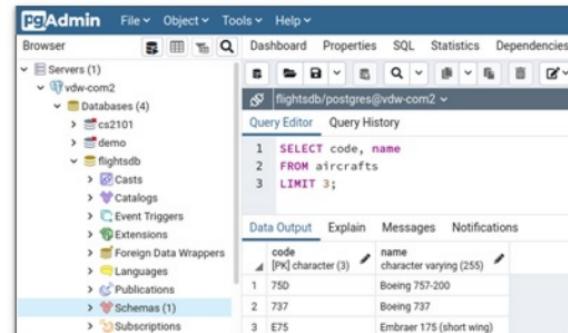
psql

```
List of relations
Schema | Name   | Type  | Owner
-----+-----+-----+-----
public | aircrafts | table | postgres
public | airports  | table | postgres
public | countries | table | postgres
public | flightcodes | table | postgres
public | flights   | table | postgres
(5 rows)

flightsdb=# SELECT code, name FROM aircrafts LIMIT 3;
 code |      name
-----+
 750 | Boeing 757-200
 737 | Boeing 737
 E75 | Embraer 175 (short wing)
(3 rows)

flightsdb=#
```

Graphical User Interface  
→ pgAdmin



| This will be our preferred mode

# SQL Overview

SQL  
Using SQL  
Commands

## Types of Commands/Statements

SQL

~~DDL~~

Data Definition

- CREATE
- ALTER
- DROP
- RENAME
- TRUNCATE

~~DML~~

Data Manipulation

- INSERT
- UPDATE
- DELETE
- MERGE

~~DQL~~

Data Query

- SELECT

DCL



Data Control

- GRANT
- REVOKE

TCL

Transaction Control

- BEGIN
- COMMIT
- ROLLBACK
- SAVEPOINT

# Creating/Insert/Deleting

---

# Creating/Insert/Delete

CREATE

- Syntax

- Semantics

- Example

Data Types

INSERT

DELETE

## CREATE TABLE

### Syntax

```
CREATE TABLE <table_name> (
    <attr1> <type1> [<column_constraint>],
    <attr2> <type2> [<column_constraint>],
    :
    <attrn> <typen> [<column_constraint>],
    [<table_constraint>], -- comment
    [<table_constraint>], /* comment */
    :
    [<table_constraint>] -- no comma
);
```

### Notes

- SQL is **case-insensitive**
  - You can use uppercase or lowercase
- **Convention:**
  - Keywords should be in uppercase
  - Try to align elements



cAn MiX uPpeR AnD LoWEr

# Creating/Insert/Delete

## CREATE

- Syntax

- Semantics

- Example

## Data Types

## INSERT

## DELETE

## CREATE TABLE

### Syntax

```
CREATE TABLE <table_name> (
    <attr1> <type1> [<column_constraint>],
    <attr2> <type2> [<column_constraint>],
    :
    <attrn> <typen> [<column_constraint>],
    [<table_constraint>], -- comment
    [<table_constraint>], /* comment */
    :
    [<table_constraint>] -- no comma
);
```

### Semantics

attr1	attr2	...	attrn

```
<table_name>(
    <attr1> : <type1>,
    <attr2> : <type2>,
    :
    <attrn> : <typen>
)
```

# Creating/Insert/Delete

## CREATE

- Syntax
- Semantics
- Example

## Data Types

## INSERT

## DELETE

## CREATE TABLE

### Example

Create the relation *Employees*(*id*: **INT**, *name*: **TEXT**, *age*: **INT**, *role*: **TEXT**).



# Creating/Insert/Delete

CREATE  
Data Types  
INSERT  
DELETE

## Data Types

### Basic Data Types

Type	Description
BOOLEAN	Logical Boolean ( <i>true/false</i> )
INTEGER (INT)	Signed 4-bytes integer
FLOAT8	Double precision 8-bytes floating-point
NUMERIC[( <i>p,s</i> )]	Exact numeric of selectable precision
CHAR( <i>n</i> )	Fixed-length character string
VARCHAR( <i>n</i> )	Variable-length character string
TEXT	Variable-length character string
DATE	Calendar date ( <i>year, month, day</i> )
TIMESTAMP	Date and time

### Extended Data Types

Kinds	Types
Document	XML JSON
Spatial	Point Line Polygon Circle Box Path
Special	Money/Currency MAC/IP Address

### User-Defined Types

[CREATE TYPE](#)

# Creating/Insert/Delete

CREATE  
Data Types  
**INSERT**  
- *Syntax*  
- *Semantics*  
- *Default*  
DELETE

## INSERT INTO

### Syntax

```
INSERT INTO <table_name> [(attr_1, attr_2, ..., attr_n)]  
VALUES  
    (<val_1_1>, <val_2_1>, ..., <val_n_1>),  
    (<val_1_2>, <val_2_2>, ..., <val_n_2>),  
    :  
    (<val_1_m>, <val_2_m>, ..., <val_n_m>);
```

### Notes

- Either **all** inserted or **none** inserted
- Attributes can be specified out-of-order (*optionally*)
- Missing values are replaced with **NULL** (*if allowed*) or default values (*if specified*)



# Creating/Insert/Delete

CREATE  
Data Types  
**INSERT**  
- Syntax  
- Semantics  
- Default  
DELETE

## INSERT INTO

### Semantics

Table "Employees"

id	name	age	role

```
INSERT INTO Employees
VALUES (101, 'Sarah', 25, 'dev');
```

id	name	age	role
101	'Sarah'	25	'dev'

```
INSERT INTO Employees (name, id)
VALUES ('Judy', 102), ('Max', 103);
```

id	name	age	role
101	'Sarah'	25	'dev'
102	'Judy'	NULL	NULL
103	'Max'	NULL	NULL

# Creating/Insert/Delete

CREATE  
Data Types  
**INSERT**  
- Syntax  
- Semantics  
- Default  
DELETE

## INSERT INTO

### Default

```
CREATE TABLE Employees (
    id    INTEGER,
    name  VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50) DEFAULT 'sales'
);
```

Table "Employees"

id	name	age	role

### Insertion

```
INSERT INTO Employees
VALUES (101, 'Sarah', 25, 'dev');
INSERT INTO Employees (name, id)
VALUES ('Judy', 102), ('Max', 103);
```

id	name	age	role
101	'Sarah'	25	'dev'
102	'Judy'	NULL	'sales'
103	'Max'	NULL	'sales'

# Creating/Insert/Delete

CREATE  
Data Types  
INSERT  
DELETE  
- Syntax  
- Semantics

## DELETE FROM

### Syntax

```
DELETE FROM <table_name>
[ WHERE <condition> ];
```



WHERE TRUE

### Note

- Condition is *optional*
  - If unspecified, equivalent to always true
  - If specified, condition can be arbitrarily complex

- ① Boolean
- ② Only involve
  - > this row
  - b) this file



# Creating/Insert/Delete

CREATE  
Data Types  
INSERT  
**DELETE**  
- Syntax  
- Semantics

## DELETE FROM

### Semantics

#### Principle of Acceptance

- Perform the operation if the condition evaluates to True\*.
- Used in WHERE clause (*and relational algebra*)

Table "Employees"

id	name	age	role
101	'Sarah'	25	'dev'
102	'Judy'	NULL	'sales'
103	'Max'	NULL	NULL



### Example

```
DELETE FROM Employees;  
-- delete all!
```

id	name	age	role

```
DELETE FROM Employees  
WHERE role <> 'dev';
```

id	name	age	role
101	'Sarah'	25	'dev'
103	'Max'	NULL	NULL

\*Remember, we have three-valued logic!

# Integrity Constraints

---

# Integrity Constraints

## Preliminary

- *Principles*
- *Three-Valued*
- *Types*

Not-NULL

Unique

Primary Key

Foreign Key

General

Summary

## Preliminary

Principles

### Principle of Rejection

**Reject** the insertion if the condition **evaluates to False**.

- Used in integrity constraints

# Integrity Constraints

## Preliminary

- Principles
- Three-Valued
- Types

Not-NULL

Unique

Primary Key

Foreign Key

General

Summary

## Preliminary



Three-Valued Logic

IS [NOT] NULL

$x \equiv \text{NULL}$

$x \neq \text{NULL}$

x	x IS NULL	x IS NOT NULL
not-NULL	False	True
NULL	True	False

IS [NOT] DISTINCT FROM

x	y	x IS DISTINCT FROM y	x IS NOT DISTINCT FROM y
not-NULL	not-NULL	$x \neq y$	$x = y$
not-NULL	NULL	True	False
NULL	not-NULL	True	False
NULL	NULL	False	True

# Integrity Constraints

## Preliminary

- Principles
- Three-Valued
- **Types**

Not-NULL  
Unique  
Primary Key  
Foreign Key  
General  
Summary

## Preliminary

### Types of Integrity Constraints

1. Not-NULL Constraints
2. Unique Constraints
3. Primary Key Constraints
4. Foreign Key Constraints
5. General Constraints

→ Look at another table

## Specification

### Where

- **Column Constraint:**
  - Applies to a single column\*
  - Specified at column definition
- **Table Constraint:**
  - Applies to one or more columns
  - Specified after column definition

### Naming

- **Named Constraint:**
  - Name is specified by user
- **Unnamed Constraint:**
  - Name is automatically assigned by DBMS

# Integrity Constraints

Preliminary

Not-NULL

- Column Constraint

- Exercise

Unique

Primary Key

Foreign Key

General

Summary

## Not-NULL Constraint

Column Constraint

Unnamed

```
CREATE TABLE Employees (
    id INT NOT NULL,
    name TEXT NOT NULL,
    age INT,
    role TEXT
);
```

Named

```
CREATE TABLE Employees (
    id INT CONSTRAINT nn_id NOT NULL,
    name TEXT CONSTRAINT nn_name NOT NULL,
    age INT,
    role TEXT
);
```

(1, NULL) X  
(NULL, Alice) X

# Integrity Constraints

Preliminary

**Not-NULL**

- Column Constraint

- Exercise

Unique

Primary Key

Foreign Key

General

Summary

## Not-NULL Constraint

Schema

Table "Employees"

```
CREATE TABLE Employees (
    id      INT NOT NULL,
    name    TEXT NOT NULL,
    age     INT,
    role    TEXT
);
```

	<b>id</b>	<b>name</b>	<b>age</b>	<b>role</b>	<b>Row</b>
	101	'Sarah'	25	'dev'	Row 1
	NULL	'Judy'	32	'sales'	Row 2
	102	"	27	'hr'	Row 3
	103	'Max'	NULL	NULL	Row 4

### Exercise #1

Consider the schema on the left and the table at the top. Which rows violate the **NOT NULL** constraint?

	<b>Choice</b>	<b>Comment</b>	
A	Row 1		?
B	Row 2		?
C	Row 3		?
D	Row 4		?

# Integrity Constraints

Preliminary

Not-NULL

**Unique**

- Column Constraint

- Table Constraint

- Exercise

Primary Key

Foreign Key

General

Summary

## Unique Constraint

Column Constraint

Unnamed

```
CREATE TABLE Employees (
    id      INT  UNIQUE,
    name   TEXT,
    age    INT,
    role   TEXT
);
```

Named

```
CREATE TABLE Employees (
    id      INT  CONSTRAINT emp_id UNIQUE,
    name   TEXT ,
    age    INT ,
    role   TEXT
);
```

$$\forall r_1, r_2 : r_1 \equiv r_2 \vee (\exists \underline{a} : r_1.a \neq r_2.a)$$

↓

*unique(Emp)*

# Integrity Constraints

Preliminary

Not-NULL

**Unique**

- Column Constraint

- Table Constraint

- Exercise

Primary Key

Foreign Key

General

Summary

## Unique Constraint

Table Constraint

Unnamed

```
CREATE TABLE Teams (
    eid INT,
    pname TEXT,
    hours INT,
    UNIQUE (eid, pname)
);
```

~~→~~ 

Named

```
CREATE TABLE Teams (
    eid INT,
    pname TEXT,
    hours INT,
    CONSTRAINT team_id UNIQUE (eid, pname)
);
```

(4, 'A', 10)

(1, 'B', 10)      (2, 'A', 10)

# Integrity Constraints

Preliminary

Not-NULL

Unique

- Column Constraint

- Table Constraint

- Exercise

Primary Key

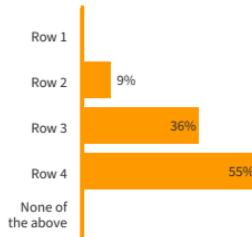
Foreign Key

General

Summary

Post locked. Responses not accepted.

Which rows are not UNIQUE?



Total Results: 11

## Constraints

Schema

Table "Teams"

```
CREATE TABLE Teams (
    eid    INT,
    pname  TEXT,
    hours  INT,
    UNIQUE (eid, pname)
);
```

eid	pname	hours	Row
101	NULL	25	Row 1
101	NULL	25	Row 2
102	'Judy'	32	Row 3
102	'Judy'	32	Row 4

NULL  $\leftrightarrow$  NULL

$\Rightarrow$  NULL

Exercise #2

Consider the schema on the left and the table at the top. Which rows are not UNIQUE?

## Choice

A	Row 1	
B	Row 2	
C	Row 3	
D	Row 4	

## Comment

?
?
?
?

# Integrity Constraints

Preliminary

Not-NULL

Unique

**Primary Key**

- *Recap*

- *Column Constraint*

- *Table Constraint*

- *Equivalent*

Foreign Key

General

Summary

## Primary Key Constraint

Recap

### Definition

A **primary key** is a selected candidate keys.

- Uniquely identifies a tuple in a relation
- Cannot be NULL

In other words, **UNIQUE** and **NOT NULL**

Indexing

# Integrity Constraints

Preliminary

Not-NULL

Unique

**Primary Key**

- Recap

- *Column Constraint*

- *Table Constraint*

- *Equivalent*

Foreign Key

General

Summary

## Primary Key Constraint

Column Constraint

Unnamed

```
CREATE TABLE Employees (
    id      INT  PRIMARY KEY,
    name   TEXT,
    age    INT,
    role   TEXT
);
```

Named

```
CREATE TABLE Employees (
    id      INT  CONSTRAINT emp_pk PRIMARY KEY,
    name   TEXT,
    age    INT,
    role   TEXT
);
```

# Integrity Constraints

Preliminary

Not-NULL

Unique

**Primary Key**

- Recap

- Column Constraint

- **Table Constraint**

- Equivalent

Foreign Key

General

Summary

## Primary Key Constraint

Table Constraint

Unnamed

```
CREATE TABLE Teams (
    eid    INT,
    pname TEXT,
    hours INT,
    PRIMARY KEY (eid, pname)
);
```

Named

```
CREATE TABLE Teams (
    eid    INT,
    pname TEXT,
    hours INT,
    CONSTRAINT team_pk PRIMARY KEY (eid, pname)
);
```

# Integrity Constraints

Preliminary

Not-NULL

Unique

**Primary Key**

- Recap

- Column Constraint

- Table Constraint

- **Equivalent**

Foreign Key

General

Summary

## Primary Key Constraint

Equivalent Constraint

Unnamed

```
CREATE TABLE Teams (
    eid INT NOT NULL,
    pname TEXT NOT NULL,
    hours INT,
    UNIQUE (eid, pname)
);
```

Named

```
CREATE TABLE Teams (
    eid INT CONSTRAINT nn_eid NOT NULL,
    pname TEXT CONSTRAINT nn_pname NOT NULL,
    hours INT,
    CONSTRAINT team_pk PRIMARY KEY (eid, pname)
);
```

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- *Recap*  
- *Constraint*  
- *Condition*  
- *Problems*  
- *Actions*  
- *Example*  
- *Consideration*  
- *Exercise*  
General Summary

## Foreign Key Constraint

Recap

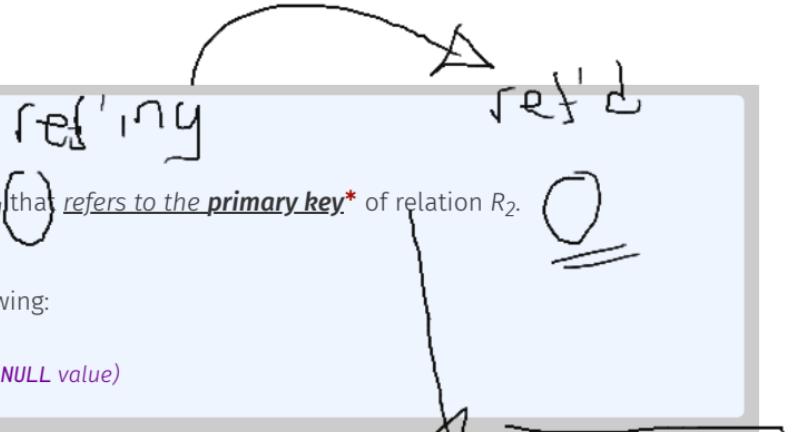
### Definition

A **foreign key** is a subset of attributes of relation  $R_1$  that refers to the primary key\* of relation  $R_2$ .

### Requirement:

Each foreign key in  $R_1$  must satisfy one of the following:

- Appear as **primary key** in  $R_2$
- Be a **NULL** value (or a tuple containing at least one **NULL** value)



What about  
NULL?

any unique  
constraint

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- Problems  
- Actions  
- Example  
- Consideration  
- Exercise  
General  
Summary

## Foreign Key Constraint

Constraint on Referencing Relation

```
CREATE TABLE Teams (
    eid      INT,
    pname   TEXT REFERENCES Projects,
    hours   INT,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id)
);
```



→ good practice

## Referenced Relations

```
CREATE TABLE Employees (
    id      INT PRIMARY KEY,
    name   TEXT,
    age    INT,
    role   TEXT
);
```

```
CREATE TABLE Projects (
    name      TEXT,
    start_year INT,
    end_year  INT,
    PRIMARY KEY (name)
);
```

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- Problems  
- Actions  
- Example  
- Consideration  
- Exercise  
General  
Summary

## Foreign Key Constraint

### Condition

A referential constraint is satisfied if one of the following conditions is true, depending on the <match option> specified in the <referential constraint definition>:

- If no <match type> was specified then, for each row R1 of the referencing table, either at least one of the values of the referencing columns in R1 shall be a null value, or the value of each referencing column in R1 shall be equal to the value of the corresponding referenced column in some row of the referenced table.
- If MATCH FULL was specified then, for each row R1 of the referencing table, either the value of every referencing column in R1 shall be a null value, or the value of every referencing column in R1 shall not be null and there shall be some row R2 of the referenced table such that the value of each referencing column in R1 is equal to the value of the corresponding referenced column in R2.

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- **Problems**  
- Actions  
- Example  
- Consideration  
- Exercise  
General Summary

## Foreign Key Constraint

### Problems

- What if the first tuple on "Movies" is **deleted**?
- What if Sigourney Weaver *id* is **updated** from 20 to 30?

### Idea

Extend the syntax to specify the *behavior* when data in referenced table is deleted or updated using *optional* specification:

- ON DELETE <action>
- ON UPDATE <action>

Table "Movies"

<u><i>id</i></u>	<u><i>title</i></u>	<u><i>genre</i></u>	<u><i>opened</i></u>
101	Aliens	Action	1986
102	Logan	Drama	2017
103	Heat	Crime	1995
104	Terminator	Action	1984

Table "Cast"

<u><i>movie_id</i></u>	<u><i>actor_id</i></u>	<u><i>role</i></u>
101	20	Ellen Ripley
101	23	Private Hudson
101	54	Corporal Hicks
102	21	Logan
104	23	Punk Leader

Table "Actors"

<u><i>id</i></u>	<u><i>name</i></u>	<u><i>dob</i></u>
20	Sigourney Weaver	08-10-1949
21	Hugh Jackman	12-10-1968
22	Tom Hanks	09-07-1956
23	Bill Paxton	17-05-1955

# Integrity Constraints

Preliminary

Not-NULL

Unique

Primary Key

**Foreign Key**

- Recap

- Constraint

- Condition

- Problems

**- Actions**

- Example

- Consideration

- Exercise

General

Summary

## Foreign Key Constraint

### Actions

Keyword	Action
NO ACTION	<u>Reject</u> delete/update if it violates constraint ( <i>default value</i> )
RESTRICT	Similar to "NO ACTION" except that check of constraint <u>cannot be deferred</u> ( <i>deferrable constraints are discussed in a bit</i> )
CASCADE	<u>Propagates</u> delete/update to the referencing tuples
SET DEFAULT	<u>Updates</u> the foreign key of the referencing tuples to some default value ( <b>Important:</b> <i>default value must be a primary key in the referencing table</i> )
SET NULL	<u>Updates</u> the foreign key of the referencing tuples to NULL value ( <b>Important:</b> <i>corresponding column must be allowed to contain NULL values</i> )

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- Problems  
- Actions  
- **Example**  
- Consideration  
- Exercise  
General  
Summary

## Foreign Key Constraint

### Example

```
CREATE TABLE Teams (
    eid      INT,
    pname   TEXT DEFAULT 'FastCash',
    hours   INT,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE NO ACTION ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON DELETE SET DEFAULT ON UPDATE CASCADE
);
```

### Effects

- Updates on *Employees.id* and *Projects.name* are propagated
- Deleting a project will set *Teams.pname* to default value (i.e., *FastCash*)
- Deleting an employee will raise an error if that employee is still assigned to a team

# Integrity Constraints

Preliminary

Not-NULL

Unique

Primary Key

**Foreign Key**

- Recap

- Constraint

- Condition

- Problems

- Actions

- Example

- Consideration

- Exercise

General

Summary

## Foreign Key Constraint

### Example

```
CREATE TABLE Teams (
    eid      INT,
    pname   TEXT DEFAULT 'FastCash',
    hours   INT,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE NO ACTION ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON DELETE SET DEFAULT ON UPDATE CASCADE
);
```

name	start_year	end_year
BigAI	2020	2025
FastCash	2018	2025
:	:	:

eid	pname	hours
101	BigAI	10
102	GlobalDB	20
:	:	:

```
UPDATE Projects
SET    name = 'SmartAI'
WHERE  name = 'BigAI'
```

name	start_year	end_year
SmartAI	2020	2025
FastCash	2018	2025
:	:	:

eid	pname	hours
101	SmartAI	10
102	GlobalDB	20
:	:	:

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- Problems  
- Actions  
- **Example**  
- Consideration  
- Exercise  
General Summary

## Foreign Key Constraint

### Example

```
CREATE TABLE Teams (
    eid      INT,
    pname   TEXT DEFAULT 'FastCash',
    hours   INT,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE NO ACTION ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON DELETE SET DEFAULT ON UPDATE CASCADE
);
```

name	start_year	end_year
BigAI	2020	2025
FastCash	2018	2025
:	:	:

```
DELETE FROM Projects
WHERE name = 'BigAI';
```

eid	pname	hours
101	BigAI	10
102	GlobalDB	20
:	:	:

name	start_year	end_year
FastCash	2018	2025
:	:	:



eid	pname	hours
101	FastCash	10
102	GlobalDB	20
:	:	:

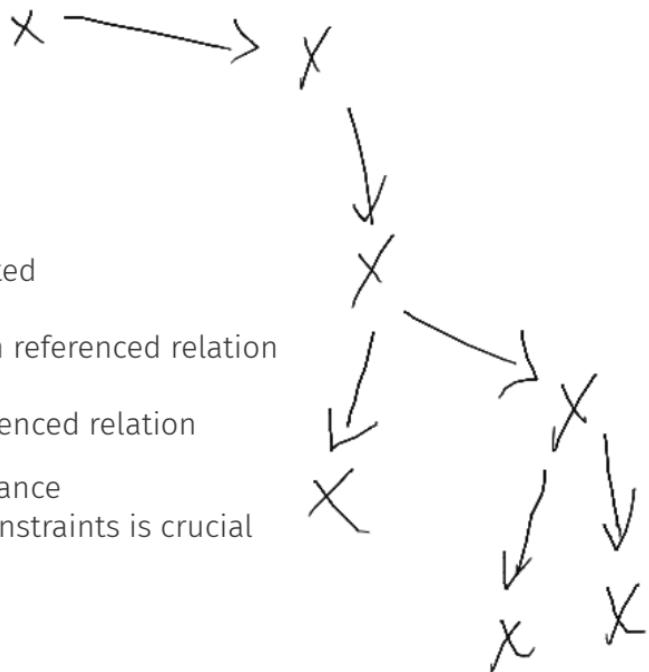
# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- Problems  
- Actions  
- Example  
- Consideration  
- Exercise  
General  
Summary

## Foreign Key Constraint

### Practical Considerations

- Specified constraints might not behave as expected
  - SET NULL issue with prime attributes
  - SET DEFAULT issue with default value not in referenced relation
  - CASCADE may create a chain of propagation
    - If the referencing relation is also a referenced relation
- CASCADE may significantly affect overall performance
- Careful design and specification of foreign key constraints is crucial



# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
**Foreign Key**  
- Recap  
- Constraint  
- Condition  
- Problems  
- Actions  
- Example  
- Consideration  
- **Exercise**  
General  
Summary

## Foreign Key Constraint

### Schema

- $R(A, B)$
- $S(C, D)$

A	B
1	2
2	1
1	1
1	3

C	D
3	NULL
NULL	1
1	3
NULL	NULL

Row
Row 1
Row 2
Row 3
Row 4

### Exercise #3

Consider the foreign key constraints  $(S.C, S.D) \rightarrow (R.A, R.B)$ . Which rows violate the FOREIGN KEY constraint?

Choice	Comment	
A	Row 1	?
B	Row 2	?
C	Row 3	?
D	Row 4	?

# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
Foreign Key  
**General**  
- *CHECK*  
- *Complex*  
- *Assertion*  
Summary

## General Constraint

### CHECK Constraint

- Most basic general constraint (*i.e., not a structural integrity constraint*)
- Allows us to specify that column values must satisfy an arbitrary Boolean expression
- **Scope:** one table, one row

### Single Column

```
CREATE TABLE Projects (
    eid      INT,
    pname    TEXT,
    hours   INT CHECK (hours > 0),
    PRIMARY KEY (eid, pname)
);
```

### Multiple Columns

```
CREATE TABLE Teams (
    name      TEXT PRIMARY KEY,
    start_year INT,
    end_year   INT,
    CHECK (start_year <= end_year)
);
```

# Integrity Constraints

Preliminary

Not-NULL

Unique

Primary Key

Foreign Key

**General**

- *CHECK*

- *Complex*

- *Assertion*

Summary

## General Constraint

Complex Boolean Constraint

### Example

Create the Table "Teams" where the minimum hours for 'CoreOS' is at least 30 hours but there are no such restrictions for other projects.

```
CREATE TABLE Projects (
    eid      INT,
    pname   TEXT,
    hours   INT CHECK (hours > 0),
    PRIMARY KEY (eid, pname),
    CHECK (
        (pname = 'CoreOS' AND hours >= 30)
        OR
        (pname <> 'CoreOS' AND hours > 0)
    )
);
```

#Is this good enough? Can you do better?

# Integrity Constraints

Preliminary

Not-NULL

Unique

Primary Key

Foreign Key

**General**

- *CHECK*

- *Complex*

- ***Assertion***

Summary

## General Constraint

### Assertion

- CREATE ASSERTION statement (*since SQL-92*)
- Used for (*almost*) arbitrary constraints (*multiple table, multiple rows*)
  - **Example:** "Each project must have at least one team member being 30 or older"

### In Practice

- Assertions may have various potential side-effects and/or limitations
  - Cannot modify the data
  - No proper error handling
  - Not linked to a specific table (*e.g., dropping a table does not affect assertion*)
- Most RDBMS do **NOT** support assertions
  - **Alternative: Triggers**



# Integrity Constraints

Preliminary  
Not-NULL  
Unique  
Primary Key  
Foreign Key  
General  
**Summary**

## Summary

### Types

Type	Column	Table	Condition
Not-NULL	NOT NULL	-	IS NOT NULL
Unique	UNIQUE	UNIQUE( $A_1, A_2, \dots$ )	$x.A_i \neq y.A_i$
Primary Key	PRIMARY KEY	PRIMARY KEY( $A_1, A_2, \dots$ )	UNIQUE & NOT NULL
Foreign Key	REFERENCES R <sub>1</sub> (B)	FOREIGN KEY ( $A_1, A_2, \dots$ ) REFERENCES R <sub>1</sub> (B <sub>1</sub> , B <sub>2</sub> , ...)	The tuple exists in R <sub>1</sub> or the tuple contains NULL value
General	CHECK (c)	CHECK (c)	Condition c does not evaluate to False

# Modifying Database

---

# Modifying Database

ALTER

- *Syntax*

- Single Column

- Add/Drop

DROP

## ALTER TABLE

### Syntax

```
ALTER TABLE <table_name>
[ALTER / ADD / DROP] [COLUMN / CONSTRAINT] <name>
<changes>;
```

### Common Changes:

- Adding/dropping columns
- Adding/dropping constraints
- Changing the specification of a single column:
  - Data type
  - Default values

# Modifying Database

ALTER

- *Syntax*

- *Single Column*

- *Add/Drop*

DROP

## ALTER TABLE

Altering a Single Column

Changing Data Type

```
ALTER TABLE Projects ALTER COLUMN name TYPE VARCHAR(200);
```

Changing Default Values

```
ALTER TABLE Projects ALTER COLUMN start_year SET DEFAULT 2021;
```

Removing Default Values

```
ALTER TABLE Projects ALTER COLUMN name DROP DEFAULT;
```

# Modifying Database

ALTER

- Syntax
- Single Column

- Add/Drop

DROP

## ALTER TABLE

Add/Drop

### Adding/Dropping Columns

```
ALTER TABLE Projects ADD COLUMN budget NUMERIC DEFAULT 0.0;  
ALTER TABLE Projects DROP COLUMN budget;
```

### Adding/Dropping Constraints

```
ALTER TABLE Teams ADD CONSTRAINT  
eid_fkey FOREIGN KEY (eid) REFERENCES Employees (id);  
ALTER TABLE Teams DROP CONSTRAINT eid_fkey;
```

- Name of constraints need to be known (*may be retrieved from metadata*)

<sup>#</sup><https://www.postgresql.org/docs/14/sql-altertable.html>

# Modifying Database

ALTER  
DROP  
-Syntax

## DROP TABLE

### Syntax

```
DROP TABLE
[IF EXISTS]    -- no error if table does not exist
<table_name>[, <table_name> [, <table_name> [...]]] -- multiple table
[CASCADE];     -- also delete referencing table
```

# Deferrable Constraints

---

# Deferrable Constraints

## Motivation

- *Behavior*
- *Transition*
- *Circular*
- Deferment Considerations

## Motivation

### Definition

An SQL **statement** is any SQL code that ends with a **semi-colon** (*i.e.*, ;)

### Default Behavior

- Constraints are checked **immediately** at the end of SQL statement
  - This is true even for transaction containing multiple SQL statement A violation will cause the statement/transaction to be **rolled back**

### Relaxed Constraint

- Check only at the end of a transaction
- Available for: **UNIQUE**, **PRIMARY KEY**, **FOREIGN KEY**

# Deferrable Constraints

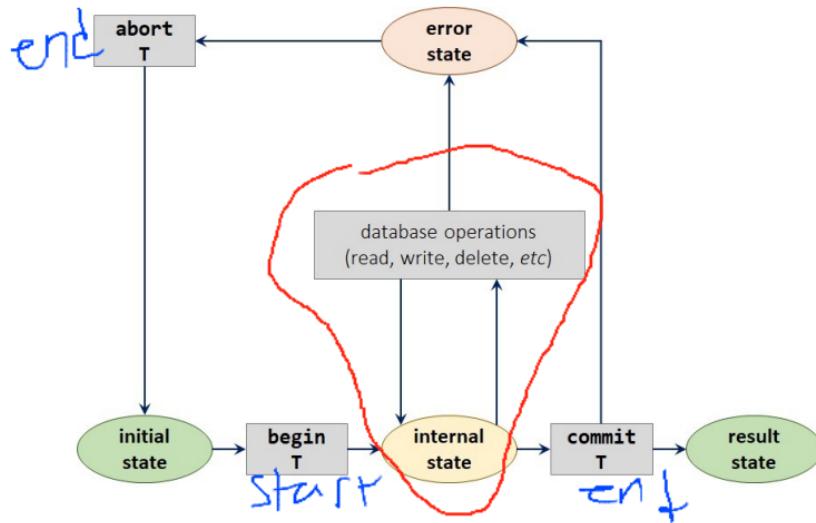
## Motivation

- Behavior
- Transition
- Circular

Deferment Considerations

## Motivation

### Transition Graph



Deferrable constraints may (*temporarily*) be violated within the scope of a transaction in the "internal state"

# Deferrable Constraints

## Motivation

- Behavior
- Transition
- Circular

## Deferment Considerations

## Motivation

A Curious Case of Circular Reference

### Employee

An employee must work in a department

```
CREATE TABLE Employee (
    eid INT PRIMARY KEY,
    name TEXT,
    did INT
);
```

### Department

A department must be managed by an employee

```
CREATE TABLE Department (
    did INT PRIMARY KEY,
    name TEXT,
    eid INT
);
```

```
ALTER TABLE Employee ADD CONSTRAINT
    did_fk FOREIGN KEY (did)
    REFERENCES Department (did);
```

```
ALTER TABLE Department ADD CONSTRAINT
    eid_fk FOREIGN KEY (eid)
    REFERENCES Employee (eid);
```

# Deferrable Constraints

Motivation

## Deferment

- NOT DEFERRABLE
  - INITIALLY DEFERRED
  - INITIALLY IMMEDIATE
- Considerations

## Deferment

NOT DEFERRABLE

```
BEGIN; -- start of transaction  
  
INSERT INTO Employee VALUES (101, 'Sarah', 1001);  
INSERT INTO Department VALUES (1001, 'dev', 101);  
COMMIT; -- successful end of transaction
```

- **Line 3:** Foreign Key Constraint Violation
- **Line 4:** Foreign Key Constraint Established

**ABORT**

*not executed*

```
ALTER TABLE Employee ADD CONSTRAINT  
    did_fk FOREIGN KEY (did)  
    REFERENCES Department (did)  
    NOT DEFERRABLE;
```

```
ALTER TABLE Department ADD CONSTRAINT  
    eid_fk FOREIGN KEY (eid)  
    REFERENCES Employee (eid)  
    NOT DEFERRABLE;
```

# Deferrable Constraints

Motivation

## Deferment

- NOT DEFERRABLE
  - INITIALLY DEFERRED
  - INITIALLY IMMEDIATE
- Considerations

## Deferment

INITIALLY DEFERRED

```
BEGIN; -- start of transaction  
  
INSERT INTO Employee VALUES (101, 'Sarah', 1001);  
INSERT INTO Department VALUES (1001, 'dev', 101);  
COMMIT; -- successful end of transaction
```

- **Line 3:** Foreign Key Constraint Violation
- **Line 4:** Foreign Key Constraint Established

*constraint check deferred*  
**SUCCESS**

```
ALTER TABLE Employee ADD CONSTRAINT  
    did_fk FOREIGN KEY (did)  
    REFERENCES Department (did)  
    DEFERRABLE INITIALLY DEFERRED;
```

```
ALTER TABLE Department ADD CONSTRAINT  
    eid_fk FOREIGN KEY (eid)  
    REFERENCES Employee (eid)  
    DEFERRABLE INITIALLY DEFERRED;
```

# Deferrable Constraints

Motivation

Deferment

- NOT DEFERRABLE
- INITIALLY DEFERRED
- INITIALLY IMMEDIATE

Considerations

Deferment

INITIALLY IMMEDIATE

```
BEGIN; -- start of transaction
SET CONSTRAINT Eid_fk DEFERRED;
INSERT INTO Employee VALUES (101, 'Sarah', 1001);
INSERT INTO Department VALUES (1001, 'dev', 101);
COMMIT; -- successful end of transaction
```

missing 5

- Line 3: Foreign Key Constraint Violation
- Line 4: Foreign Key Constraint Established

constraint check deferred  
**SUCCESS**

```
ALTER TABLE Employee ADD CONSTRAINT
    did_fk FOREIGN KEY (did)
    REFERENCES Department (did)
    DEFERRABLE INITIALLY IMMEDIATE;
```

```
ALTER TABLE Department ADD CONSTRAINT
    eid_fk FOREIGN KEY (eid)
    REFERENCES Employee (eid)
    DEFERRABLE INITIALLY DEFERRED;
```

# Deferrable Constraints

Motivation  
Deferment  
**Considerations**

## Considerations

### Benefits

- No need to care about the order of SQL statements within a transaction  
*(i.e., intermediate state may temporarily violate constraints)*
- Allows for cyclic foreign key constraints
- Performance boost when constraint checks are bottleneck  
*(e.g., batch insert of large number of tuples)*

### Potential Downsides

- Troubleshooting can be more difficult
- Data definition no longer unambiguous
- May incur performance penalty when performing queries
  - Certain checks may need to be done at run-time especially during the time when the constraint check is deferred

# Summary

---

# Summary

## SQL

- *What*

- *Table*

- *Rows*

Constraints

## SQL

### What Is It?

- The standard language for RDBMS
  - Different language groups: DDL, DML, DQL, DCL, TCL
- **DDL:** CREATE TABLE, ALTER TABLE, DROP TABLE
- **DML:** INSERT, UPDATE, DELETE

### Key Challenge

- Specification of integrity constraints
  - NOT NULL, UNIQUE, PRIMARY KEY, FOREIGN KEY, CHECK
- Specification actions in case of foreign key constraint violations (*ON UPDATE/ON DELETE*)
- Relaxed checks of violations with deferrable constraints

# Summary

## SQL

- *What*
  - **Table**
  - **Rows**
- Constraints

## SQL

### Table Syntax

#### CREATE

```
CREATE TABLE <table_name> (
    <attr> <type> [<column_constraint>],
    <attr> <type> [<column_constraint>],
    :
    [<table_constraint>],
    [<table_constraint>],
    :
);
;
```

#### DROP

```
DROP TABLE [IF EXISTS] <table_name>;
```

# Summary

## SQL

- *What*

- *Table*

- *Rows*

Constraints

## SQL

Rows Syntax

### INSERT

```
INSERT INTO <table_name> [<attr>, ...]
    VALUES (<values>, ...) [, (<values>, ...)];
```

### UPDATE

```
UPDATE <table_name>
    SET <attr> = <value> [, <attr> = <value>]
    [WHERE <condition>];
```

### DELETE

```
DELETE FROM <table_name> [WHERE <condition>];
```

# Summary

## SQL Constraints

### Integrity Constraints

#### Types

Type	Column	Table	Condition
Not-NULL	NOT NULL	-	IS NOT NULL
Unique	UNIQUE	UNIQUE(A <sub>1</sub> ,A <sub>2</sub> ,...)	x.A <sub>i</sub> <> y.A <sub>i</sub>
Primary Key	PRIMARY KEY	PRIMARY KEY(A <sub>1</sub> ,A <sub>2</sub> ,...)	UNIQUE & NOT NULL
Foreign Key	REFERENCES R <sub>1</sub> (B)	FOREIGN KEY (A <sub>1</sub> ,A <sub>2</sub> ,...) REFERENCES R <sub>1</sub> (B <sub>1</sub> ,B <sub>2</sub> ,...)	The tuple exists in R <sub>1</sub> or the tuple contains NULL value
General	CHECK (c)	CHECK (c)	Condition c does not evaluate to False

```
postgres=# exit
```

```
Press any key to continue . . .
```

# Solutions

## Exercise #1

Exercise #2

Exercise #3

## Not-NULL Constraint

Schema

Table "Employees"

```
CREATE TABLE Employees (
    id      INT NOT NULL,
    name    TEXT NOT NULL,
    age     INT,
    role    TEXT
);
```

	<b>id</b>	<b>name</b>	<b>age</b>	<b>role</b>	<b>Row</b>
	101	'Sarah'	25	'dev'	Row 1
	NULL	'Judy'	32	'sales'	Row 2
	102	"	27	'hr'	Row 3
	103	'Max'	NULL	NULL	Row 4

## Exercise #1

Consider the schema on the left and the table at the top. Which rows violate the **NOT NULL** constraint?

	<b>Choice</b>	<b>Comment</b>	
A	Row 1	NO: there is no <b>NULL</b> values	<span style="color:red;">✗</span>
B	Row 2	YES: <b>id</b> is <b>NULL</b>	<span style="color:green;">✓</span>
C	Row 3	NO: empty string is not <b>NULL</b>	<span style="color:red;">✗</span>
D	Row 4	NO: these attributes can be <b>NULL</b>	<span style="color:red;">✗</span>

# Solutions

Exercise #1

Exercise #2

Exercise #3

## Constraints

Schema

Table "Teams"

```
CREATE TABLE Teams (
    eid    INT,
    pname  TEXT,
    hours  INT,
    UNIQUE (eid, pname)
);
```

	eid	pname	hours	Row
Row 1	101	NULL	25	
Row 2	101	NULL	25	
Row 3	102	'Judy'	32	
Row 4	102	'Judy'	32	

Exercise #2

Consider the schema on the left and the table at the top. Which rows are not **UNIQUE**?

	Choice	Comment	
A	Row 1	NO: NULL <> NULL evaluates to NULL	<span style="color: red;">✗</span>
B	Row 2	NO: NULL <> NULL evaluates to NULL	<span style="color: red;">✗</span>
C	Row 3	YES: equal to Row 4	<span style="color: green;">✓</span>
D	Row 4	YES: equal to Row 3	<span style="color: green;">✓</span>

# Solutions

Exercise #1  
Exercise #2  
**Exercise #3**

## Foreign Key Constraint

### Schema

- $R(A, B)$
- $S(C, D)$

A	B	C	D	Row
1	2	3	NULL	Row 1
2	1	NULL	1	Row 2
1	1	1	3	Row 3
1	3	NULL	NULL	Row 4

### Exercise #3

Consider the foreign key constraints  $(S.C, S.D) \rightarrow\!\!\! \rightarrow (R.A, R.B)$ . Which rows violate the FOREIGN KEY constraint?

	Choice	Comment	
A	Row 1	NO: it has <b>NULL</b> value for column D	✗
B	Row 2	NO: it has <b>NULL</b> value for column D	✗
C	Row 3	NO: the tuple (1, 3) appears in R	✗
D	Row 4	NO: it is both <b>NULL</b>	✗