

CS2102: Database Systems

SQL (Part 2)

Clarifications

Clarifications

Candidate Keys
Weak Equivalence

Candidate Keys

NULL or Not NULL?

Consideration

- Consider $R(A,B,C,D)$
 - Let the primary key be $\{A\}$
 - Let $\{B,C\}$ be a candidate key

Line 1
Line 2

A	B	C	D
1	NULL	1	0
2	NULL	1	1

$\{B, C\}$

Problems?

- Does $\{A\}$ uniquely identify $\{D\}$?
- Does $\{B,C\}$ uniquely identify $\{D\}$?

Notes

This problem arise because postgresql does not consider $(NULL, 1)$ to be different from $(NULL, 1)$ using $=$.

In postgresql, can be inserted after Line 2. Therefore, candidate key $\{B, C\}$ does not uniquely identify D.

Clarifications

Candidate Keys
Weak Equivalence

Weak Equivalence

Mathematical Example

- Consider the sequence $1 + q + q^2 + q^3 + \dots + q^n$ for $q \leq 1$
- Is it equivalent to $(1 - q^n) / (1 - q)$?
 - If $q \neq 1$ then yes
 - If $q = 1$ then $(1 - q^n) / (1 - q)$ is *undefined* (i.e., error)

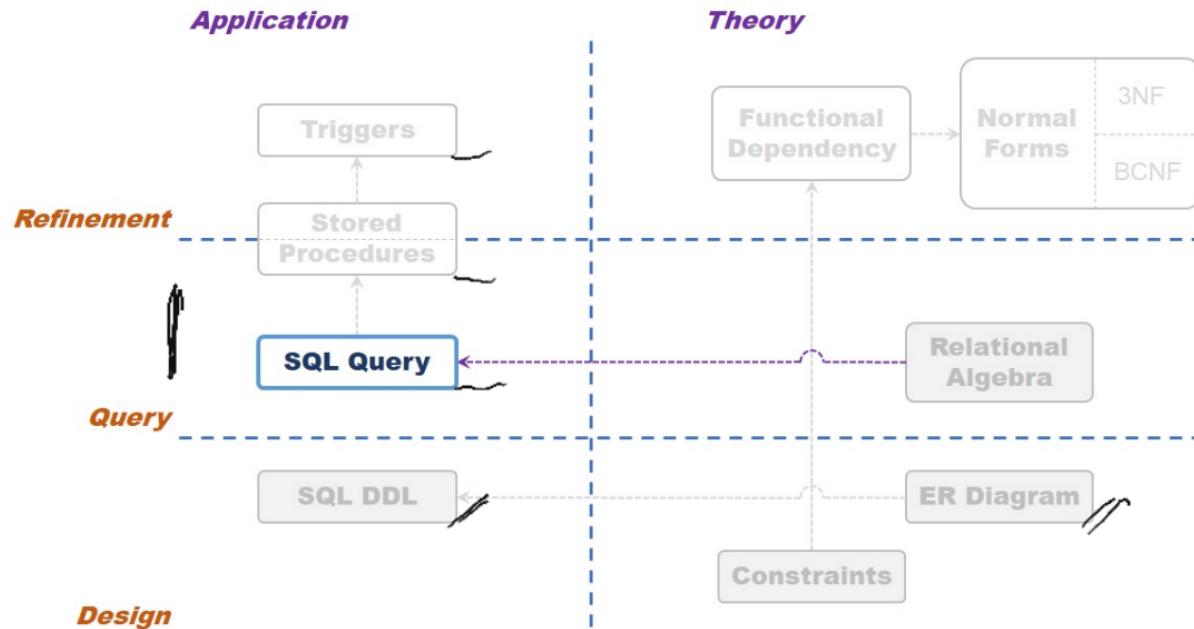
Relational Algebra

- Note that using our construct, relational algebra expression is either *always error* (i.e., *invalid expression*) or *never error*
 - Still useful for generic expressions where some elements are left unspecified (e.g., $\sigma_{[C]}(R - S)$ is only weakly equivalent to $\sigma_{[C]}(R) - S$)



- SQL query has a construct that is only checked at run-time (e.g., *scalar subquery*) which may or may not produce an error depending on the input relation

Roadmap



Roadmap

Overview

Overview

SQL vs. Relational Algebra

Basic SQL Queries

- Simple Queries
- Set Operations
- Join Queries

Composing

- Subqueries

Ordering

- Ordering
- Ranking

SQL vs. Relational Algebra

SQL vs. Relational Algebra

Comparison

Basic Form
Evaluation
Strategy
Database

Comparison

Comparison Table

	Relational Algebra	SQL
Paradigm	Imperative <i>(query language using operators to perform queries)</i>	Declarative <i>(query language built on top of relational algebra)</i>
Semantic	Set <i>(i.e., no duplicate rows/tuples)</i>	Multiset/Bag <i>(i.e., allows duplicate rows/tuples)</i>
Query	Relational algebra expression	<u>SELECT statement</u>

SELECT Statement

```
SELECT [ DISTINCT ] <target-list>
FROM   <relation-list>
[ WHERE <conditions> ];
```



SQL vs. Relational Algebra

Comparison

Basic Form

Evaluation

Strategy

Database

Basic Form

SELECT-FROM-WHERE



```
SELECT [ DISTINCT ] <target-list>
      FROM <relation-list>
      WHERE <conditions>
```

=
IS DISTINCT FROM

-- SELECT clause ✓
-- FROM clause ✓
-- WHERE clause ✓

Notes

- [DISTINCT]
By default, duplicates are **NOT** eliminated. Explicit specification of DISTINCT is required to eliminate duplicates.
- <target-list>
List of attributes/columns of the output result.
- <relation-list>
List of relations that are relevant to answer the given query.
- <conditions>
Conditions that specify a valid result row/tuple
(based on principle of acceptance).
- **Order-Independent:** The order of rows in the result does not matter (*unsorted/unranked*)

SQL vs. Relational Algebra

Comparison

Basic Form

Evaluation

Strategy

Database

Basic Form

SELECT-FROM-WHERE

```
SELECT  [ DISTINCT ] <target-list>
        FROM    <relation-list>
        WHERE   <conditions>
```

-- SELECT clause
-- FROM clause
-- WHERE clause

Notes

- [DISTINCT]
- <target-list>
- <relation-list>
- <conditions>
- **Order-Independent:** The order of rows in the result does not matter(*unsorted/unranked*)

By default, duplicates are **NOT** eliminated. Explicit specification of DISTINCT is required to eliminate duplicates.

List of attributes/columns of the output result.

List of relations that are relevant to answer the given query.

Conditions that specify a valid result row/tuple
(based on *principle of acceptance*).

[#]We will not penalize if you are using DISTINCT keyword everywhere as long as the final results satisfy the requirement.

SQL vs. Relational Algebra

Comparison
Basic Form
Evaluation
- Preprocessing
- Steps
Strategy
Database

Conceptual Evaluation

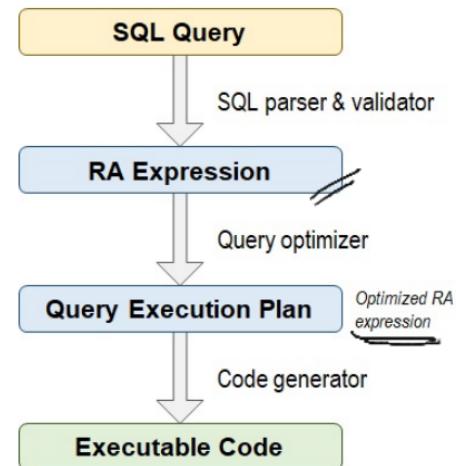
Preprocessing: SQL Query to Executable Code

SELECT Statement

```
SELECT DISTINCT a1, a2, ..., am
FROM      r1, r2, ..., rn
WHERE      c
```

Relational Algebra

$$\pi[a_1, a_2, \dots, a_m](\sigma[c](r_1 \times r_2 \times \dots \times r_n))$$



SQL vs. Relational Algebra

Comparison
Basic Form
Evaluation
- Preprocessing
- Steps
Strategy
Database

Conceptual Evaluation

Evaluation Steps

SELECT Statement

```
SELECT DISTINCT a1, a2, ..., am
FROM   r1, r2, ..., rn
WHERE  c
```

1. Compute cross product of all tables (r_1, r_2, \dots, r_n)

Relational Algebra

$$\pi[a_1, a_2, \dots, a_m](\sigma[c](r_1 \times r_2 \times \dots \times r_n))$$

SQL vs. Relational Algebra

Comparison

Basic Form

Evaluation

- Preprocessing

- Steps

Strategy

Database

Conceptual Evaluation

Evaluation Steps

SELECT Statement

```
SELECT DISTINCT a1, a2, ..., am  
FROM   r1, r2, ..., rn  
WHERE  c
```

Relational Algebra

$$\pi[a1, a2, \dots, am](\sigma_{[c]}(r1 \times r2 \times \dots \times rn))$$

1. Compute cross product of all tables ($r1, r2, \dots, rn$)
2. Keep only rows where condition is true (c)



SQL vs. Relational Algebra

Comparison

Basic Form

Evaluation

- Preprocessing

- Steps

Strategy

Database

Conceptual Evaluation

Evaluation Steps

SELECT Statement

```
SELECT DISTINCT a1, a2, ..., am
FROM   r1, r2, ..., rn
WHERE  c
```



1. Compute cross product of all tables ($r1, r2, \dots, rn$)
2. Keep only rows where condition is true (c)
3. Keep only attributes specified ($a1, a2, \dots, am$)

Relational Algebra

$$\pi [a1, a2, \dots, am] (\sigma[c](r1 \times r2 \times \dots \times rn))$$

SQL vs. Relational Algebra

Comparison

Basic Form

Evaluation

- Preprocessing

- Steps

Strategy

Database

Conceptual Evaluation

Evaluation Steps

SELECT Statement

```
SELECT DISTINCT a1, a2, ..., am  
FROM   r1, r2, ..., rn  
WHERE  c
```

Relational Algebra

$$\pi[a_1, a_2, \dots, a_m](\sigma[c](r_1 \times r_2 \times \dots \times r_n))$$

1. Compute cross product of all tables (r_1, r_2, \dots, r_n)
2. Keep only rows where condition is true (c)
3. Keep only attributes specified (a_1, a_2, \dots, a_m)
4. Remove duplicate rows $(\text{only if distinct is specified})$

SQL vs. Relational Algebra

Comparison

Basic Form

Evaluation

- Preprocessing

- Steps

Strategy

Database

Conceptual Evaluation

Evaluation Steps

SELECT Statement

```
SELECT DISTINCT a1, a2, ..., am  
FROM   r1, r2, ..., rn  
WHERE  c
```

Relational Algebra

$$\pi[a_1, a_2, \dots, a_m](\sigma[c](r_1 \times r_2 \times \dots \times r_n))$$

1. Compute cross product of all tables (r_1, r_2, \dots, r_n)
2. Keep only rows where condition is true (c)
3. Keep only attributes specified (a_1, a_2, \dots, a_m)
4. Remove duplicate rows $(\text{only if distinct is specified})$

Notes

This relational algebra mapping is a conceptual mapping without consideration for performance

SQL vs. Relational Algebra

Comparison
Basic Form
Evaluation
Strategy
Database

Basic Strategy

Problem Solving

Subproblems

- Can you find a simpler problem to solve or solve a partial problem?
 - If so, solve the simpler/partial problem as subqueries
 - Then solve the main problem

Problem Deconstruction

- Answer the following questions
 - What columns/attributes to find?
 - Which relations containing the attributes?
 - What are the conditions to be satisfied?

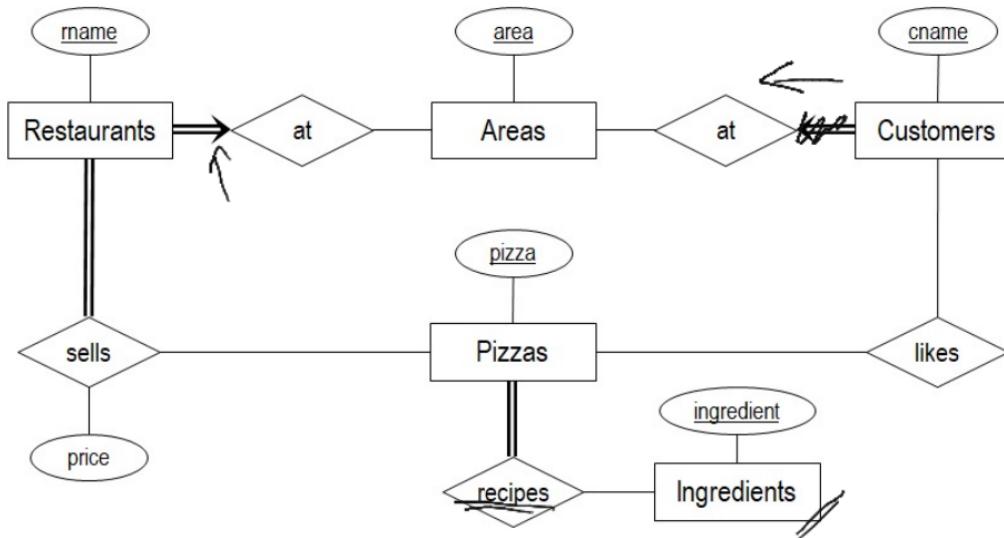


SQL vs. Relational Algebra

Comparison
Basic Form
Evaluation
Strategy
Database
- ER Diagram
- Schema
- Sample

Example Database

ER Diagram



[#]This is only an *approximate* to help visualize the schema.

SQL vs. Relational Algebra

Comparison
Basic Form
Evaluation
Strategy
Database

- ER Diagram
- Schema
- Sample

Example Database

Schema



Relations

- Restaurants(rname, area)
- Customers(cname, area)
- Pizzas(pizza)
- Ingredients(ingredient)
- Sells(rname, pizza, price)
- Likes(cname, pizza)
- Recipes(pizza, ingredient)

Foreign Key

- (Sells.rname) ↗ (Restaurants.rname)
- (Sells.pizza) ↗ (Pizzas.pizza)
- (Likes.cname) ↗ (Customers.cname)
- (Likes.pizza) ↗ (Pizzas.pizza)
- (Recipes.pizza) ↗ (Pizzas.pizza)
- (Recipes.ingredient) ↗ (Ingredients.ingredient)

SQL (L05.sql)

```
CREATE TABLE Pizzas (
    pizza      VARCHAR NOT NULL PRIMARY KEY
);
```

Universal Schema
Assumption



SQL vs. Relational Algebra

Comparison
Basic Form
Evaluation
Strategy
Database
- ER Diagram
- Schema
- *Sample*

Example Database

Sample Database

Table "Sells"

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mammas Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Table "Restaurants"

rname	area
Corleone Corner	North
Gambino Oven	Central
Lorenzo Tavern	Central
Mammas Place	South
Pizza King	East

Table "Pizzas"

pizza
Diavola
Funghi
Hawaiian
Margherita
Marinara
Siciliana

Table "Recipes"

pizza	ingredient
Diavola	Cheese
Diavola	Chilli
Diavola	Salami
Funghi	Ham
Funghi	Mushroom
Hawaiian	Ham
Hawaiian	Pineapple
Margherita	Cheese
Margherita	Tomato
Marinara	Seafood
Marinara	Chilli
Siciliana	Anchovies
Siciliana	Caspers
Siciliana	Cheese

Table "Customers"

cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Alice	Central
Bob	Central
Willie	North
Adi	NULL
Yoga	NULL

Table "Likes"

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Siciliana
Ralph	Diavola
Alice	Siciliana
Alice	Funghi
Bob	Siciliana

Basic SQL Queries

Basic SQL Queries

Simple Queries

- *Simple Condition*
- *Complex Condition*
- *All Attributes*

SELECT Clause

WHERE Clause

Set Operations

Join Operations

Simple Queries

Simple Condition

Question

name

pizza

Find the restaurants and the pizzas that they sell where the price is under \$20.

Problem Deconstruction

- What columns/attributes to find?
- Which relations have the attributes?
- What are the conditions to be satisfied?

Basic SQL Queries

Simple Queries

- Simple Condition
- Complex Condition
- All Attributes

SELECT Clause

WHERE Clause

Set Operations

Join Operations

Simple Queries

Simple Condition

Question

Find the restaurants and the pizzas that they sell where the price is under \$20.

Problem Deconstruction

- What columns/attributes to find?
- Which relations have the attributes?
- What are the conditions to be satisfied?

rname, pizza
Sells
price under \$20

— SELECT
— FROM
— WHERE

price < 20

Basic SQL Queries

Simple Queries

- Simple Condition

- Complex Condition

- All Attributes

SELECT Clause

WHERE Clause

Set Operations

Join Operations

Simple Queries

Simple Condition

Question

Find the restaurants and the pizzas that they sell where the price is under \$20.

Problem Deconstruction

- What columns/attributes to find?
- Which relations have the attributes?
- What are the conditions to be satisfied?

rname, pizza

Sells

price under \$20

SQL

```
SELECT rname, pizza  
FROM   Sells  
WHERE  price < 20;
```

Loop is applied.

rname	pizza
Corleone Corner	Margherita
Gambino Oven	Siciliana
Pizza King	Diavola

Basic SQL Queries

Simple Queries

- Simple Condition
- Complex Condition
- All Attributes

SELECT Clause

WHERE Clause

Set Operations

Join Operations

Simple Queries

Complex Condition

Exercise #1

Find the restaurants and the pizzas that they sell where either (i) the price is over \$20 or the pizza is Siciliana and (ii) the pizza is not Diavola.

```
1 -- What columns/attributes to find?  
2 -- Which relations have the attributes?  
3 -- What are the conditions to be satisfied?  
4  
5 (price > 20 OR pizza = 'Siciliana') AND (pizza <> 'Diavola')
```

rname	pizza
Corleone Corner	Hawaiian
Gambino Oven	Siciliana
Lorenzo Tavern	Funghi
Mammas Place	Marinara
Pizza King	Hawaiian

Basic SQL Queries

Simple Queries

- Simple Condition
- Complex Condition
- All Attributes

SELECT Clause
WHERE Clause
Set Operations
Join Operations

Simple Queries

All Attributes

Question

Find the restaurants, the pizzas that they sell, and the price of the pizza where the price is under \$20.

Wild Card

- We can use wild card * to include all attributes

SQL

```
SELECT *
FROM Sells
WHERE price < 20;
```

rname	pizza	price
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Pizza King	Diavola	17

Basic SQL Queries

Simple Queries

SELECT Clause

- Expression

- Duplicates

WHERE Clause

Set Operations

Join Operations

SELECT Clause

Expressions

Common Use Cases

- Combine and process attribute values
- Rename columns

Question

Find the restaurants, the pizzas that they sell, and the price of the pizza in SGD (*current exchange rate is 1 USD = 1.36 SGD*).

SQL

```
SELECT rname, pizza,  
       '$$' || (price * 1.36) AS sgd  
FROM Sells;
```

rname	pizza	sgd
Corleone Corner	Diavola	\$\\$2.64
Corleone Corner	Hawaiian	\$\\$4.00
: Corner	:	:
Pizza King	Hawaiian	\$\\$28.56

Basic SQL Queries

Simple Queries

SELECT Clause

- Expression

- Duplicates

WHERE Clause

Set Operations

Join Operations



SELECT Clause

Expressions

Common Use Cases

- Combine and process attribute values
- Rename columns

Question

Find the restaurants, the pizzas that they sell, and the price of the pizza in SGD (*current exchange rate is 1 USD = 1.36 SGD*).

SQL

```
SELECT rname, pizza,
       '$$' || (price * 1.36) AS sgd
  FROM Sells;
```

Operations

- Mathematical

+, *, -, %, /, ^, |/, etc

- String

|| (concatenate), LOWER(s),

UPPER(s), etc

- Date Time

+, NOW(), etc

Renaming

- column AS alias

Basic SQL Queries

Simple Queries

SELECT Clause

- Expression

- Duplicates

WHERE Clause

Set Operations

Join Operations

SELECT Clause

Duplicates

Multiset Semantics

- By default, SQL does **NOT** eliminate duplicates
- Keyword **DISTINCT** to enforce duplicate elimination

Question

Find all pizza names that is sold by at least one restaurant.

SQL

```
SELECT pizza  
FROM Sells;
```

pizza
Diavola
Hawaiian
:
Diavola
Hawaiian

8 rows

Notes

The order of the output when using **DISTINCT** may be different. This depends on how duplicates are removed.

```
SELECT DISTINCT pizza  
FROM Sells;
```

pizza
Hawaiian
Margherita
:
Siciliana

pizza
Hawaiian
Margherita
:
Siciliana

6 rows

Basic SQL Queries

Simple Queries

SELECT Clause

- Expression

- Duplicates

WHERE Clause

Set Operations

Join Operations

SELECT Clause

Duplicates

NULL Values

- DISTINCT keyword checks for *distinct* rows using IS DISTINCT FROM

```
SELECT area  
FROM Customers;
```

x	y	x IS DISTINCT FROM y	x IS NOT DISTINCT FROM y
not-NULL	not-NULL	$x \neq y$	$x = y$
not-NULL	NULL	True	False
NULL	not-NULL	True	False
NULL	NULL	False	True

area
West
South
East
Central
Central
Central
Central
North
NULL
NULL

```
SELECT DISTINCT area  
FROM Customers;
```

area
NULL
East
South
West
North
Central

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

- **NULL**

- *Pattern Matching*

Set Operations

Join Operations

WHERE Clause

Conditions for NULL Values

- Finding tuples with NULL or not-NUL as attribute values
 - $x \text{ IS NULL}$ vs $x = \text{NULL}$
 - $x \text{ IS NOT NULL}$ vs $x \neq \text{NOT NULL}$

Question

Find all customer names with unknown area.

SQL

IS NULL

```
SELECT cname  
FROM Customers  
WHERE area IS NULL;
```

cname
Adi
Yoga

2 rows

Notes

There is no (*syntax*) error when using = or \neq operator. But the output may not be what you expect!

$\text{NULL} = \text{NULL}$

→ evaluates to NVLL

NVLL IS NULL

→ evaluates to True.

= NULL

```
SELECT cname  
FROM Customers  
WHERE area = NULL;
```

cname

0 rows

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

- *NULL*

- *Pattern Matching*

Set Operations

Join Operations

WHERE Clause

Pattern Matching

Basic Pattern Matching ([NOT] LIKE)

- _ matches any single character
- % matches any sequence of 0 or more characters

Question

Find all pizza that start with "Ma" and end with "a".

SQL

```
SELECT pizza  
FROM Pizzas  
WHERE pizza LIKE 'Ma%a';
```

Notes

Some pattern may actually be simplified. For example, %_ can be simplified to %.

numeric, alphabet,
→ symbol

Advanced Patter Matching

- *Regular expression*
(out of scope)

Ma^q ✓ Ma X

pizza
Marinara
Margherita

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

- *NULL*

- *Pattern Matching*

Set Operations

Join Operations

WHERE Clause

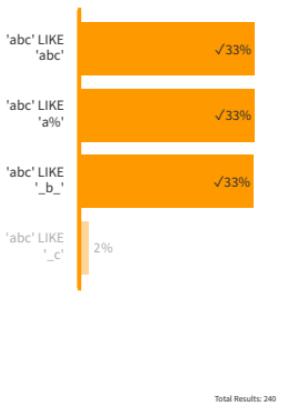
Pattern Matching

Basic Pattern Matching ([NOT] LIKE)

- _ matches *any* single character
- % matches any sequence of *0 or more* characters

Exercise 2

Which of the following conditions below returns TRUE?



Choice

Comment

A	'abc' LIKE 'abc'	
B	'abc' LIKE 'a%'	
C	'abc' LIKE '_b_'	
D	'abc' LIKE '_c'	

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

- Operations

- Duplication

- Example

- Exercises

Join Operations

Set Operations

Operations

Let Q_1 and Q_2 be two SQL queries that yield **union-compatible** tables

- $Q_1 \text{ UNION } Q_2 = Q_1 \cup Q_2$
- $Q_1 \text{ INTERSECT } Q_2 = Q_1 \cap Q_2$
- $Q_1 \text{ EXCEPT } Q_2 = Q_1 - Q_2$

Important



Eliminate duplicate rows from result

- Column names follow the left relation (*i.e.*, Q_1)

- The ordering of the output result may be different

Recap

Two relations R and S are **union-compatible** if it satisfies both of the following:

- R and S have the same number of attributes
- The corresponding attributes have the same or compatible domains

Table "R"		Table "S"	
v1	1	v2	2
	2		3
	2		

UNION	INTERSECT	EXCEPT
v1 1 2 3	v1 2	v1 1
X		

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

- Operations

- Duplication

- Example

- Exercises

Join Operations

Set Operations

Duplication

The following operations do **NOT** remove duplicates

- $Q_1 \text{ UNION } Q_2$
- $Q_1 \text{ INTERSECT } Q_2$
- $Q_1 \text{ EXCEPT } Q_2$

Important

- May produce unintuitive result (*but explainable*)
 - Basically, each element is treated as distinct element
 - Intersection keeps/removes only distinct elements

Table "R"

v1
1
2
2

Table "S"

v2
2
3

UNION

ALL

v1
1
2
2
2
3

INTERSECT

ALL

v1
2

EXCEPT

ALL

v1
1
2

S 只有 1 个 2,
只 Intersect with
1 个 2 in R.

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

- Operations

- Duplication

- Example

- Exercises

Join Operations

Set Operations

Example

Find all pizza names with both cheese and chilli as the ingredients.

Cheese Only

```
SELECT pizza
FROM Recipes
WHERE ingredient
= 'Cheese';
```

pizza
Diavola
Margherita
Siciliana

Chilli Only

```
SELECT pizza
FROM Recipes
WHERE ingredient
= 'Chilli';
```

pizza
Diavola
Marinara

Cheese and Chilli

```
SELECT pizza
FROM Recipes
WHERE ingredient
= 'Cheese';
INTERSECT
SELECT pizza
FROM Recipes
WHERE ingredient
= 'Chilli';
```

pizza
Diavola

Notes

This is an example of *decomposition* of problem into smaller problems.

1. Find all ... with cheese only
2. Find all ... with chilli only



Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

- Operations

- Duplication

- Example

- Exercises

Join Operations

Set Operations

Exercise #3

Find all pizza names with *either* cheese or chilli (*or both*) as the ingredients.

```
1 -- What operations do you need?  
2 SELECT pizza  
3 FROM Recipes  
4 WHERE ingredient  
5 .... = 'Cheese'  
6 UNION  
7 SELECT pizza  
8 FROM Recipes  
9 WHERE ingredient  
10 .... = 'Chilli';  
11
```



Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

- Operations

- Duplication

- Example

- Exercises

Join Operations

Set Operations

Exercise #4

Find all customer names that does not like any pizza.

```
1 -- What operations do you need?  
2  
3 SELECT cname FROM Customers  
4 EXCEPT  
5 SELECT cname FROM Likes;
```

small world assumption:

If you can't find it in the table, it must be false.

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

Join Operations

- *Restriction*

- *Multi-Relational*

- *Natural Join*

- *Outer Join*

Join Operations

A SQL query goes into a bar,
walks up to two tables and asks...

“Can I join you?”

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

Join Operations

- *Restriction*

- *Multi-Relational*

- *Natural Join*

- *Outer Join*

Join Operations

Restriction

- A table name (*after renaming*) can appear at most once in the **FROM** clause
- We can rename a table to allow the same table to appear more than once in the **FROM** clause (*but with different name*)

Example

Question

Find all pair of restaurant (R_1, R_2) that are in the same area such that $R_1 < R_2$.

Error

```
SELECT rname, rname
FROM Restaurants, Restaurants
WHERE rname < rname
AND area = area;
```

No Error

```
SELECT r1.rname, r2.rname
FROM Restaurants AS r1, Restaurants r2
WHERE r1.rname < r2.rname
AND r1.area = r2.area;
```

[#]The keyword **AS** here is *optional*.

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

Join Operations

- Restriction

- Multi-Relational

- Natural Join

- Outer Join

Join Operations

Multi-Relational Queries

- Most common in practice: cross product + selection condition + attribute selection

Example

For all pizzas with cheese as its ingredient, find the restaurant that sells the pizza as well as the price of the pizza.

Equivalent Queries

```
SELECT DISTINCT rname, price  
FROM Sells S, Recipes R  
WHERE S.pizza = R.pizza  
AND R.ingredient = 'Cheese';
```

(Note: The original query has a circled 'AND' removed.)

```
SELECT DISTINCT rname, price  
FROM Sells S JOIN Recipes R  
ON S.pizza = R.pizza  
WHERE R.ingredient = 'Cheese';
```

(Note: A red arrow points from the circled 'ON' to the circled 'AND' in the first query, with the handwritten note "can be 'and'" below it.)

rname	price
Corleone Corner	19
Corleone Corner	24
Gambino Oven	16
Pizza King	17

[#]The operation **JOIN** is equivalent to the more explicit operation **INNER JOIN**. In other words, the use of **INNER** is *optional*.

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

Join Operations

- *Restriction*

- *Multi-Relational*

- **Natural Join**

- *Outer Join*

Join Operations

Natural Join

- Check for equivalence of common attributes

Example

For all pizzas with cheese as its ingredient, find the restaurant that sells the pizza as well as the price of the pizza.

Natural Join

```
SELECT DISTINCT rname, price
FROM   Sells S NATURAL JOIN Recipes R
-- no need S.pizza = R.pizza, why?
WHERE  R.ingredient = 'Cheese';
```

rname	price
Corleone Corner	19
Corleone Corner	24
Gambino Oven	16
Pizza King	17

| This is why *universal schema assumption* is useful!

Basic SQL Queries

Simple Queries

SELECT Clause

WHERE Clause

Set Operations

Join Operations

- *Restriction*

- *Multi-Relational*

- *Natural Join*

- *Outer Join*

Join Operations

Outer Join

- Recall that sometimes the "dangling tuples" are of interest

Example

For all customer names that does not like any pizza.

Outer Join

```
SELECT DISTINCT C cname
FROM Customers C LEFT JOIN Likes L
ON C cname = L cname
WHERE L pizza IS NULL; -- keep only dangling tuples
```

cname
Adi
Willie
Yoga

[#]Of course, using EXCEPT is probably easier. This shows that there are multiple solutions for the same problem.

Composing

$$P \rightarrow q \equiv \neg P \vee q.$$

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

Subqueries / Nested Queries

Preliminaries

- Subqueries are queries *within* another queries
 - Can appear in **SELECT** clause
 - Can appear in **FROM** clause (*for next week*)
 - Can appear in **WHERE** clause
- Can be used to construct complex queries
 - Subquery expressions:

IN/NOT IN	subqueries
EXISTS/NOT EXISTS	subqueries
ANY/ <u>SOME</u>	subqueries (<i>SOME is not the keyword in postgresql</i>)
ALL	subqueries

Q1 UNION Q2

Q1, Q2 are not subqueries



Composing

Subqueries

Scalar

- Definition

- Examples

[NOT] IN
ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

Scalar Subqueries

Definition

A **scalar subquery** is a query that returns at most a single value (i.e., a table with 1 row and 1 column).

Note:

- This is dynamically checked at run-time
- Can be used as a value in queries
- If the result returns 0 row, then it is treated as NULL

Examples

Scalar



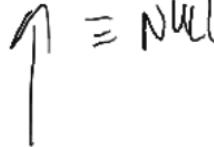
attr
CS2102



Scalar



attr



Not Scalar

attr
1
2



Not Scalar

attr1	attr2
CS2102	A+



Composing

Subqueries

Scalar

- Definition

- Examples

[NOT] IN
ANY/SOME

ALL
[NOT] EXISTS

Scoping
Nested
Remarks

Scalar Subqueries

Example #1

Find all area, the pizza sold in the area, and all the prices of the pizza in that area.

Using Scalar Subquery

```
SELECT (SELECT area
        FROM Restaurants R
        WHERE R.rname = S.rname),
       pizza, price
      FROM Sells S;
```

Note

Alternatives using joins are also possible.

area	pizza	price
North	Diavola	24
North	Hawaiian	25
North	Margherita	19
Central	Siciliana	16
Central	Funghi	23
South	Marinara	22
East	Diavola	17
East	Hawaiian	21

Composing

Subqueries

Scalar

- Definition

- Examples

[NOT] IN
ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

Scalar Subqueries

Example #2

Find all restaurant names, pizza, and the price of the pizza that sells pizza cheaper than the price of Margherita sold by Corleone Corner.

Using Scalar Subquery

```
SELECT *
FROM   Sells
WHERE  price < (SELECT S.price
                  FROM   Sells S
                  WHERE  rname = 'Corleone Corner'
                  AND   pizza = 'Margherita');
```

rname	pizza	price
Gambino Oven	Siciliana	16
Pizza King	Diavola	17

Composing

Subqueries

Scalar

[NOT] IN

- Syntax/Semantics

- Examples

ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

[NOT] IN Subqueries

Syntax

:

WHERE <expr> [NOT] IN [<subquery> | <tuple>];

Note:

- <tuple> specification is (value₁, value₂, ..., value_n)

Semantics

- For the query to work, the subquery **must** return exactly one column
- IN returns TRUE if <expr> matches any subquery row

- NOT IN returns TRUE if <expr> matches no subquery row
 - $\exists \text{row } \in \text{<subquery>} : \text{row} = \text{<expr>}$
 - $\forall \text{row } \in \text{<subquery>} : \text{row} \neq \text{<expr>}$

= <>

IS DISTINCT FROM

*Remember, we have *three-valued logic*.

$\exists s \in \phi$
 $\forall s \in \phi$

Vacuously
True/False

Composing

Subqueries

Scalar

[NOT] IN

- Syntax/Semantics

- Examples

ANY/SOME

ALL

[NOT] EXISTS

Scoping



Nested

Remarks

[NOT] IN Subqueries

Example #1

Find all restaurant names that sells pizza with Cheese as one of its ingredients.

Outer Query

```
SELECT DISTINCT rname  
FROM   Sells  
WHERE  pizza IN
```

Inner Query

all pizza w/
cheese

```
(SELECT pizza  
FROM   Recipes  
WHERE  ingredient = 'Cheese');
```

Result

rname
Corleone Corner
Gambino Oven
Pizza King

Quiz #1

Which pizza is being referred to in the inner query? Is it
(A) Sells.pizza or (B) Recipes.pizza?

Composing

Subqueries

Scalar

[NOT] IN

- Syntax/Semantics

- Examples

ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

[NOT] IN Subqueries

Example #2

Find all customer names and area such that the customer do not like any pizza.

Outer Query

```
SELECT cname, area  
FROM Customers  
WHERE cname NOT IN
```

Inner Query

```
(SELECT cname  
FROM Likes);
```



Result

cname	area
Willie	North
Adi	NULL
Yoga	NULL

Rule of Thumb

- IN subqueries can typically be replaced with inner joins
- NOT IN subqueries can typically be replaced with outer joins

Composing

Subqueries

Scalar

[NOT] IN

- Syntax/Semantics

- Examples

ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

[NOT] IN Subqueries

Example #3

Find all pizza and price sold by either Gambino Oven or Pizza King.

Explicit Tuple

```
SELECT pizza, price  
FROM Sells  
WHERE rname IN ('Gambino Oven', 'Pizza King');
```



Result

pizza	price
Siciliana	16
Diavola	17
Hawaiian	21



Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

- Syntax/Semantics

- Examples

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

ANY/SOME Subqueries

Syntax

:

WHERE <expr> <op> ANY <subquery>;

Note:

- $<\text{op}>$ are relational operators (e.g., $=$, \neq , $<$, \leq , $>$, \geq)
- SOME is used in some other DBMS (PostgreSQL uses ANY)

price < ANY (

↓

Q₁

)

↓

20

(20, 21, 19)

$(20 < 20) \vee (20 < 21) \vee (20 < 19)$

Semantics

- For the query to work, the subquery **must** return exactly one column
- The expression $<\text{expr}>$ is compared to each row from subquery using the operator $<\text{op}>$
- ANY returns TRUE if comparison evaluates to TRUE for at least one row in the subquery
 - $\exists \text{row} \in \text{subquery}. (<\text{expr}> <\text{op}> \text{row}) = \text{TRUE}$

*Remember, we have **three-valued logic**.

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

- Syntax/Semantics

- Examples

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

ANY/SOME Subqueries

Example

Find all restaurants that sells any pizza cheaper than any pizza sold by Lorenzo Tavern

(also possible, sells any pizza cheaper than the most expensive pizza sold by Lorenzo Tavern).

Outer Query

```
SELECT DISTINCT rname  
FROM   Sells  
WHERE  price <
```

Result

rname
Pizza King
Mammas Place
Gambino Oven
Corleone Corner

Inner Query

```
ANY (SELECT price  
      FROM   Sells  
     WHERE  rname = 'Lorenzo Tavern');
```



Composing

Subqueries

Scalar

[NOT] IN
ANY/SOME

ALL

- Syntax/Semantics

- Examples

[NOT] EXISTS

Scoping

Nested

Remarks

ALL Subqueries

Syntax

:

WHERE <expr> <op> ALL <subquery>;

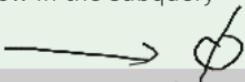
Note:

- <op> are relational operators (e.g., =, <>, <, <=, >, >=)



Semantics

- For the query to work, the subquery **must** return exactly one column
- The expression <expr> is compared to each row from subquery using the operator <op>
- ALL returns TRUE if comparison evaluates to TRUE for all row in the subquery
 - $\forall \text{row} \in \text{subquery} : (\text{expr} \text{ } \text{op} \text{ } \text{row}) = \text{TRUE}$



NULL

[#]Remember, we have *three-valued logic*.

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

- Syntax/Semantics

- Examples

[NOT] EXISTS

Scoping

Nested

Remarks

ALL Subqueries

Example

Find all restaurants that sells any pizza more expensive than all pizzas sold by Mammas Place (*also possible, sells any pizza more expensive than the most expensive pizza sold by ...*).

Outer Query

```
SELECT DISTINCT rname  
FROM   Sells  
WHERE  price >
```

Result

rname
Lorenzo Tavern
Corleone Corner

Inner Query

```
ALL (SELECT price  
      — FROM   Sells  
      WHERE  rname = 'Mammas Place');
```



Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

-Syntax/Semantics

-Examples

Scoping

Nested

Remarks

[NOT] EXISTS Subqueries

Syntax

:

WHERE [NOT] EXISTS <subquery>;

Semantics

- The subquery may return any number of columns
- EXISTS returns TRUE if the subquery returns at least one row
 - $\exists \text{ row} \in \text{subquery} : \text{TRUE}$
- NOT EXISTS returns TRUE if the subquery returns no row
 - $\forall \text{ row} \in \text{subquery} : \text{FALSE}$

$\phi \rightarrow \text{False}$

$\phi \rightarrow \text{True}$

*You read that right, there is no operation within the "forall" (i.e., \forall) and "there exists" (i.e., \exists)

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

- Syntax/Semantics

- Examples

Scoping

Nested

Remarks

[NOT] EXISTS Subqueries

Example #1

Find all areas for which Diavola is sold by at least one restaurant.

Outer Query

```
SELECT DISTINCT area  
FROM Restaurants R  
WHERE
```

Inner Query

```
EXISTS (SELECT 1 - ???  
        FROM Sells S  
        WHERE S.rname = R.rname  
          AND S.pizza = 'Diavola');
```

Result

area
East
North

Quiz #2

Why do we use SELECT 1?

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

- Syntax/Semantics

- Examples

Scoping

Nested

Remarks

[NOT] EXISTS Subqueries

Example #1

Find all areas for which Diavola

Outer Query

```
SELECT DISTINCT area  
FROM Restaurants R  
WHERE
```

Result

area
East
North

Notes

Let's fix R.rname to Pizza King. What is the result of the inner query?

```
SELECT 1 FROM Sells S  
WHERE S.rname = 'Pizza King'  
AND S.pizza = 'Diavola';
```

Inner Query

IN ('Carteone Corner',
'Pizza King')

EXISTS (SELECT 1 -- ???

```
FROM Sells S  
WHERE S.rname = R.rname  
AND S.pizza = 'Diavola');
```

r name
Sells
pizza =
'Diavola'

Quiz #2

Why do we use SELECT 1?

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

- Syntax/Semantics

- Examples

Scoping

Nested

Remarks

[NOT] EXISTS Subqueries

Example #2

Find all customer names and areas for which there is no restaurant in the area.

Outer Query

```
SELECT *
FROM Customers C
WHERE
```

Inner Query

```
NOT EXISTS (SELECT 1
              FROM Restaurants R
              WHERE C.area = R.area);
```

Result

cname	area
Homer	West
Adi	NULL
Yoga	NULL

Quiz #3

Why is NULL part of the result? Can we remove this and if so, how? (*Hint: it's not IS NOT DISTINCT FROM*)

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

Scoping

Correlated Subqueries

- Inner query uses value from outer query
 - Uncorrelated subqueries for [NOT] EXISTS are usually either wrong or unnecessary
- Result of inner query typically depends on the value of the outer query
 - This may result in potentially slow performance

Naming Ambiguities

- Same attribute names in both inner and outer queries (e.g., price)
- Resolve ambiguities using table aliases (e.g., S1 and S2)
 - Otherwise, scoping rules applies

```
SELECT DISTINCT S1.rname
FROM   Sells S1
WHERE  price < ANY
       (SELECT price
        FROM   Sells S2
        WHERE  S2.rname = 'Lorenzo Tavern');
```



Composing

Subqueries

Scalar

[NOT] IN
ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

Scoping

Scoping Rules

- A table alias declared in a (*sub*-)query Q can only be used in Q or subqueries nested within Q
- If the same table alias is declared in both subquery Q_1 and in outer query Q_0 (*or not at all*), the declaration in Q_1 is applied

General Rule:

- From inner to outer queries in case of multiple nestings

SFW Q_1
 (\underline{SFW})



Composing

Subqueries

Scalar

[NOT] IN
ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

Scoping

Scoping Rules

- A table alias declared in a (*sub*)-query Q can only be used in Q or subqueries nested within Q
- If the same table alias is declared in both subquery Q_1 and in outer query Q_0 (*or not at all*), the declaration in Q_1 is applied

General Rule:

- From inner to outer queries in case of multiple nestings

Examples

```
SELECT DISTINCT S1.rname
FROM   Sells S1
WHERE  price < ANY
       (SELECT price
        FROM   Sells S2
        WHERE  S2.rname = 'Lorenzo Tavern');
```

Notes

Table aliases are used

- $S1.rname$ from $S1$
- $S2.rname$ from $S2$



Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

Scoping

Scoping Rules

- A table alias declared in a (*sub*)-query Q can only be used in Q or subqueries nested within Q
- If the same table alias is declared in both subquery Q₁ and in outer query Q₀ (*or not at all*), the declaration in Q₁ is applied

General Rule:

- From inner to outer queries in case of multiple nestings

Examples

```
SELECT DISTINCT S1.rname
FROM   Sells S1
WHERE  price < ANY -- $1.price
       (SELECT price -- $2.price
        FROM   Sells S2
        WHERE  S2.rname = 'Lorenzo Tavern');
```

Notes

Table aliases are **NOT** used and should refer to the "closest" scope (*from inside to outside*)

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

Confusion

- Mistakes

Nested

Remarks

Scoping

Possible Confusion

Example #1

For each restaurant, find the *most expensive pizza* (*if there are multiple, return all*).

Correct

```
SELECT *
FROM   Sells S1
WHERE  S1.price >= ALL
       (SELECT S2.price
        FROM   Sells S2
        WHERE  S2.rname = S1.rname);
```

rname	pizza	price
Corleone Corner	Hawaiian	25
Gambino Oven	Siciliana	16
:	:	:

Incorrect

```
SELECT *
FROM   Sells S1
WHERE  S1.price >= ALL
       (SELECT S2.price
        FROM   Sells S2
        WHERE  S2.rname = S1.rname);
```

rname	pizza	price
Corleone Corner	Hawaiian	25

| Where are the rest?

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

Scoping

Possible Confusion

Example #1

For each restaurant, find the *most expensive pizza* (*if there are multiple, return all*).

Correct

```
SELECT *
FROM   Sells S1
WHERE  S1.price >= ALL
       (SELECT S2.price
        FROM   Sells S2
        WHERE  S2.rname = S1.rname);
```

Where Are The Rest?

`S2.rname = rname` is basically `S2.rname = S2.rname` if we follow the scoping rules. But `S2.rname = S2.rname` is always true for non-NULL values.

Incorrect

```
SELECT *
FROM   Sells S1
WHERE  S1.price >= ALL
       (SELECT S2.price
        FROM   Sells S2
        WHERE  S2.rname = rname);
```

rname	pizza	price
Corleone Corner	Hawaiian	25
Gambino Oven	Siciliana	16
:	:	:

rname	pizza	price
Corleone Corner	Hawaiian	25

| Where are the rest?

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

East → (East)

Scoping

Possible Mistakes

Example #2

Find all customer names with a known area located in an area *without* a restaurant.

Good

```
SELECT T1 cname  
FROM Customers T1  
WHERE T1.area NOT IN  
(SELECT T2.area  
FROM Restaurants T2);
```

cname
Homer

Okay-ish (can be confusing)

```
SELECT T cname  
FROM Customers T  
WHERE T.area NOT IN  
(SELECT T.area  
FROM Restaurants T);
```

cname
Homer

Wrong

```
SELECT T1 cname  
FROM Customers T1  
WHERE T1.area NOT IN  
(SELECT T1.area  
FROM Restaurants T2);
```

cname

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

- Correlated

- Rules

- Confusion

- Mistakes

Nested

Remarks

Scoping

Possible Mistakes

Example #2

Find all customer names with a known area located in an area *without* a restaurant.

Good

```
SELECT T1.cname  
FROM Customers T1  
WHERE T1.area NOT IN  
(SELECT T2.area  
FROM Restaurants T2);
```

Okay-ish (*can be confusing*)

```
SELECT T.cname  
FROM Customers T  
WHERE T.area NOT IN  
(SELECT T.area  
FROM Restaurants T);
```

Wrong

```
SELECT T1.cname  
FROM Customers T1  
WHERE T1.area NOT IN  
(SELECT T1.area  
FROM Restaurants T2);
```

cname
Homer

cname
Homer

cname

Quiz #4

What happened to the wrong code?

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

Multiple Nested Subqueries

Example

Find all customers in Central area that likes at least one pizza sold by Corleone Corner
(exclude customers that do not like any pizza).

Outer Query

```
-- Customers in central
SELECT C.cname
FROM Customers C
WHERE area = 'Central'
AND C.cname IN
```

Result

cname
Ralph

Inner Query #1

```
-- Likes at least one pizza
(SELECT L cname
FROM Likes L
WHERE L.pizza IN
);
```

Inner Query #2

```
-- Sold by Corleone Corner
(SELECT S pizza
FROM Sells S
WHERE S.rname
= 'Corleone Corner')
```

Composing

Subqueries

Scalar

[NOT] IN

ANY/SOME

ALL

[NOT] EXISTS

Scoping

Nested

Remarks

Remarks

Required Constructs

- Not all constructs are *absolutely* required

IN \equiv =ANY

:

WHERE <expr> IN <subquery>

\equiv

:

WHERE <expr> = ANY <subquery>

IN

=ANY

ANY

\equiv

EXISTS

\equiv

:
WHERE <e1> <op> ANY (
 SELECT <e2>
 FROM <rel>
 WHERE <cond>
)

:

WHERE EXISTS(
 SELECT *
 FROM <rel>
 WHERE <cond> AND <e1> <op> <e2>
)

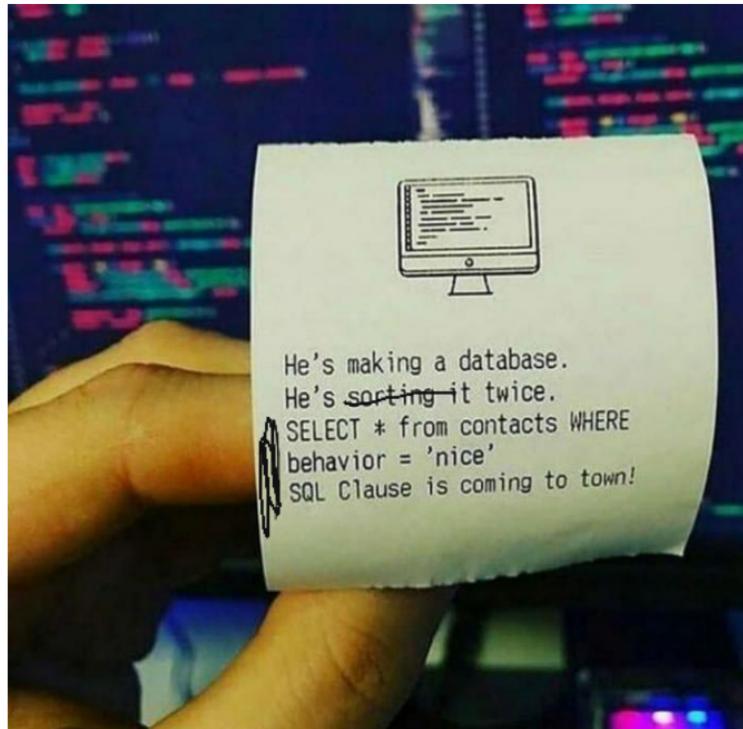
Ordering

Ordering

ORDER BY

- Preliminary
 - Multi-Attribute
 - Usage
- LIMIT & OFFSET

ORDER BY



Ordering

ORDER BY

- Preliminary

- Multi-Attribute

- Usage

LIMIT & OFFSET

ORDER BY

Preliminary

- By default, order of rows in a table is left up to the DBMS
- Sorting of rows by a single attribute can be done by ORDER BY
 - ORDER BY <attribute> ASC : Ascending Order (*default, ASC can be omitted*)
 - ORDER BY <attribute> DESC: Descending Order
- If duplicate removal is needed (*e.g., DISTINCT*), the attribute being sorted **must** appear in SELECT clause
- Sorting with respect to multiple attributes and different ordering are supported

Example

Find all restaurants and their pizzas and prices sorted in descending order of price.

SQL

```
SELECT *
FROM   Sells
ORDER BY price DESC;
```

Ordering

ORDER BY

- Preliminary
 - Multi-Attribute
 - Usage
- LIMIT & OFFSET

ORDER BY

Multi-Attribute Sorting

Example

Find all pizza with recipes containing Cheese sorted in descending order of pizza and for each pizza sort in ascending order of ingredient.

SQL

```
SELECT *
FROM Recipes
WHERE pizza IN
(SELECT pizza
 FROM Recipes
 WHERE ingredient = 'Cheese')
ORDER BY pizza DESC,
          ingredient ASC;
```

Result

pizza	ingredient
Siciliana	Anchovies
Siciliana	Caspers
Siciliana	Cheese
Margherita	Cheese
Margherita	Tomato
Diavola	Cheese
Diavola	Chilli
Diavola	Salami

Ordering

ORDER BY

- Preliminary
- Multi-Attribute
- Usage

LIMIT & OFFSET

ORDER BY

Real-World Usage

Ranking

Pos	Club	P	W	D	L	GD	Pts
1	Liverpool	28	21	6	1	49	69
2	Manchester City	28	22	2	4	55	68
3 *	Tottenham Hotspur	28	20	0	8	26	60
4	Arsenal	28	17	5	6	22	56
5	Manchester United	28	16	7	5	19	55
6 *	Chelsea	27	16	5	6	18	53
7 *	Wolverhampton Wanderers	28	11	7	10	0	40
8 *	Watford	28	11	7	10	-1	40
9 *	Everton	28	10	6	12	0	36
10 *	West Ham United	28	10	6	12	-6	36
11	Leicester City	28	10	5	13	-5	35

Ordering

ORDER BY
LIMIT & OFFSET

- Ranking
- Offset
- Usage

LIMIT & OFFSET

Ranking

- Returning only a portion of the result table
 - LIMIT k Return the "first"* k rows of the result table
 - OFFSET i Specify the position of the "first" row to be considered
- LIMIT and OFFSET are typically meaningful *only in combination with ORDER BY*

Example

Find the 3 cheapest pizza, the restaurant selling them, and the price.

```
SELECT *
FROM Sells
ORDER BY price ASC
LIMIT 3;
```

Result

rname	pizza	price
Gambino Oven	Siciliana	16
Pizza King	Diavola	17
Corleone Corner	Margherita	19

*If there are multiple "first" k (e.g., same ranking), then the order within this ordering is undetermined.

Ordering

ORDER BY
LIMIT & OFFSET

- Ranking
- Offset
- Usage

LIMIT & OFFSET

Offset

Example

Find the "next" 3 cheapest pizza, the restaurant selling them, and the price.

```
SELECT *
FROM   Sells
ORDER BY price ASC
OFFSET 3
LIMIT  3;
```

Result

rname	pizza	price
Pizza King	Hawaiian	21
Mammas Place	Marinara	22
Lorenzo Tavern	Funghi	23

Quiz #5

Can you find the 3 cheapest pizza **BUT** ordered in descending order of price?

Ordering

ORDER BY LIMIT & OFFSET

- Ranking
- Offset
- Usage

LIMIT & OFFSET

Real-World Usage

Pagination



A screenshot of a database object list titled "Objects". The list contains six rows of objects, each represented by a small icon and some text. To the right of each row is a three-dot menu icon. At the bottom of the list, there is a navigation bar with numbers 1 through 10. The number '1' is highlighted with a blue background and white text. To the right of the navigation bar, there is a "Show:" dropdown menu set to "5 rows". Above the list, there is a search bar with the placeholder "Search objects..." and a "New Object" button.

Summary

Summary

Summary

Summary

Query

- Declaration query language
 - Describe **what** the output looks like and **NOT how** to compute the output
- Build on top of Relational Algebra

Topics

- Basic SELECT ... FROM ... WHERE
- Set UNION, INTERSECT, EXCEPT (*and the ALL variants*)
- Join [INNER] JOIN, [LEFT | RIGHT | FULL] [OUTER] JOIN (*and NATURAL*)
- Subqueries [NOT] IN, ANY, ALL, [NOT] EXISTS
- Ordering ORDER BY, LIMIT, OFFSET

Notes

Basic, set, and join queries have a *direct mapping* to RA expression (*apart from set vs multiset*).

Aggregates

Common Table Expression

Universal Quant

```
postgres=# exit
```

```
Press any key to continue . . .
```

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Simple Queries

Complex Condition

Exercise #1

Find the restaurants and the pizzas that they sell where either (i) the price is over \$20 or the pizza is Siciliana and (ii) the pizza is not Diavola.

rname	pizza
Corleone Corner	Hawaiian
Gambino Oven	Siciliana
Lorenzo Tavern	Funghi
Mammas Place	Marinara
Pizza King	Hawaiian

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

WHERE Clause

Pattern Matching

Basic Pattern Matching ([NOT] LIKE)

- `_` matches any single character
- `%` matches any sequence of 0 or more characters

Exercise 2

Which of the following conditions below returns TRUE?

Choice	Comment	
A 'abc' LIKE 'abc'	YES: exact match!	✓
B 'abc' LIKE 'a%'	YES: but can also accept 'abcd'	✓
C 'abc' LIKE '_b_'	YES: 3 characters where middle is b	✓
D 'abc' LIKE '_c'	NO: this is only two characters	✗

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Set Operations

Exercise #3

Find all pizza names with *either* cheese or chilli (*or both*) as the ingredients.

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Set Operations

Exercise #4

Find all customer names that does not like any pizza.

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #1

Which `pizza` is being referred to in the *inner query*? Is it (A) `Sells.pizza` or (B) `Recipes.pizza`?

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Outer Scope

```
SELECT DISTINCT rname
  FROM Sells
 WHERE pizza IN
    ( SELECT pizza
      FROM Recipes
     WHERE ingredient = 'Cheese' );
```

The outer query refers to
`Sells.pizza`

Inner Scope

```
SELECT DISTINCT rname
  FROM Sells
 WHERE pizza IN
    ( SELECT pizza
      FROM Recipes
     WHERE ingredient = 'Cheese' );
```

The inner query refers to
`Recipes.pizza`

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Quiz #2

Quiz #2

Why do we use `SELECT 1`?

Solution

`EXISTS` and `NOT EXISTS` only check if the table has any row or not. It does not matter what the content of the table is. Try the following:

```
SELECT 1 FROM Restaurants;
```

What do you see?

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Quiz #3

Quiz #3

Why is NULL part of the result? Can we remove this and if so, how? (*Hint: it's not IS NOT DISTINCT FROM*)

Solution

As stated, we cannot simply replace = with IS NOT DISTINCT FROM. So the following does not work.

```
SELECT * FROM Customers C  
WHERE NOT EXISTS (SELECT 1 FROM Restaurants R WHERE C.area = R.area);
```

Instead, we should add another condition to ensure that when `C.area` is `NULL`, it is excluded from the result (*since we use NOT EXISTS, we should put OR C.area IS NULL*).

```
SELECT * FROM Customers C  
WHERE NOT EXISTS (SELECT 1 FROM Restaurants R WHERE C.area = R.area  
OR C.area IS NULL);
```

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Quiz #4

Quiz #4

What happened to the wrong code?

Solution

The code is reproduced on the right. Since we use **T1.area** in the *inner query*, it refers directly to the **area** of **Customers**. Notice also that at Line 3, we check if **T1.area NOT IN (...)**. We can then see that **x IN (x)** is always true and **x NOT IN (x)** is always false.

Because **x NOT IN (x)** is always false, then the **WHERE** clause always evaluate to false. This means, the **WHERE** clause will not match anything and the result will be empty.

```
SELECT T1.cname  
FROM Customers T1  
WHERE T1.area NOT IN  
(SELECT T1.area  
FROM Restaurants T2);
```

Solutions

Exercise #1

Exercise #2

Exercise #3

Exercise #4

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Quiz #5

Quiz #5

Quiz #5

Can you find the 3 cheapest pizza **BUT** ordered in descending order of price?

Solution

One way to solve this is really to just first find all the 3 cheapest pizza as a subquery. Then we check if the pizza is part of this solution (*using IN or EXISTS*). We can now sort this outer query in a different ordering from the inner query!

```
SELECT *
FROM   Sells
WHERE  price IN ( SELECT price FROM Sells ORDER BY price ASC OFFSET 3 LIMIT 3 )
ORDER BY price DESC;
```

Note that this is an incomplete solution and breaks down if there are duplicate price. A complete solution will be shown next week when we put the subquery in the **FROM** clause.