

CS2102: Database Systems

SQL (Part 3)

Announcements

Announcement

Midterm

- **Week 7:** Monday, 27 February 2023 (12:30 - 13:30)
- MPSH 2A & MPSH 2B
- Seating arrangement will be sent nearer to the assessment time
- **Softwares**
 - **Examplify:** for the assessment question and (*most*) answers
 - Practice using **Practice Exam - Non-Secure Block Internet**
 - Practice accessing the attached files
 - There can be attached file to the assessment (*i.e., global*)
 - There can be attached file to the question (*i.e. local*)
 - **PostgreSQL:** for testing your answer before submitting to Examplify
 - We recommend **psql** but you may use other software to access PostgreSQL too
 - Please check compatibility with Examplify
 - We will be using the same schema as in **Assignment #1** so please get yourself familiar with it

Announcement

Midterm (Rough Timeline)

1. **Download** the assessment before coming to MPSH
 - You will not be given additional time to download
 - Only download on the machine you intend to use for the assessment (*only 1 download*)
2. You can enter the venue by 12:15 (*please sit down on your assigned seat*)
 - There are seats dedicated for charging, these are limited so please be courteous
3. Open **psql** first (*may require some initial connection that may be blocked by Examplify once assessment starts*)
4. Open **Examplify** and login using your ID
 - Choose **National University of Singapore - APAC (nus)**
5. We will release the password around 12:29
 - Type in the password and start the assessment promptly
 - Except for technical error, no additional time will be granted for students who started late
 - Some consideration is given for slow load as long as you click start as soon as the password is released

Announcement

Midterm (End of Assessment)

- You have until the end of the timer to do your work
 - Unless you started late for non-technical reason
- You should wait on the **Green Screen** of successful submission
 - Do not close Examplify until you have seen this message
- Please wait in your seat until you are released
 - We will need to tally all the submissions

Announcement

Midterm (End of Assessment)

- **Topic**
 - L01 to L06 (*yes, this lecture included*)
- **Types of (possible) Questions**
 - **MCQ:** only one correct answer (*choose the most appropriate*)
 - **MRQ:** may have more than one correct answers (*choose all correct answers*)
 - **FITB:** fill in the blank
 - **SQL:** write your query using the VIEW provided in the question .pdf
 - You should check using PostgreSQL that your code runs
 - Code that cannot be run may receive an immediate 0
 - **ERD:** draw the ERD in the paper provided and submit them to one of the invigilators at the end
~~—~~ If this question appear, the actual answer will not be on Examplify
- Please do not forget to take your belonging with you

- A
- B
- C
- all of
- none

Announcement

Assignment #1

- Week 6.5: Saturday, 25 February 2023 (12:00 Noon)
- Submit only the .sql file containing the VIEW (*DDL and Data file will not be used*)
- **Late Submission Penalty**
 - 2 marks (*out of 5*) will be deducted
 - Late submission more than 1 day will not be accepted
- Solution will be posted on Sunday, 26 February 2023 to facilitate revision for midterm

$$\overline{\delta_c}(R-S) \approx \delta_c(R) - \underline{\delta_c(S)}$$

Recap

$$C := \underline{\underline{a}} < 10$$

$R(a)$

$S(b)$

$$\overline{\delta_c}(R-S) \equiv \overline{\delta_c(R)} - S \quad \begin{matrix} \uparrow \\ \cancel{\times} \end{matrix}$$

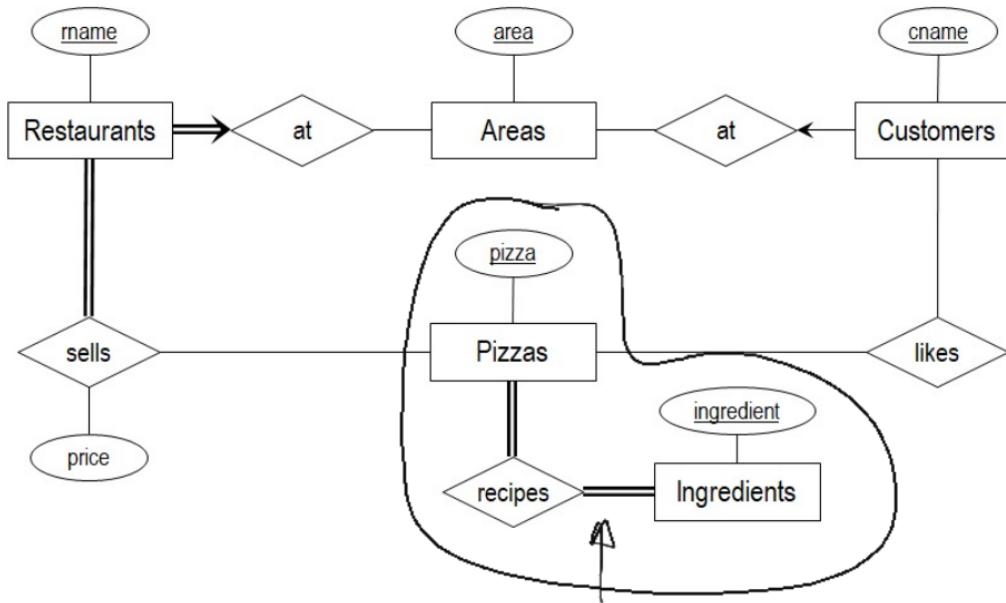
Recap

Database

- ER Diagram
- Schema
- Sample

Example Database

ER Diagram



#This is only an *approximate* to help visualize the schema.

Recap

Database

- ER Diagram
- Schema
- Sample

Example Database

Schema

Relations

- Restaurants(rname, area)
- Customers(cname, area)
- Pizzas(pizza)
- Ingredients(ingredient)
- Sells(rname, pizza, price)
- Likes(cname, pizza)
- Recipes(pizza, ingredient)

Foreign Key

- (Sells.rname) ↗ (Restaurants.rname)
- (Sells.pizza) ↗ (Pizzas.pizza)
- (Likes.cname) ↗ (Customers.cname)
- (Likes.pizza) ↗ (Pizzas.pizza)
- (Recipes.pizza) ↗ (Pizzas.pizza)
- (Recipes.ingredient) ↗ (Ingredients.ingredient)

SQL (Lo6.sql)

```
CREATE TABLE Pizzas (
    pizza      VARCHAR NOT NULL PRIMARY KEY
);
```

Recap

Database

- ER Diagram
- Schema
- Sample

Example Database

Sample Database

Table "Sells"

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Gambino Oven	Hawaiian	0
Lorenzo Tavern	Funghi	23
Mammas Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21



Table "Restaurants"

rname	area
Corleone Corner	North
Gambino Oven	Central
Lorenzo Tavern	Central
Mammas Place	South
Pizza King	East

Table "Customers"

cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Alice	Central
Bob	Central
Willie	North
Adi	NULL
Yoga	NULL

Table "Pizzas"

pizza
Diavola
Funghi
Hawaiian
Margherita
Marinara
Siciliana

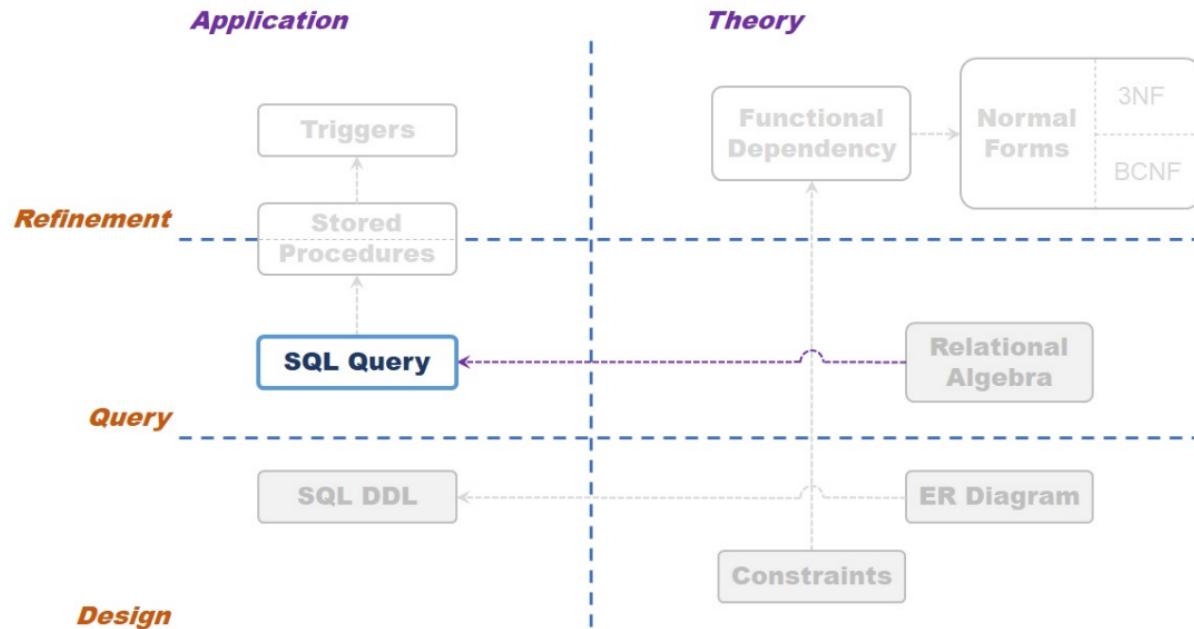
Table "Recipes"

pizza	ingredient
Diavola	Cheese
Diavola	Chilli
Diavola	Salami
Funghi	Ham
Funghi	Mushroom
Hawaiian	Ham
Hawaiian	Pineapple
Margherita	Cheese
Margherita	Tomato
Marinara	Seafood
Marinara	Chilli
Siciliana	Anchovies
Siciliana	Caspers
Siciliana	Cheese

Table "Likes"

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Siciliana
Ralph	Diavola
Alice	Siciliana
Alice	Funghi
Alice	Siciliana

Roadmap



Roadmap

Overview

Overview

Common SQL Constructs

- Aggregation
- Grouping
- Conditional Expressions

Structuring Queries

- Common Table Expressions
- Views

Extended Concepts

- Universal Quantification
- Recursive Queries

→ TO 2

Common SQL Constructs

Common SQL Constructs

Aggregation

- Functions

- Interpretation

- Signatures

Grouping

HAVING

Conditional

Aggregation

Aggregate Functions

- Aggregate functions compute a **single value** from a set of tuples
 - Example: MIN(), MAX(), AVG(), COUNT(), SUM()

Question

Find the lowest and highest price pizza among all the restaurants that sells at least one pizza.

```
SELECT MIN(price) AS lowest,  
       MAX(price) AS highest  
FROM   Sells;
```

lowest	highest
0	25

Common SQL Constructs



Aggregation

- Functions

- Interpretation

- Signatures

Grouping

HAVING

Conditional

Aggregation

Interpretation

| Let R be a non-empty relation with attribute A.

...	A	...
...	3	...
...	42	...
...	NULL	...
...	0	...
...	3	...

Query	Interpretation	Result
SELECT MIN(A) FROM R;	Minimum <u>non-NULL</u> value in A	0
SELECT MAX(A) FROM R;	Maximum <u>non-NULL</u> value in A	42
SELECT AVG(A) FROM R;	Average of <u>non-NULL</u> values in A	12 (i.e., <u>48 / 4</u>)
SELECT SUM(A) FROM R;	Sum of <u>non-NULL</u> values in A	48
SELECT COUNT(A) FROM R;	Count of <u>non-NULL</u> values in A	4
SELECT COUNT(*) FROM R;	Count of rows in A <u>in R</u>	5
SELECT AVG(DISTINCT A) FROM R;	Average of distinct non-NULL value in A	15
SELECT SUM(DISTINCT A) FROM R;	Sum of distinct non-NULL value in A	45
SELECT COUNT(DISTINCT A) FROM R;	Count of distinct non-NULL value in A	3

Common SQL Constructs

Aggregation

- Functions
- Interpretation
- Signatures

Grouping

HAVING

Conditional

Aggregation

Interpretation

Let R be an empty relation with attribute B .

...	B	...
-----	---	-----

...	A	...
...	NULL	...
...	NULL	...
:	:	:



Query	Result
SELECT MIN(B) FROM R;	NULL
SELECT MAX(B) FROM R;	NULL
SELECT AVG(B) FROM R;	NULL
SELECT SUM(B) FROM R;	NULL
SELECT COUNT(B) FROM R;	0
SELECT COUNT(*) FROM R;	0

Query	Result
SELECT MIN(A) FROM S;	NULL
SELECT MAX(A) FROM S;	NULL
SELECT AVG(A) FROM S;	NULL
SELECT SUM(A) FROM S;	NULL
SELECT COUNT(A) FROM S;	0
SELECT COUNT(*) FROM S;	n

Common SQL Constructs

Aggregation

- Functions
- Interpretation
- Signatures

Grouping
HAVING
Conditional

Aggregation

Function Signatures

Function	Input Type	Output Type
MIN	any comparable type	same as input
MAX	any comparable type	same as input
SUM	Numeric data (e.g., INT, BIGINT, REAL, etc)	$\text{SUM}(\text{INT}) \rightarrow \underline{\text{BIGINT}}$; $\text{SUM}(\text{REAL}) \rightarrow \underline{\text{REAL}}$
COUNT	any data	<u>BIGINT</u>

Common SQL Constructs

Aggregation

Grouping

- Aggregation

- Partition

- Groups

- Restrictions

- Primary Key

HAVING

Conditional

Grouping

Grouping and Aggregation

- So far, application of aggregate functions are over all tuples of a relation
 - Resulting relation only has one tuple

GROUP BY

Syntax

```
GROUP BY attr1, attr2, ...
```

Notes:

- Logical partition of relation into groups based on values for specified attributes
- In principle, always applied together with aggregation
(GROUP BY without aggregation is valid but typically not meaningful)
- Application of aggregation functions are now over each group
 - One result tuple for each group

Common SQL Constructs

Aggregation

Grouping

- Aggregation

- Partition

- Groups

- Restrictions

- Primary Key

HAVING

Conditional

Grouping

Logical Partitioning

Question

For each restaurant, find the lowest and highest price pizza.

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Gambino Oven	Siciliana	0
Lorenzo Tavern	Funghi	23
Mammas Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Query

```
SELECT rname,  
       MIN(price),  
       MAX(price)  
FROM Sells  
GROUP BY rname;
```

rname	min	max
Lorenzo Tavern	23	23
Pizza King	17	21
Mammas Place	22	22
Gambino Oven	0	16
Corleone Corner	19	25

Common SQL Constructs

Aggregation

Grouping

- Aggregation

- Partition

- Groups

- Restrictions

- Primary Key

HAVING

Conditional

Grouping

Defining Groups

- Given **GROUP BY a₁, a₂, ..., a_n** two tuple t and t' belong to the same group if
 $(t.a_1 \text{ IS NOT DISTINCT FROM } t'.a_1) \text{ AND}$
 $(t.a_2 \text{ IS NOT DISTINCT FROM } t'.a_2) \text{ AND}$
... AND ...
 $(t.a_n \text{ IS NOT DISTINCT FROM } t'.a_n)$ evaluates to True

Example

Table
 $R(A,B,C)$

A	B	C
NULL	4	19
6	1	NULL
20	2	10
1	1	2
1	18	2
NULL	21	19
6	20	NULL

```
SELECT *  
FROM R  
GROUP BY A,C;
```

A	B	C
NULL	4	19
NULL	21	19
6	1	NULL
6	20	NULL
20	2	10
1	1	2
1	18	2

Common SQL Constructs

Aggregation
Grouping

- Aggregation
- Partition
- Groups
- Restrictions**
- Primary Key
- HAVING
- Conditional

Grouping

Restrictions to SELECT Clause

Restriction

If column A_j of table R appears in the **SELECT** clause, one of the following conditions must hold

1. A_j appears in the **GROUP BY** clause
2. A_j appears as input of an aggregation function in the **SELECT** clause
3. The primary key ~~or a candidate key*~~ of R appears in the **GROUP BY** clause

Example (Valid)

```
SELECT rname, SUM(price)
FROM   Sells
GROUP BY rname;
```

Quiz #1

What is the problem with the invalid query?

Example (Invalid)

```
SELECT rname, pizza, SUM(price)
FROM   Sells
GROUP BY rname;
```

*This is valid in standard SQL but not supported in PostgreSQL. In this module, we follow PostgreSQL's tighter restriction.

Common SQL Constructs

Aggregation
Grouping

- Aggregation
- Partition
- Groups
- Restrictions
- Primary Key

HAVING

Conditional

Grouping

Grouping Over Primary Key

Consider the table "Countries" on the right.

Valid Query

```
SELECT name, population, COUNT(*)  
FROM Countries  
GROUP BY iso2;
```

Quiz #2

What is the "problem" with the *valid* query?

```
CREATE TABLE Countries (  
    iso2      CHAR(2) PRIMARY KEY,  
    name      VARCHAR(255) UNIQUE,  
    population INTEGER,  
    gdp       NUMERIC,  
    continent VARCHAR(255)  
);
```

Invalid Query

```
SELECT name, population, COUNT(*)  
FROM Countries  
GROUP BY name;
```

[#]The invalid query is *valid* in SQL standard but *invalid* in PostgreSQL even if we add NOT NULL constraint on name.

Common SQL Constructs

Aggregation Grouping

- Aggregation
- Partition
- Groups
- Restrictions
- Primary Key
- HAVING
- Conditional

Grouping

Valid Query

```
SELECT R.rname, R.area,
       COUNT (C.cname)
  FROM Restaurants R, Customers C
 WHERE R.area = C.area
 GROUP BY R.rname;
```

Quiz #3

Will it be valid if we do SELECT R.rname, C.area ...
instead of SELECT R.rname, R.area ...?

Invalid Query

```
SELECT R.rname, C cname,
       COUNT (C.cname)
  FROM Restaurants R, Customers C
 WHERE R.area = C.area
 GROUP BY R.rname;
```

rname	area	cname	area
Corleone Corner	North	Willie	North
Gambino Oven	Central	Bob	Central
Gambino Oven	Central	Alice	Central
Gambino Oven	Central	Ralph	Central
Gambino Oven	Central	Moe	Central
:	:	:	:
Pizza King	East	Maggie	East

Common SQL Constructs

Aggregation
Grouping
HAVING
- Syntax
- Example
- Exercise
- Restrictions
- Evaluation
Conditional

HAVING Clause

$\text{avg} > \text{overall avg}$
Where: remove the row where ' $\text{avg} \leq \text{overall}$ '.

Syntax

`GROUP BY attr1, attr2, ...`

`HAVING <condition>`

but then the avg changes

\Rightarrow infinite reasoning

Notes:

- Conditions check for each group defined by GROUP BY clause
- HAVING clause cannot be used without GROUP BY clause
- Conditions typically involve aggregate functions

HAVING: filter on groups

Common SQL Constructs

Aggregation

Grouping

HAVING

- Syntax

- Example

- Exercise

- Restrictions

- Evaluation

Conditional

HAVING Clause

Example

Question

Find all restaurants and the number of pizza they sells such that the restaurant sells more than 1 pizza.

Query

```
SELECT rname, COUNT(pizza)
FROM   Sells
GROUP BY rname
HAVING COUNT(pizza) > 1;
```

Result

rname	count
Pizza King	2
Gambino Oven	2
Corleone Corner	3

“Where” cannot be together with aggregation.

Common SQL Constructs

Aggregation
Grouping
HAVING
- Syntax
- Example
- Exercise
- Restrictions
- Evaluation
Conditional

HAVING Clause

Exercise #1

Question

Find all restaurants that sells any pizza cheaper than any pizza sold by Lorenzo Tavern.

Query

```
1 SELECT S.rname
2 FROM   Sells S
3 GROUP BY S.rname
4 HAVING MIN(price) < (
5     SELECT MAX(price)
6     FROM   Sells S1
7     WHERE  rname = 'Lorenzo Tavern'
8 );
```

min(pizza)

max(pizza)

Result

rname
Pizza King
Mammas Place
Gambino Oven
Corleone Corner

Common SQL Constructs

Aggregation
Grouping
HAVING
- Syntax
- Example
- Exercise
- **Restrictions**
- Evaluation
Conditional

HAVING Clause

Restrictions to HAVING Clause

Restriction

If column A_j of table R appears in the **HAVING** clause, one of the following conditions must hold

1. A_j appears in the **GROUP BY** clause
2. A_j appears as input of an aggregation function in the **HAVING** clause
3. The primary key ~~or a candidate key*~~ of R appears in the **GROUP BY** clause

Valid

```
SELECT rname, COUNT(*) FROM Sells  
GROUP BY rname HAVING AVG(price) > 20;
```

```
SELECT rname, COUNT(*) FROM Sells  
GROUP BY rname, pizza HAVING price > 20;
```

Invalid

```
SELECT rname, COUNT(*) FROM Sells  
GROUP BY rname HAVING price > 20;
```

Quiz #4

What is the result of the query above?

II
where

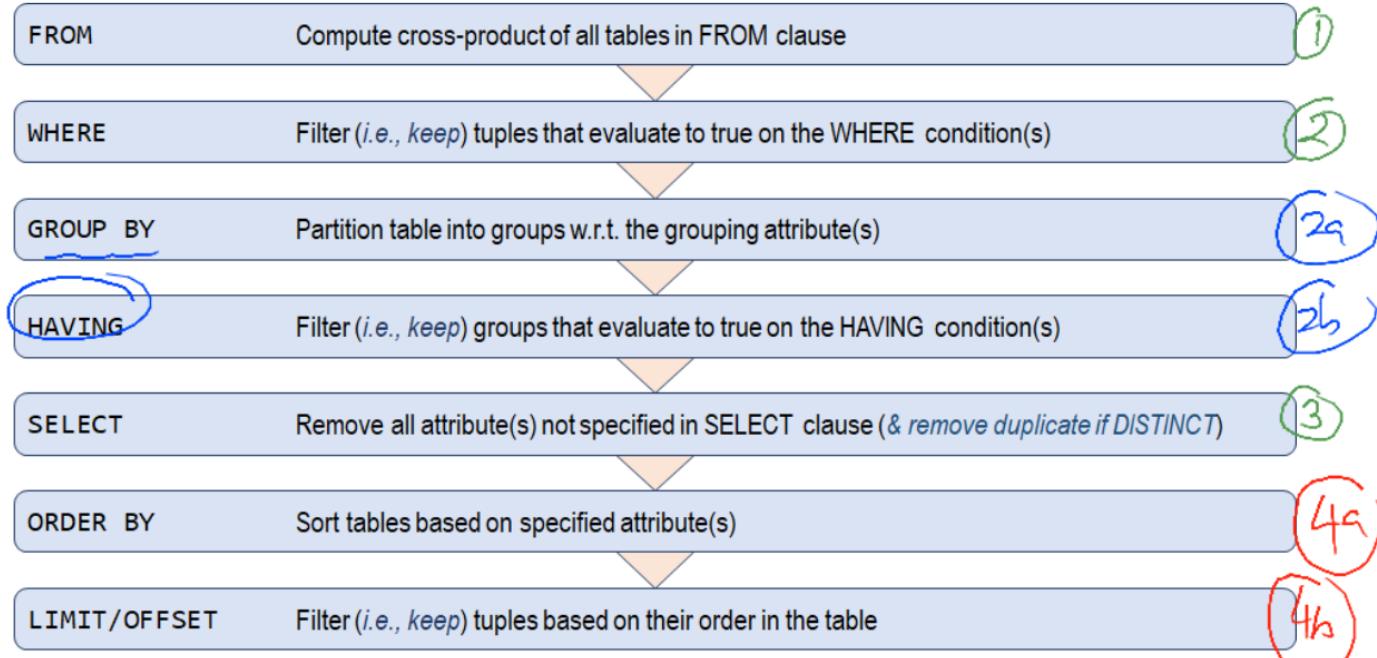
Common SQL Constructs

Aggregation
Grouping
HAVING
- Syntax
- Example
- Exercise
- Restrictions
- Evaluation
Conditional

HAVING Clause

Conceptual Evaluation of Queries

$$\pi_{[1]} \left(\sigma_{[c]} (R_1 \times R_2 \times \dots \times R_n) \right)$$





Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

CASE Expression

Syntax

```
CASE
    WHEN <condition_1> THEN <result_1>
    WHEN <condition_2> THEN <result_2>
    :
    WHEN <condition_n> THEN <result_n>
    ELSE <result_0>
END
```

Notes:

- Only one of the results will be returned
- If `<expression>` is specified, similar to `<condition>` being defined as `<expression> = <value>`
- `ELSE` is *optional*
 - Return `NULL` if no `ELSE` is specified and no condition matched

```
CASE <expression>
    WHEN <value_1> THEN <result_1>
    WHEN <value_2> THEN <result_2>
    :
    WHEN <value_n> THEN <result_n>
    ELSE <result_0>
END
```

Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

CASE Expression

Question

For each restaurant, we categorize the restaurant by the average price of their pizza.

Query

```
SELECT rname, CASE
    WHEN AVG(price) >= 23 THEN 'Expensive'
    WHEN AVG(price) >= 18 THEN 'Reasonable'
    ELSE 'Cheap'
END
FROM Sells
GROUP BY rname;
```

Category	Price
Expensive	price \geq 23
Reasonable	23 > price \geq 18
Cheap	18 > price

Table

rname	case
Lorenzo Tavern	Expensive
Pizza King	Reasonable
Mammas Place	Reasonable
Gambino Oven	Cheap
Corleone Corner	Reasonable



Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

COALESCE Expression

Syntax

```
COALESCE(<value_1>, <value_2>, <value_3>, ...)
```

Notes:

- Returns the first non-NULL value in the list of input arguments (*order of <value> matters!*)
- Returns NULL if all values in the list of input arguments are NULL

Equivalence

```
COALESCE(<value_1>,
          <value_2>,
          <value_3>)
```

≡

```
CASE
    WHEN <value_1> IS NOT NULL THEN <value_1>
    WHEN <value_2> IS NOT NULL THEN <value_2>
    ELSE <value_3>
END
```

Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

COALESCE Expression

Scenario

A student may attempt a quiz up to 3 times until they pass. Each potential attempt is recorded as follows:

- On passing
- On failing
- On no attempt

the value is recorded as 'pass'; cannot attempt another
the value is recorded as 'fail'; may attempt another
the value is recorded as NULL; may attempt another

Question

For each student, find the result of their last attempted quiz.

name	first	second	third
Alice	pass	NULL	NULL
Bob	fail	pass	NULL
Carol	fail	fail	pass
Dave	fail	NULL	fail
Eve	fail	fail	NULL

Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

COALESCE Expression

Scenario

A student may attempt a quiz up to 3 times until they pass. Each potential attempt is recorded as follows:

- On passing
- On failing
- On no attempt

the value is recorded as 'pass'; cannot attempt another
the value is recorded as 'fail'; may attempt another
the value is recorded as NULL; may attempt another

Question

For each student, find the result of their last attempted quiz.

```
SELECT name,  
       COALESCE (third, second, first) AS result -- the order matter!  
FROM Quiz;
```

? 'fail'

name	first	second	third
Alice	pass	NULL	NULL
Bob	fail	pass	NULL
Carol	fail	fail	pass
Dave	fail	NULL	fail
Eve	fail	fail	NULL

name	result
Alice	pass
Bob	pass
Carol	pass
Dave	fail
Eve	fail

Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

NULLIF Expression

Syntax

NULLIF(<value_1>, <value_2>)

Notes:

- Returns NULL if $<\text{value_1}> = <\text{value_2}>$, otherwise returns $<\text{value_1}>$

Examples

```
SELECT  
    NULLIF(1,1)  
    AS val;
```

val
NULL

```
SELECT  
    NULLIF(2,1)  
    AS val;
```

val
2

```
SELECT  
    NULLIF(NULL,1)  
    AS val;
```

val
NULL

```
SELECT  
    NULLIF(2,NULL)  
    AS val;
```

val
2

Common SQL Constructs

Aggregation

Grouping

HAVING

Conditional

- CASE

- COALESCE

- NULLIF

Conditional Expressions

NULLIF Expression

Question

Find the minimum and average price of pizza across different restaurants. However, since *unknown values* are represented as 0 (*instead of NULL*), exclude the unknown values from the computation.

Incorrect

```
SELECT MIN(price) AS min_price,  
       AVG(price) AS avg_price  
FROM   Sells;
```

min_price	avg_price
0	18.55555555555555555555

Correct

```
SELECT MIN(NULLIF(price,0)) AS min_price,  
       AVG(NULLIF(price,0)) AS avg_price  
FROM   Sells;
```

min_price	avg_price
16	20.875000000000000000

Structuring Queries

Structuring Queries

Table Expression

- Motivation

- Common

- Syntax

- Example

VIEWS

Table Expressions

Motivational Example

Question

Find the 3 cheapest pizza **BUT** ordered in descending order of price?

Incomplete Solution

```
SELECT *
FROM   Sells
WHERE  price IN (
    SELECT price FROM Sells ORDER BY price ASC LIMIT 3
)
ORDER BY price DESC;
```

- The problem is we are only looking at the price which may contains duplicate
 - Even worse, limiting this to 3 does not even give us the 2 cheapest!

Table "Sells"

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	19
Mammas Place	Marinara	19
Pizza King	Diavola	17
Pizza King	Hawaiian	19

Structuring Queries

Table Expression

- Motivation
 - Common
 - Syntax
 - Example
- VIEWS

Table Expressions

Motivational Example

Question

Find the 3 cheapest pizza **BUT** ordered in descending order of price?

Complete Solution

```
SELECT *
FROM (
    SELECT * FROM Sells
    ORDER BY price ASC LIMIT 3
) T1
ORDER BY price DESC;
```

- This is like a "temporary" table
(of course, no actual tables are actually created)
- This is not yet "common" table expression because it can only be used by outer query

Table "Sells"

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	19
Mammas Place	Marinara	19
Pizza King	Diavola	17
Pizza King	Hawaiian	19

Structuring Queries

Table Expression

- Motivation
 - Common
 - Syntax
 - Example
- VIEWS

Table Expressions

Common Computation

Question

For each pizza with Cheese as one of its ingredient , find the restaurant names that sells it and the customer names that likes it such that the customer and restaurant are in the same area.

Common Table Expressions

Compute the highlighted part as a "temporary" table.

Then we can use this "temporary" table as if it is a table.

WITH

```
PizzaWithCheese (pizza) AS (
    SELECT DISTINCT pizza FROM Recipes
    WHERE ingredient = 'Cheese'
)
SELECT S.pizza, R.rname, C cname
FROM   Sells S, Likes L, Restaurants R, Customers C
WHERE  S.rname = R.rname AND C cname = L cname
      AND S.pizza = L.pizza AND R.area = C.area
      AND S.pizza IN ( SELECT * FROM PizzaWithCheese );
```



Structuring Queries

Table Expression

- Motivation
- Common
- **Syntax**
- Example

VIEWS

Table Expressions

Common Table Expression Syntax

Syntax

WITH

```
CTE_1 AS ( Q_1 ),  
CTE_2 AS ( Q_2 ),  
    ;,  
CTE_n AS ( Q_n )  
Q_0  
-- main SELECT statement
```

Usage:

- Each CTE_i is the **name** of a temporary table defined by query Q_i
- Each CTE_i can reference any other CTE_j that has been declared before CTE_i (*i.e.*, j < i)
- The main SELECT statement Q₀ can reference *any* possible subset of all CTE_i
 - It does not have to reference all, but if there are any CTE_i not referenced by any, it can be deleted

Note:

- The goal of using CTEs is *not* to write less code
- CTEs help to improve readability, debugging, and maintenance
- The resulting column names can also be specified as
CTE_i (<attr>, <attr>, ..., <attr>) AS (Q_i)

Important Notes

You are only allowed **up to 2** CTEs for each question for assignment and midterm.

Structuring Queries

Table Expression

- Motivation
 - Common
 - Syntax
 - Example
- VIEWS

Table Expressions

Extended Example

Question

For each pizza with Cheese as one of its ingredient, find the restaurant names that sells it and the customer names that likes it such that the customer and restaurant are in the same area.

WITH

```
PizzaWithCheese (pizza) AS (
    SELECT DISTINCT pizza FROM Recipes
    WHERE ingredient = 'Cheese'
),
RestaurantWithCheese (rname, pizza, area) AS (
    SELECT DISTINCT rname, pizza, area
    FROM Sells NATURAL JOIN Restaurants
    WHERE pizza IN ( SELECT * FROM PizzaWithCheese )
),
CustomerLikesCheese (cname, pizza, area) AS (
    SELECT DISTINCT cname, pizza, area
    FROM Likes NATURAL JOIN Customers
    WHERE pizza IN ( SELECT * FROM PizzaWithCheese )
)
SELECT pizza, rname, cname
FROM RestaurantWithCheese
NATURAL JOIN
CustomerLikesCheese;
```

Structuring Queries

Table Expression

IEWS

- *Virtual Tables*

- *Syntax*

- *Example*

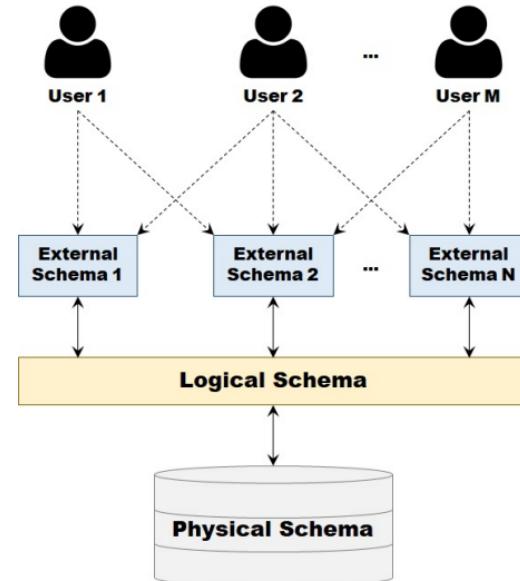
- *Usability*

IEWS

Virtual Tables

Common Observations

- Often only parts of a table (*i.e., rows/columns*) are of interest
- Often not all parts of a table (*i.e., rows/columns*) should be accessible to all users (*e.g., PDPA, GDPR, etc*)
- Often the same queries and/or subqueries are regularly and frequently used



Structuring Queries

Table Expression

IEWS

- *Virtual Tables*

- *Syntax*

- *Example*

- *Usability*

IEWS

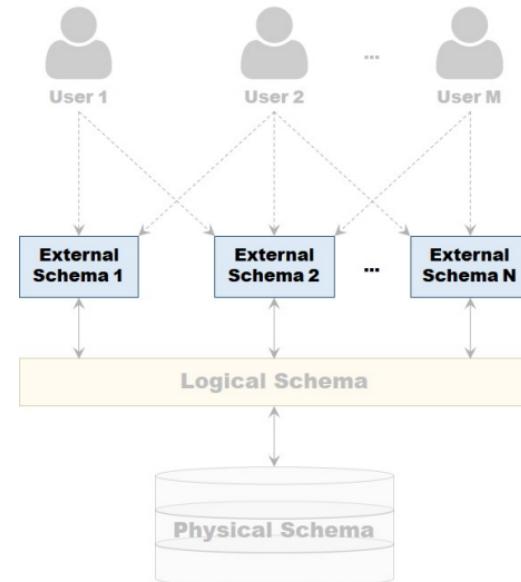
Virtual Tables

Common Observations

- Often only parts of a table (*i.e., rows/columns*) are of interest
- Often not all parts of a table (*i.e., rows/columns*) should be accessible to all users (*e.g., PDPA, GDPR, etc*)
- Often the same queries and/or subqueries are regularly and frequently used

Note

Logical schemas are tables we have created using SQL DDL (*i.e., CREATE TABLE*). External schema can be represented using VIEWS.



Structuring Queries

Table Expression

VIEWs

- Virtual Tables
- Syntax
- Example
- Usability

VIEWs

CREATE VIEW

Syntax

```
CREATE VIEW <name> AS  
    <query> -- SELECT statement  
;
```

Notes:

- A VIEW is a permanently named query (*also known as virtual table*)
 - Conceptually, **NO** actual table is created
 - The computation *may* be done each time the virtual table is accessed (*i.e., the result is conceptually not permanently stored*)
- A VIEW can be used almost like a normal table
 - Some restrictions apply (*to be discussed later*)
- The resulting column names can also be specified as
`CREATE VIEW <name> (<attr>, <attr>, ..., <attr>) AS ...`

Important Notes

You must wrap the solution of each question with the respective **VIEW** for both assignment and midterm.



Structuring Queries

Table Expression

IEWS

- Virtual Tables
- Syntax
- Example
- Usability

VIEWS

Previous Example

Question

For each **pizza with Cheese** as one of its ingredient, find the restaurant names that sells it and the customer names that likes it such that the customer and restaurant are in the same area.

- Assume finding **pizza with Cheese** is a very frequent query

```
CREATE VIEW PizzaWithCheese (pizza) AS  
    SELECT DISTINCT pizza FROM Recipes  
    WHERE ingredient = 'Cheese';
```

WITH

```
RestaurantWithCheese (rname, pizza, area) AS (  
    SELECT DISTINCT rname, pizza, area  
    FROM Sells NATURAL JOIN Restaurants  
    WHERE pizza IN ( SELECT * FROM PizzaWithCheese )  
)  
CustomerLikesCheese (cname, pizza, area) AS (  
    SELECT DISTINCT cname, pizza, area  
    FROM Likes NATURAL JOIN Customers  
    WHERE pizza IN ( SELECT * FROM PizzaWithCheese )  
)  
SELECT pizza, rname, cname  
FROM RestaurantWithCheese  
    NATURAL JOIN  
CustomerLikesCheese;
```

Structuring Queries

Table Expression

VIEWs

- Virtual Tables
- Syntax
- Example
- Usability

VIEWs

Usability

SELECT Statements

- No restrictions when used in SQL queries (*i.e.*, *SELECT statement*)

Updatable View

- Restrictions for INSERT, UPDATE, and/or DELETE statements
- Requirements:
 - A. Only one entry in FROM clause (*table or updatable view*)
 - B. No WITH, DISTINCT, GROUP BY, HAVING, LIMIT, or OFFSET
 - C. No UNION, INTERSECT, EXCEPT or the ALL variants
 - D. No aggregate functions
 - E. etc

[#]More information can be found at <https://www.postgresql.org/docs/14/sql-createview.html>

Extended Concepts

Extended Concepts

Universal

- Motivation
- Transformation
- Double Negation
- Cardinality
- Recursive

Universal Quantification

Motivational Example

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

Problems

- SQL queries is evaluated for each row
- SQL conditions in WHERE clause is checked row-by-row
 - But this is a property that may need to be satisfied by different rows

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mammas Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

cname	pizza
Homer	Hawaiian
Homer	Margherita
:	:



Notes

SQL directly supports only existential quantification (**EXISTS**)



Extended Concepts

Universal

- Motivation
- Transformation
- Double Negation
- Cardinality
- Recursive

Universal Quantification

Query Transformation

restaurant that sells all pizzas liked by 'Homer'

Note

In other words (*a logical approximation of why, can be captured as*)

$$\forall x : \text{Exist}(x) \Leftrightarrow \neg\neg\forall x : \text{Exist}(x) \Leftrightarrow \neg\exists x : \neg\text{Exist}(x)$$

\Rightarrow there does not exist pizza that 'Homer' likes and not sold by the restaurant

Useful Query

all pizzas that 'Homer' likes and restaurant with rname = R does not sell

```
SELECT L.pizza
FROM Likes L
WHERE L cname = 'Homer'
AND NOT EXISTS (
    SELECT 1
    FROM Sells S
    WHERE S pizza = L pizza
    AND S rname = 'R'
);
```

Try It Out!

Replace R with

- Corleone Corner
- Gambino Oven
- Lorenzo Tavern
- Mamas Place
- Pizza King

Extended Concepts

Universal

- Motivation
- Transformation
- **Double Negation**
- Cardinality
- Recursive

Universal Quantification

Motivational Example

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

Double Negation Solution

```
SELECT DISTINCT S1.rname
FROM   Sells S1
WHERE  NOT EXISTS (
    SELECT 1 FROM Likes L
    WHERE L cname = 'Homer'
    AND NOT EXISTS (
        SELECT 1 FROM Sells S
        WHERE S pizza = L pizza
        AND S rname = S1 rname
    )
);
```

We need to replace 'R' with all possible restaurants, so we put another query outside and connect it to the innermost subquery (*i.e.*, $S.rname = S1.rname$).

#While not overly common, SQL queries requiring universal quantification can get "ugly".

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

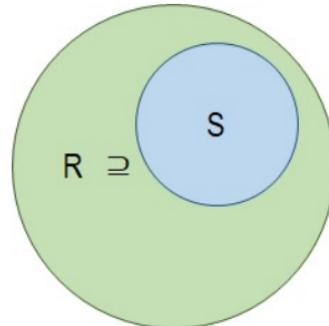
Cardinality Solution

Question

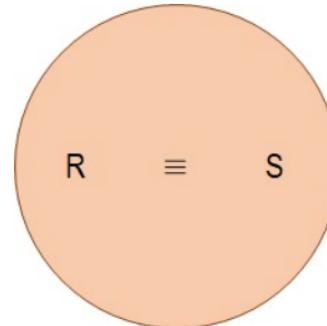
Find the name of a restaurant that sells all pizzas liked by 'Homer'.

- Let R be the set of all pizzas sold by a restaurants (e.g., 'Corleone Corner')
- Let L be the set of all pizzas liked by a customer (e.g., 'Homer')

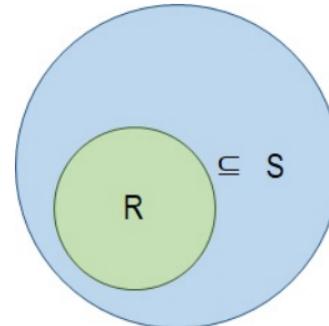
Superset



Equal



Subset



Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

- Let R be the set of all pizzas sold by a restaurants (e.g., 'Corleone Corner')
- Let S be the set of all pizzas liked by a customer (e.g., 'Homer')

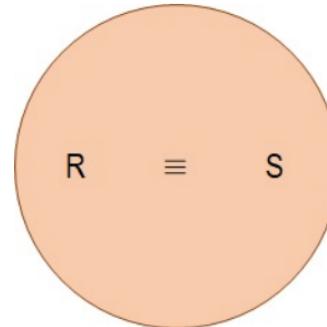
Superset

- $R \supseteq S$

$$\leftrightarrow R \cup S = R$$

$$\leftrightarrow |R \cup S| = |R|$$

Equal



Subset

- $R \subseteq S$

$$\leftrightarrow R \cap S = R$$

$$\leftrightarrow |R \cap S| = |R|$$

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

- Let R be the set of all pizzas sold by a restaurants (e.g., 'Corleone Corner')
- Let L be the set of all pizzas liked by a customer (e.g., 'Homer')

Superset

- $R \supseteq S$

$$\leftrightarrow R \cup S = R$$

$$\leftrightarrow |R \cup S| = |R|$$



Equal

- $R \equiv S$

$$\leftrightarrow R \supseteq S \wedge R \subseteq S$$

$$\leftrightarrow |R \cup S| = |R \cap S|$$



Subset

- $R \subseteq S$

$$\leftrightarrow R \cap S = R$$

$$\leftrightarrow |R \cap S| = |R|$$

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
        SELECT COUNT(DISTINCT pizza)
        FROM (
                SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
                UNION
                SELECT pizza FROM Likes WHERE cname = 'Homer'
            ) AS T1
    ) = (
        SELECT COUNT(DISTINCT pizza)
        FROM   Sells S1 WHERE S1.rname = S.rname
    );
```

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
    SELECT COUNT(DISTINCT pizza)
    FROM   (
        SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
        UNION
        SELECT pizza FROM Likes WHERE cname = 'Homer'
    ) AS T1
) = (
    SELECT COUNT(DISTINCT pizza)
    FROM   Sells S1 WHERE S1.rname = S.rname
);
```

Note

R

```
SELECT pizza
FROM   Sells S1
WHERE  S1.rname
      =
      S.rname
```

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
        SELECT COUNT(DISTINCT pizza)
        FROM (
                SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
                UNION
                SELECT pizza FROM Likes WHERE cname = 'Homer'
            ) AS T1
        ) = (
        SELECT COUNT(DISTINCT pizza)
        FROM   Sells S1 WHERE S1.rname = S.rname
    );
```

Note

S

```
SELECT pizza
FROM   Likes
WHERE  cname
      =
      'Homer'
```

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
    SELECT COUNT(DISTINCT pizza)
    FROM (
        SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
        UNION
        SELECT pizza FROM Likes WHERE cname = 'Homer'
    ) AS T1
) = (
    SELECT COUNT(DISTINCT pizza)
    FROM   Sells S1 WHERE S1.rname = S.rname
);
```

Note

R U S

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
    SELECT COUNT(DISTINCT pizza)
    FROM (
        SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
        UNION
        SELECT pizza FROM Likes WHERE cname = 'Homer'
    ) AS T1
) = (
    SELECT COUNT(DISTINCT pizza)
    FROM   Sells S1 WHERE S1.rname = S.rname
);
```

Note

|R U S|

Renaming to T1 is
needed for subquery

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
        SELECT COUNT(DISTINCT pizza)
        FROM (
                SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
                UNION
                SELECT pizza FROM Likes WHERE cname = 'Homer'
            ) AS T1
    ) = (
        SELECT COUNT(DISTINCT pizza)
        FROM   Sells S1 WHERE S1.rname = S.rname
    );
```

Note

|R|

Extended Concepts

Universal

- Motivation
 - Transformation
 - Double Negation
 - **Cardinality**
- Recursive

Universal Quantification

Cardinality Solution

Question

Find the name of a restaurant that sells all pizzas liked by 'Homer'.

Scalar {

```
SELECT DISTINCT rname
FROM   Sells S
WHERE  (
        SELECT COUNT(DISTINCT pizza)
        FROM  (
                SELECT pizza FROM Sells S1 WHERE S1.rname = S.rname
                UNION
                SELECT pizza FROM Likes WHERE cname = 'Homer'
            ) AS T1
        ) = 0
        SELECT COUNT(DISTINCT pizza)
        FROM   Sells S1 WHERE S1.rname = S.rname
    );
```

Scalar {

Note

$$|R \cup S| = |R|$$

Extended Concepts

Universal

Recursive

- MRT

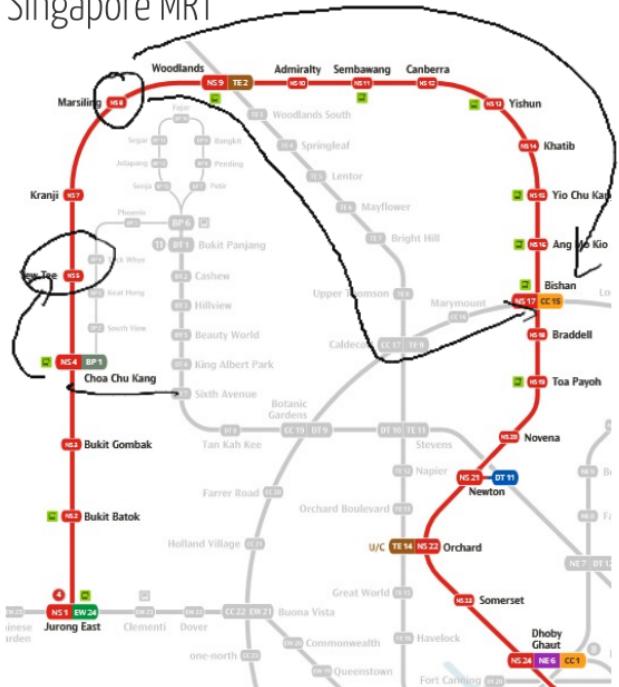
- Query

- Recursive CTE

- Lazy Evaluation

Recursive Queries

Singapore MRT



Schema

```
CREATE TABLE MRT (
    → fr_stn CHAR(5),
    → to_stn CHAR(5),
    PRIMARY KEY (fr_stn, to_stn)
);
```

Sample

fr_stn	to_stn
NS4	NS5
NS5	NS7
EW5	EW6
EW6	EW7

Extended Concepts

Universal

Recursive

- MRT

- Query

- Recursive CTE

- Lazy Evaluation

Recursive Queries

Interesting Query

Question

Find all MRT station that can be reached from NS1 in 1-stop.

Potential Solution

```
SELECT DISTINCT to_stn  
FROM    MRT  
  
WHERE   fr_stn = 'NS1';
```



Table

to_stn
NS2

1 row

Extended Concepts

Universal

Recursive

- MRT

- Query

- Recursive CTE

- Lazy Evaluation

Recursive Queries

Interesting Query

Question

Find all MRT station that can be reached from NS1 in 2-stops.

Potential Solution

```
SELECT DISTINCT M2.to_stn
FROM   MRT M1,
       MRT M2
WHERE  M1.fr_stn = 'NS1'
AND    M1.to_stn = M2.fr_stn;
```



Table

to_stn
NS3

1 row

Extended Concepts

Universal

Recursive

- MRT

- Query

- Recursive CTE

- Lazy Evaluation

Recursive Queries

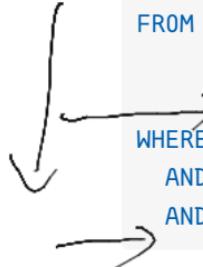
Interesting Query

Question

Find all MRT station that can be reached from NS1 in 3-stops.

Potential Solution

```
SELECT DISTINCT M3.to_stn
FROM   MRT M1,
       MRT M2,
       MRT M3
WHERE  M1.fr_stn = 'NS1'
      AND M1.to_stn = M2.fr_stn
      AND M2.to_stn = M3.fr_stn;
```



Table

to_stn
NS4

1 row

Note

What about n -stops? What is the maximum n ?

Extended Concepts

Universal Recursive

- MRT
- Query
- **Recursive CTE**
- Lazy Evaluation

Recursive Queries

Recursive CTE

Syntax

`WITH RECURSIVE`

`CTE_name AS (`

`Q_1`

`UNION [ALL]`

`Q_2 (CTE_name)`

`Q_0 (CTE_name)`

Notes:

- `Q_1` is non-recursive
- `Q_2` is recursive and can reference `CTE_name`
- Query is evaluated "lazily", stops when a *fixed-point* is reached
(i.e., no additional rows need to be printed)

Extended Concepts

Universal

Recursive

- MRT
- Query
- **Recursive CTE**
- Lazy Evaluation

Recursive Queries

Recursive CTE

Question

Find all MRT station that can be reached from NS1 in at most 3-stops.



Recursive Solution

```
WITH RECURSIVE
    Linker(to_stn, stops) AS (
        SELECT to_stn, 0
        FROM   MRT
        WHERE  fr_stn = 'NS1'
        UNION ALL
        SELECT M.to_stn, L.stops + 1
        FROM   Linker L, MRT M
        WHERE  L.to_stn = M.fr_stn
        AND   L.stops < 2
    )
SELECT DISTINCT (to_stn)
FROM   Linker;
```



Table

to_stn

to_stn
NS2
NS3
NS4

3 rows

Extended Concepts

Universal

Recursive

- MRT
- Query
- Recursive CTE
- Lazy Evaluation

Recursive Queries

Lazy Evaluation

Question

Find all MRT station that can be reached from NS1 in at most 3-stops.

Recursive Solution

```
WITH RECURSIVE
    Linker(to_stn, stops) AS (
        SELECT to_stn, 0
        FROM MRT
        WHERE fr_stn = 'NS1'
        UNION ALL
        SELECT M.to_stn, L.stops + 1
        FROM Linker L, MRT M
        WHERE L.to_stn = M.fr_stn
    )
    SELECT DISTINCT (to_stn)
    FROM Linker
    WHERE stops < 3;
```

$$f(x) = x$$

$$\underbrace{f(f(f(f(x))))}_{\text{3 stops}} = x$$

Table

to_stn
NS2
NS3
NS4

3 rows

Summary

Summary

Covered
Limitations

What We Have Covered

SQL Query

- Most common vocabulary for writing queries
- Basic means to "organize" complex queries (CTEs, Views)
- Extended features
 - Universal quantification
 - Recursive queries

Advantages

- More maintainable SQL query
- Easily extendable query

Summary

Covered
Limitations

Limitations of SQL Queries

Poorly Supported Queries

- "Sorted by pizza price, are there somewhere in the ranking 3 pizzas that are sold by the same restaurant listed in a row"
- Queries/tasks common for time-series: moving average, sliding window, etc

Common Approaches

- Keep or move logic into the application
- Use features that make SQL "Turing-complete"^{*} (e.g., stored procedures)
- Use different data model / DBMS (e.g., graph database for recursive queries, or use time-series databases)

^{*}Strictly speaking, the support of recursive CTEs already made SQL "Turing-complete".

```
postgres=# exit
```

```
Press any key to continue . . .
```

Solutions

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Exercise #1

Quiz #1

Quiz #1

What is the problem with the invalid query?

Hint

Look at the table created from the query on the right. Try to do your own grouping using **rname** as the grouping criteria. What do you see? Can you answer what is the value of **pizza** for all group?

```
SELECT rname, pizza, SUM(price)  
FROM Sells;
```

Solutions

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Exercise #1

Quiz #2

Quiz #2

What is the "problem" with the *valid* query?

Comment

Since `iso2` is the PRIMARY KEY of the table `Countries`, every single row can be uniquely identified by their `iso2` column. In other words, each "group" is a group of a single row.

Which brings us to the weird behaviour by PostgreSQL on why it does not allow candidate key (*i.e.*, `UNIQUE constraint, even with NOT NULL`) to be used as a grouping.

Solutions

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Exercise #1

Quiz #3

Quiz #3

Will it be valid if we do `SELECT R.rname, C.area ...` instead of `SELECT R.rname, R.area ...?`

Comment

Test it out and see if there is an error. This is a simplified case of the rule. By `R.area = C.area`, we know that if `R.area` is unique in the result, so should `C.area`. Is PostgreSQL looking into this "fact" before they are running the code?

Solutions

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Exercise #1

Quiz #4

Quiz #4

Will it be valid if we do `SELECT R.rname, C.area ...` instead of `SELECT R.rname, R.area ...?`

Comment

Given that in the `GROUP BY` clause, `rname` and `pizza` is the primary key of `Sells`, each row is in a group of its own. Therefore, the condition `price > 20` can be applied to each row. Can you not use `HAVING` clause in this case? If so, what will be the value of `COUNT(*)`?

Solutions

Quiz #1

Quiz #2

Quiz #3

Quiz #4

Exercise #1

HAVING Clause

Exercise #1

Question

Find all restaurants that sells any pizza cheaper than any pizza sold by Lorenzo Tavern.

Query

Result

rname
Pizza King
Mammas Place
Gambino Oven
Corleone Corner