

CS2102: Database Systems

Tutorial #4: SQL (Part 2)


Week 6 Guide

AY 2022/23 Sem 2

1 Discussions

The following questions are to be discussed during tutorial. All answers will be released with explanation.

This tutorial discussion questions are based on the following pizza database schema. The ER diagram is shown below.



The ER diagram produces the following schemas:

Relation	Description
<i>Pizzas</i> (<u><i>pizza</i></u>)	All the pizzas of interest.
<i>Customers</i> (<u><i>cname</i></u> , <i>area</i>)	The name and location of each customer.
<i>Restaurants</i> (<u><i>rname</i></u> , <i>area</i>)	The name and location of each restaurant.
<i>Recipes</i> (<u><i>pizza</i></u> , <u><i>ingredients</i></u>)	The ingredients used in each pizza.
<i>Sells</i> (<u><i>rname</i></u> , <u><i>pizza</i></u> , <i>price</i>)	Pizzas sold by restaurants and the prices.
<i>Likes</i> (<u><i>cname</i></u> , <u><i>pizza</i></u>)	Pizzas that customers like.

Additionally, we have the following foreign key constraints on the database schema:

- $(Recipes.pizza) \rightsquigarrow (Pizzas.pizza)$
- $(Sells.rname) \rightsquigarrow (Restaurants.rname)$
- $(Sells.pizza) \rightsquigarrow (Pizzas.pizza)$
- $(Likes.cname) \rightsquigarrow (Customers.cname)$
- $(Likes.pizza) \rightsquigarrow (Pizzas.pizza)$

1. **(Simple Query)** For each of the following queries, write an *equivalent* SQL query that does **not** use any subquery. Note that we do not consider set operation (e.g., in $Q1 \cup Q2$, neither $Q1$ nor $Q2$ are considered subqueries).

(a) Query A

```
1 SELECT DISTINCT cname
2 FROM Likes L
3 WHERE EXISTS (
4     SELECT 1
5     FROM Sells S
6     WHERE S.rname = 'Corleone Corner'
7     AND S.pizza = L.pizza
8 );
```

(c) Query C

```
1 SELECT DISTINCT rname
2 FROM Sells
3 WHERE rname <> 'Corleone Corner'
4 AND price > ANY (
5     SELECT price
6     FROM Sells
7     WHERE rname = 'Corleone Corner'
8 );
```

(b) Query B

```
1 SELECT cname
2 FROM Customers C
3 WHERE NOT EXISTS (
4     SELECT 1
5     FROM Likes L, Sells S
6     WHERE S.rname = 'Corleone Corner'
7     AND S.pizza = L.pizza
8     AND C.cname = L.cname
9 );
```

(d) Query D

```
1 SELECT rname, pizza, price
2 FROM Sells S
3 WHERE price >= ALL (
4     SELECT S2.price
5     FROM Sells S2
6     WHERE S2.rname = S.rname
7     AND S2.price IS NOT NULL
8 );
```

Suggested Guide:

A possible way to solve this kind of question is to translate the query into an equivalent English statement and then attempt to solve problem again without subquery.

- (a) The query can be translated into English roughly as:

Find all distinct customer names that likes at least one pizza that is sold by Corleone Corner. Include only customers that likes at least one pizza. (Alternatively, exclude customers that do not like any pizza.)

This can be solved without subquery by using join operation (e.g., *Cartesian product*, *inner join*, etc).

```
1 SELECT DISTINCT cname
2 FROM Likes L, Sells S
3 WHERE S.rname = 'Corleone Corner'
4 AND S.pizza = L.pizza;
```

(b) The query can be translated into English roughly as:

Find all distinct customer names that do not like any pizza that is sold by Corleone Corner. Include customers that do not like any pizza.

This can be solved without subquery by decomposing the problem into two parts:

- Find all customers that likes at least one pizza (a).
- Remove these customers from the collection of all customers (*i.e.*, *set difference*).

```
1 SELECT cname FROM Customers
2 EXCEPT
3 SELECT cname
4 FROM Likes L, Sells S
5 WHERE S.rname = 'Corleone Corner'
6 AND S.pizza = L.pizza;
```

(c) The query can be translated into English roughly as:

Find all restaurant names (besides Corleone Corner) that sells at least one pizza more expensive than any pizza sold by Corleone Corner.

This can be solved without subquery by using join operation (*e.g.*, *Cartesian product*, *inner join*, *etc*).

```
1 SELECT DISTINCT S.rname
2 FROM Sells S, Sells S2
3 WHERE S.rname <> 'Corleone Corner'
4 AND S2.rname = 'Corleone Corner'
5 AND S.price > S2.price;
```

(d) The query can be translated into English roughly as:

Find all distinct restaurant names (besides Corleone Corner), the pizza they sell, and the price of the pizza where the pizza is at least as expensive as the most expensive pizza sold by Corleone Corner.

This can be solved without subquery by decomposing the problem into two parts:

- Find all triples (restaurant names, pizza, price) where the price is cheaper than any pizza sold by Corleone Corner.
- Remove these triples from the collection of triples (restaurant names, pizza, price).

```
1 SELECT rname, pizza, price
2 FROM Sells
3 WHERE price IS NOT NULL
4 EXCEPT
5 SELECT S.rname, S.pizza, S.price
6 FROM Sells S, Sells S2
7 WHERE S.rname = S2.rname
8 AND S.price < S2.price;
```

We need the condition `price IS NOT NULL` here because these triple will not be removed by the second query. When `price` is `NULL`, `S.price < S2.price` is `NULL` so the triple is not in the second query.

Similarly, in the original `ALL` subquery, the result of the subquery will be `NULL` so the triple containing `NULL` price is not going to be in the result.

2. (SQL Query) Write an SQL query to answer each of the following questions on the pizza database *without using aggregate functions*. Remove duplicate records from all query results.

- (a) Find pizzas that Moe likes but is not liked by Lisa.
- (b) Find pizzas that are sold by at most one restaurant in each area; exclude pizzas that are not sold by any restaurant.
- (c) Find all tuples (A, P, P_{min}) where P is a pizza that is available in area A (*i.e.*, there is some restaurant in area A selling pizza P) and P_{min} is the *lowest* price of P in area A .

Suggested Guide:

For each question, we will discuss some possible solutions (*remember, there can be multiple solutions*) and potentially some solutions that may look correct but are actually wrong.

(a) **Solution 1:**

```
1 SELECT pizza FROM Likes
2 WHERE  cname = 'Moe'
3     AND pizza NOT IN (
4         SELECT pizza FROM Likes
5         WHERE  cname = 'Lisa'
6     );
```

Solution 2:

```
1 SELECT pizza FROM Likes L1
2 WHERE  cname = 'Moe'
3     AND NOT EXISTS (
4         SELECT 1 FROM Likes L2
5         WHERE L2.cname = 'Lisa' AND L2.pizza = L1.pizza
6     );
```

Solution 3:

```
1 SELECT pizza FROM Likes
2 WHERE  cname = 'Moe'
3     AND NOT pizza = ANY (
4         SELECT pizza FROM Likes
5         WHERE  cname = 'Lisa'
6     );
```

Note that if Lisa does not like any pizza, then the ANY subquery will evaluate to FALSE and NOT FALSE will evaluate to TRUE. Thus the query will return all the pizzas that Moe likes.

Solution 4: *Probably simplest*

```
1 SELECT pizza FROM Likes WHERE  cname = 'Moe'
2 EXCEPT
3 SELECT pizza FROM Likes WHERE  cname = 'Lisa';
```

WRONG ANSWER: *The following answer is incorrect*

```
1 SELECT pizza FROM Likes
2 WHERE cname = 'Moe'
3 AND pizza <> ANY (
4     SELECT pizza FROM Likes
5     WHERE cname = 'Lisa'
6 );
```

This answer looks similar to **Solution 3** BUT it is *incorrect*. If Lisa does not like any pizza, then the ANY subquery will evaluate to FALSE and the query will return an empty set. This will be incorrect if Moe likes some pizza.

- (b) A pizza is the output if there does not exist two distinct restaurants that are located in the same area selling that pizza.

```
1 SELECT DISTINCT pizza
2 FROM Sells S3
3 WHERE NOT EXISTS (
4     SELECT 1
5     FROM Sells S, Restaurants R, Sells S2, Restaurants R2
6     WHERE S.rname = R.rname AND S2.rname = R2.rname
7     AND S.pizza = S2.pizza AND R.area = R2.area
8     AND R.rname <> R2.rname AND S.pizza = S3.pizza
9 );
```

WRONG ANSWER: *The following answer is incorrect*

```
1 SELECT DISTINCT pizza
2 FROM Sells S, Restaurants R
3 WHERE S.rname = R.rname
4 AND NOT EXISTS (
5     SELECT 1
6     FROM Sells S2, Restaurants R2
7     WHERE S2.rname = R2.rname AND S.pizza = S2.pizza
8     AND R.area = R2.area AND R.rname <> R2.rname
9 );
```

This answer computes the pizzas that are sold by at most one restaurant in *some* area, which is a weaker condition than what is required by the question.

- (c) A possible solution

```
1 SELECT DISTINCT R.area, S.pizza, S.price
2 FROM Restaurants R, Sells S
3 WHERE R.rname = S.rname
4 AND S.price <= ALL (
5     SELECT S2.price
6     FROM Restaurants R2, Sells S2
7     WHERE R2.rname = S2.rname
8     AND R2.area = R.area
9     AND S2.pizza = S.pizza
10 );
```

3. **(Equivalence)** Consider the query to find distinct restaurants that are located in the East area. The following are two possible SQL answers (*denoted by Q_1 and Q_2*) for this query.

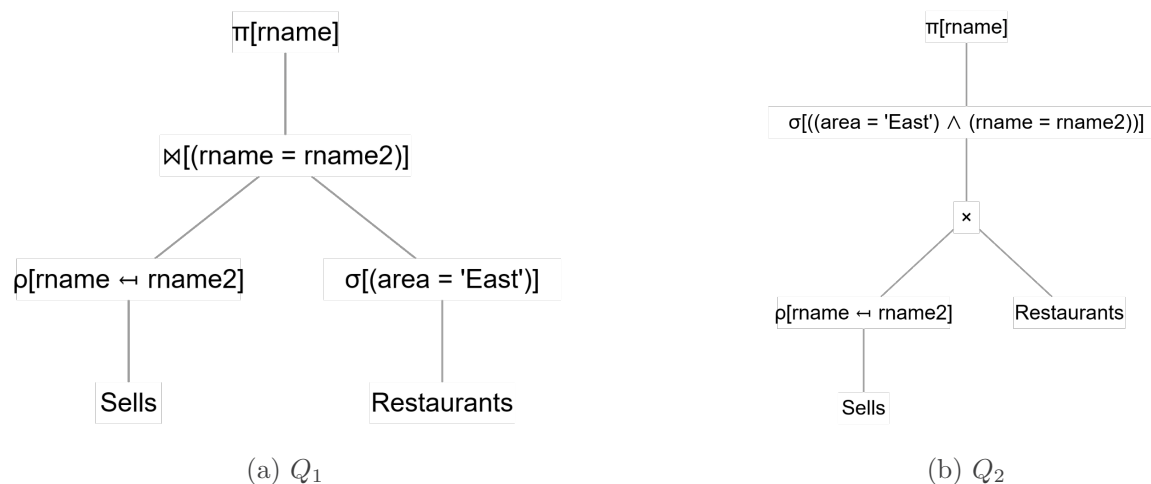
Q_1 Query 1

```
1 SELECT DISTINCT S.rname
2 FROM   Sells S JOIN Restaurants R
3     ON   S.rname = R.rname AND R.area = 'East';
```

Q_2 Query 2

```
1 SELECT DISTINCT S.rname
2 FROM   Sells S, Restaurants R
3 WHERE  S.rname = R.rname
4     AND R.area = 'East';
```

The semantics of these two SQL queries are defined by the relational algebra expressions shown below. Discuss whether Q_1 and Q_2 are equivalent queries.



Suggested Guide:

Queries Q_1 and Q_2 are **equivalent**. Whether the selection predicate “`area = 'East'`” is evaluated *before* or *after* the join/cross product operation does not change the semantics of the query.

4. **(Equivalence)** Consider the query to find distinct restaurants that are located in the East area or restaurants that sell some pizza that Lisa likes, where the restaurants that do not sell any pizza are to be excluded. The following are two possible SQL answers (*denoted by Q_1 and Q_2*) for this query.

Q_1 Query 1

```
1 SELECT DISTINCT S.rname
2 FROM   Sells S JOIN Restaurants R
3     ON   S.rname = R.rname AND R.area = 'East '
4 UNION
5 SELECT DISTINCT S.rname
6 FROM   Sells S JOIN Likes L
7     ON   S.pizza = L.pizza AND L.cname = 'Lisa';
```

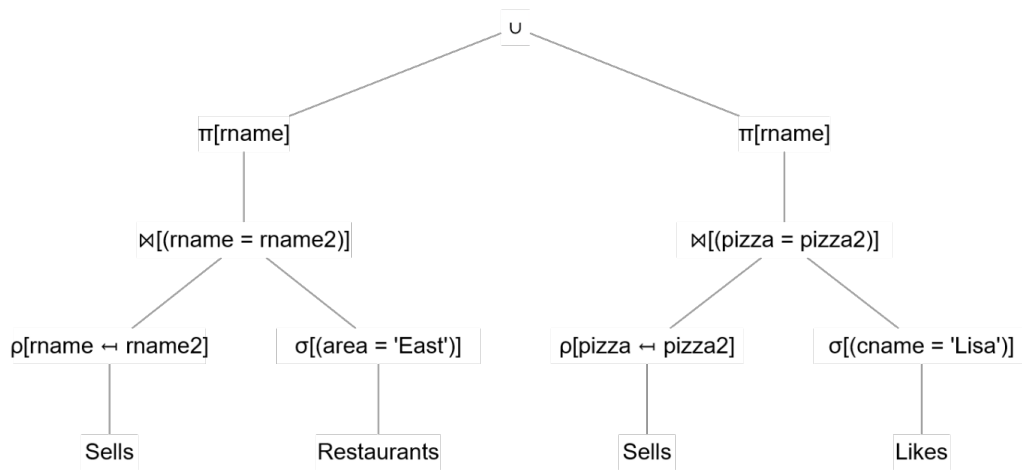
Q_2 Query 2

```

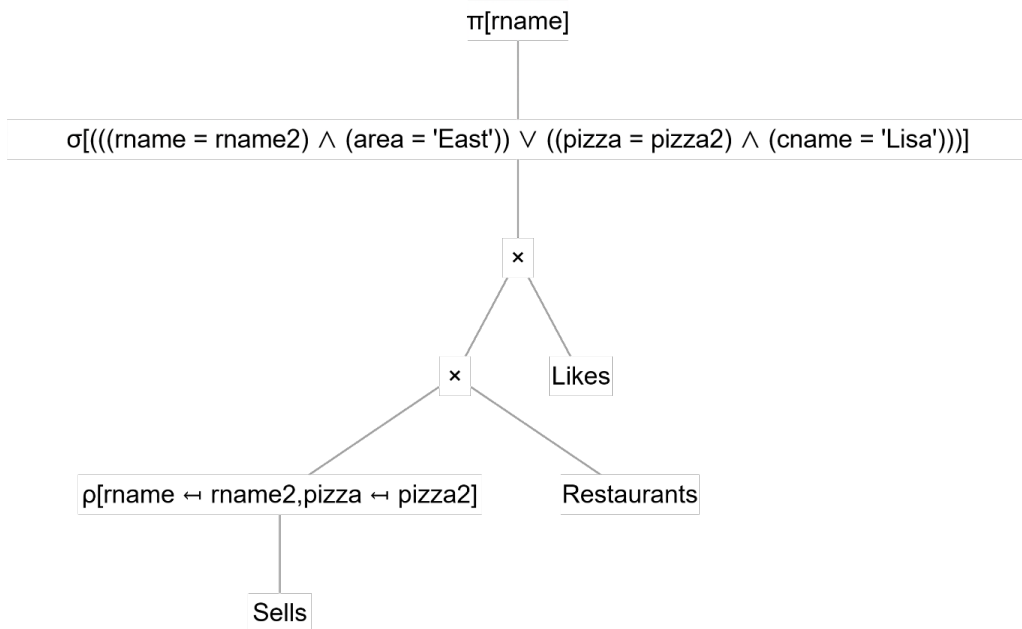
1  SELECT DISTINCT S.rname
2  FROM   Sells S, Restaurants R, Likes L
3  WHERE  (S.rname = R.rname AND R.area = 'East')
4         OR (S.pizza = L.pizza AND L.cname = 'Lisa');

```

The semantics of these two SQL queries are defined by the relational algebra expressions shown below. Discuss whether Q_1 and Q_2 are equivalent queries.



(a) Q_1



(b) Q_2

Suggested Guide:

Queries Q_1 and Q_2 are **NOT equivalent**. Observe that if the **Likes** relation is *empty*, then query Q_1 simplifies to

$$\pi_{\text{rname}}(\text{Sells} \bowtie \sigma_{\text{area}='East'}(\text{Restaurants}))$$

BUT the result of Q_2 is always an empty set due to

$$\text{Sells} \times \text{Restaurants} \times \emptyset$$

2 Challenge

The answers to the following questions is given without explanation. Please discuss them on Canvas.

1. **(SQL Query)** Write an SQL query to answer each of the following questions on the pizza database *without using aggregate functions*. Remove duplicate records from all query results.
 - (a) Find all tuples (A, P, P_{min}, P_{max}) where P is a pizza that is available in area A (*i.e.*, there is some restaurant in area A selling pizza P), P_{min} is the *lowest* price of P in area A and P_{max} is the *highest* price of P in area A .

Suggested Guide:

- (a) You should recognize that this query is simply an extension of Question 2(c) requiring an additional information (*i.e.*, highest selling price) for each area-pizza pair. For a given area-pizza pair (A, P) , the following query will compute the highest price of pizza P in area A .

```
1 SELECT DISTINCT S2.price
2 FROM   Restaurants R2, Sells S2
3 WHERE  R2.rname = S2.rname
4       AND R2.area = A AND S2.pizza = P
5       AND R2.price >= ALL (
6         SELECT S3.price
7         FROM   Restaurants R3, Sells S3
8         WHERE  R2.rname = S2.rname
9               AND R3.area = A
10              AND S3.pizza = P
11 );
```

Since the above query will return a *single one-column tuple*, it can be used as a **scalar** subquery to extend the previous question's solution as follows:

```
1 SELECT DISTINCT R.area, S.pizza, S.price AS minPrice, (
2   SELECT DISTINCT S2.price
3   FROM   Restaurants R2, Sells S2
4   WHERE  R2.rname = S2.rname
5         AND R2.area = R.area AND S2.pizza = S.pizza
6         AND R2.price >= ALL (
```



```

7      SELECT S3.price
8      FROM   Restaurants R3, Sells S3
9      WHERE  R2.rname = S2.rname
10     AND    R3.area = R.area
11     AND    S3.pizza = S.pizza
12 ) AS maxPrice
13 FROM   Restaurants R, Sells S
14 WHERE  R.rname = S.rname
15     AND S.price <= ALL (
16     SELECT S2.price
17     FROM   Restaurants R2, Sells S2
18     WHERE  R2.rname = S2.rname
19     AND    R2.area = R.area
20     AND    S2.pizza = S.pizza
21 );

```

We will learn about other (*simpler and more elegant*) solutions for such queries later in class.

2. **(Update)** Consider *again* the following relational schema discussed in Tutorial 2 Question 2 (*Discussions*).

```

1 CREATE TABLE Offices (
2     office_id    INTEGER,
3     building     TEXT NOT NULL,
4     level        INTEGER NOT NULL,
5     room_number  INTEGER NOT NULL,
6     area         INTEGER,
7     PRIMARY KEY (office_id),
8     UNIQUE (building, level, room_number)
9 );
10
11 CREATE TABLE Employees (
12     emp_id       INTEGER,
13     name         TEXT NOT NULL,
14     office_id    INTEGER NOT NULL,
15     manager_id   INTEGER,
16     PRIMARY KEY (emp_id),
17     FOREIGN KEY (office_id) REFERENCES Offices (office_id)
18     ON UPDATE CASCADE,
19     FOREIGN KEY (manager_id) REFERENCES Employees (emp_id)
20     ON UPDATE CASCADE
21 );

```

Suppose that the office with `office_id = 123` needs to be renovated. Write an SQL statement to reassign the employees located in this office to another temporary office located at room number 11 on level 5 at the building named *Tower1*.

Hint: You can use subquery in an update statement.

Suggested Guide:

SQL update statement.

```
1 UPDATE Employees
2 SET     office_id = (SELECT office_id FROM Offices
3                     WHERE building = 'Tower1'
4                       AND level = 5
5                       AND room_number = 11)
6 WHERE  office_id = 123;
```

3. **(Backward Reasoning)** You are given the following schema:

- Students(matric, sname)
- Workings(pid, matric, since)
- Projects(pid, pname)
- Categories(pid, cname)

You are also given the following foreign key:

- (Workings.matric) \rightsquigarrow (Students.matric)
- (Workings.pid) \rightsquigarrow (Projects.pid)
- (Categories.pid) \rightsquigarrow (Projects.pid)

Consider the following SQL query:

```
1 SELECT *
2 FROM   Students NATURAL JOIN Workings
3        NATURAL JOIN Projects
4        NATURAL JOIN Categories;
```

Say it produces the following result:

matric	sname	pid	since	cname
A0001	AA	P01	2002	CA
A0001	AA	P01	2002	CB
A0001	AA	P02	2004	CB
A0002	BB	P01	2003	CA
A0002	BB	P01	2003	CB
A0003	CC	P03	2004	CA
A0003	CC	P03	2004	CC
A0003	CC	P03	2004	CD
A0004	AA	P03	2004	CA
A0004	AA	P03	2004	CC
A0004	AA	P03	2004	CD

Now consider another the query to find all pair of distinct projects' pid (*i.e.*, (p1, p2)) such that the two projects have exactly the same set of categories¹. It produces the following result:

pid	pid
P01	P04
P04	P01
P03	P05
P05	P03

Simply note that P01 and P04 have exactly the same set of categories. Similarly, P03 and P05 have exactly the same set of categories. What is a possible result of the following SQL query? Show your answer using only P01, P03, P04, and P05?

```
1 SELECT pid FROM Categories
2 EXCEPT ALL
3 SELECT DISTINCT pid FROM Categories WHERE cname <> 'CA';
```

Suggested Guide:

A possible solution using only information in both tables above is:

pid
P01
P03
P03
P04
P05
P05

¹This involves some constructs that you will only learn in Lecture 06.