

CS2102: Database Systems

Tutorial #5: SQL (Part 3)

Week 7 Guide

AY 2022/23 Sem 2

1 Discussions

The following questions are to be discussed during tutorial. All answers will be released with explanation.

1. **(Equivalence)** Consider the following relational schema:

```
1 CREATE TABLE R (  
2   a  INTEGER PRIMARY KEY ,  
3   b  INTEGER ,  
4   c  INTEGER  
5 );  
6  
7 CREATE TABLE S (  
8   x  INTEGER PRIMARY KEY ,  
9   y  INTEGER REFERENCES R(a)  
10 );
```

Are these two queries *equivalent*?

(a) Q_1

```
1 SELECT  COUNT(c)  
2 FROM    R  
3 WHERE   a = 10;
```

(b) Q_2

```
1 SELECT  COUNT(c)  
2 FROM    R  
3 WHERE   a = 10  
4 GROUP BY a;
```

Suggested Guide:


Q_1 and Q_2 are **NOT** equivalent queries. Consider the case where there is no record in R with $a = 10$.

- (a) Q_1 : The WHERE clause will evaluate to an *empty table* and the aggregate function COUNT(c) will evaluate to a *single-row table* with a *single-column* value of 0.

- (b) Q_2 : The WHERE clause will evaluate to an *empty table* so the GROUP BY clause will **NOT** produce any group. As such, the COUNT(c) will **NOT** be evaluated at all. As a result, the query result will be an *empty table*.

2. (Un-Aggregate) This question is based on the pizza database schema shown below.

Answer each of the following queries using SQL. For parts (a) to (e), remove duplicate records from all query results.



- Find the most expensive pizzas and the restaurants that sell them (*at the most expensive price*). Do **NOT** use any aggregate function in your answer.
- Find all restaurant pairs (R_1, R_2) such that the price of the most expensive pizza sold by R_1 is *higher* than that of R_2 . Exclude restaurant pairs where R_2 do not sell any pizza.
- For each restaurant that sells some pizza, find the restaurant name and the average price of its pizzas if its average price is higher than \$22. Do **NOT** use the HAVING clause in your answer.
- For each restaurant R that sells some pizza, let $totalPrice(R)$ denote the total price of all the pizzas sold by R . Find all pairs $(R, totalPrice(R))$ where $totalPrice(R)$ is *higher* than the average of $totalPrice()$ over all the restaurants.
- Find the customer pairs (C_1, C_2) such that $C_1 < C_2$ and they like *exactly* the same pizzas. Exclude customer pairs that do not like any pizza. Do **NOT** use the EXCEPT operator in your answer.

Suggested Guide:

(a) Solution 1:

```

1  SELECT pizza, rname
2  FROM    Sells
3  WHERE   price >= ALL (SELECT price FROM Sells);

```

Solution 2:

```

1  SELECT pizza, rname
2  FROM    Sells S1
3  WHERE   NOT EXISTS (
4      SELECT 1
5      FROM    Sells S2
6      WHERE   S2.price > S1.price
7  );

```

(b) **Solution 1:**

```
1 SELECT R1.rname, R2.rname
2 FROM   Restaurants R1, Restaurants R2
3 WHERE  (SELECT MAX(price) FROM Sells
4         WHERE rname = R1.rname)
5         >
6         (SELECT MAX(price) FROM Sells
7         WHERE rname = R2.rname);
```

Solution 2:

```
1 WITH RestMaxPrice AS (
2     SELECT rname, (SELECT MAX(price) FROM Sells
3                     WHERE rname = R.rname) as maxPrice
4     FROM   Restaurants R
5 )
6 SELECT R1.rname, R2.rname
7 FROM   RestMaxPrice R1, RestMaxPrice R2
8 WHERE  R1.maxPrice > R2.maxPrice;
```

(c) A possible solution

```
1 WITH RestAvgPrice AS (
2     SELECT rname, AVG(price) as avgPrice
3     FROM   Sells
4     GROUP BY rname
5 )
6 SELECT *
7 FROM   RestAvgPrice
8 WHERE  avgPrice > 22;
```

(d) **Solution 1:**

```
1 WITH RestTotalPrice AS (
2     SELECT rname, SUM(price) as totalPrice
3     FROM   Sells
4     GROUP BY rname
5 )
6 SELECT rname, totalPrice
7 FROM   RestTotalPrice
8 WHERE  totalPrice > (SELECT AVG(totalPrice) FROM RestTotalPrice);
```

Solution 2:

```
1 SELECT rname, SUM(price) as totalPrice
2 FROM   Sells S
3 GROUP BY rname
4 HAVING SUM(price) > (SELECT SUM(price) / COUNT(DISTINCT rname) FROM
5                     Sells);
```

WRONG ANSWER 1: The following query is an **invalid** query

```
1 SELECT rname, SUM(price) as totalPrice
2 FROM   Sells S
3 GROUP BY rname
4 HAVING totalPrice > (SELECT SUM(price) / COUNT(DISTINCT rname) FROM
5                     Sells);
```

totalPrice is *undefined* in the HAVING clause as the SELECT clause is conceptually evaluated *after* the HAVING clause.

WRONG ANSWER 2: The following query produce the wrong result

```

1 SELECT rname, SUM(price)
2 FROM Sells
3 GROUP BY rname
4 HAVING SUM(price) > SUM(price) / COUNT(*);

```

The query above is *incorrect* because both $SUM(price)$ and $COUNT(*)$ in the HAVING clause are computed w.r.t. a group.

- (e) **Solution 1:** Customer C1 and C2 like exactly the same pizzas if and only if (a) for every pizza that C1 likes, C2 also likes that pizza, and (b) for every pizza that C2 likes, C1 also likes that pizza. Conditions (a) holds if there does not exists any pizza that C1 likes that C2 does not like.

```

1 SELECT C1.cname, C2.cname
2 FROM Customers C1, Customers C2
3 WHERE C1.cname < C2.cname
4 AND EXISTS (SELECT 1 FROM Likes L WHERE L.cname = C1.cname)
5 AND NOT EXISTS (
6     SELECT 1
7     FROM Likes L1
8     WHERE L1.cname = C2.cname
9     AND NOT EXISTS (
10         SELECT 1
11         FROM Likes L2
12         WHERE L2.cname = C2.cname
13         AND L2.pizza = L1.pizza
14     )
15 )
16 AND NOT EXISTS (
17     SELECT 1
18     FROM Likes L2
19     WHERE L2.cname = C2.cname
20     AND NOT EXISTS (
21         SELECT 1
22         FROM Likes L1
23         WHERE L1.cname = C1.cname
24         AND L1.pizza = L2.pizza
25     )
26 );

```

Solution 2: This solution is based on the property that if customer C1 likes the set of pizzas $S1$ and customer C2 likes the set of pizzas $S2$, then $S1 = S2$ iff $|S1 \cap S2| = |S1| = |S2|$ where $|X|$ denote the cardinality of set X . Alternatively, you can also use the following equality $|S1 \cap S2| = |S1 \cup S2|$.

Here, we use CTE to solve the problem but a solution without CTE is possible.

```

1 WITH NumLike AS (
2     SELECT cname, COUNT(*) AS num FROM Likes L
3     GROUP BY cname
4 ), NumBothLike AS (
5     SELECT L1.cname AS cname1, L2.cname AS cname2, COUNT(*) AS num
6     FROM Likes L1, Likes L2
7     WHERE (L1.pizza = L2.pizza) AND (L1.cname < L2.cname)
8     GROUP BY L1.cname, L2.cname
9 )
10 -- continue on the next page
11

```

```

12 SELECT cname1, cname2
13 FROM   NumBothLike B
14 WHERE  num = (SELECT num FROM NumLike N WHERE N.cname = B.cname1)
15       AND num = (SELECT num FROM NumLike N WHERE N.cname = B.cname2);

```

Solution 3: This solution is based on the property that if customers C1 and C2 like exactly the same pizzas, then for the subset of all tuples $S \subseteq \text{Likes}$ that are associated with C1 or C2, there must be exactly two tuples in S associated with each distinct pizza in S .

```

1 SELECT C1.cname, C2.cname
2 FROM   Customers C1, Customers C2
3 WHERE  C1.cname < C2.cname
4       AND C1.cname IN (SELECT cname FROM Likes)
5       AND NOT EXISTS (
6           SELECT 1
7           FROM    Likes
8           WHERE   cname IN (C1.cname, C2.cname)
9           GROUP BY pizza
10          HAVING  COUNT(*) <> 2
11       );

```

(f) **Solution:** This is simply done using the UPDATE statement with CASE expression.

```

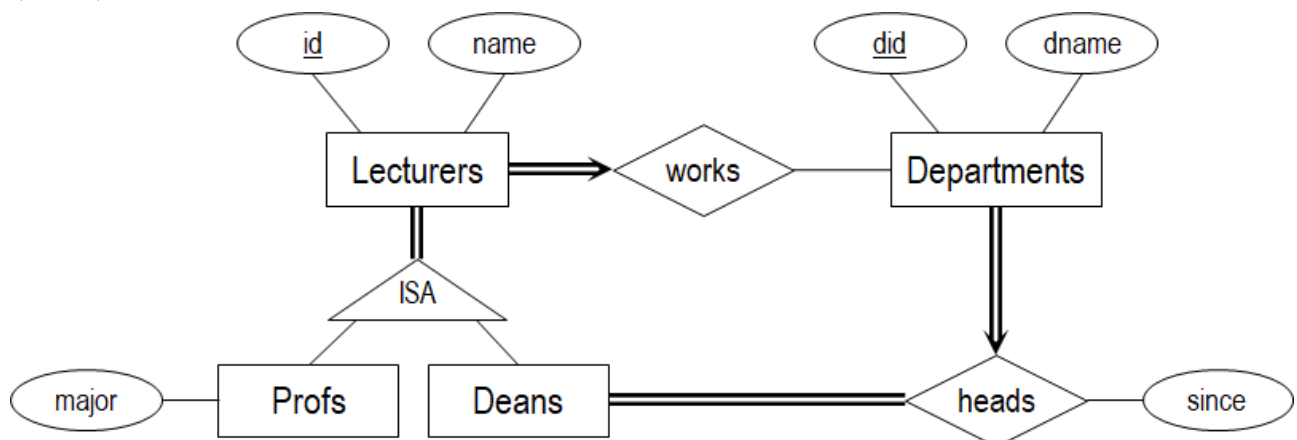
1 UPDATE Sells S
2 SET   price = CASE (
3         SELECT area
4         FROM   Restaurants
5         WHERE  rname = S.rname
6     )
7     WHEN 'Central' THEN price + 3
8     WHEN 'East' THEN price + 2
9     ELSE price + 1
10 END;

```

2 Challenge

The answers to the following questions is given without explanation. Please discuss them on Canvas.

1. (CTE) Consider the following ER diagram.



We assume that we have the following schema:

- Lecturers(id, name, did)
- Profs(id, major)
- Deans(id)
- Departments(did, dname, id)

We say that a dean is *important* if the dean heads a department where either:

- There are at least 20 professors working in the department excluding the dean, OR
- There are at least 5 professors with different majors working in the department

Find all the *non-important* dean. **Hint:** You may use CTE to simplify the problem.

Suggested Guide:

Our solution does not actually use CTE. But you may use CTE to simplify the problem.

```
1 SELECT DISTINCT D.lid, D.name
2 FROM   Deans D, Departments H, Profs P
3 WHERE  D.id = H.id AND H.did = P.did
4       AND ((SELECT COUNT(DISTINCT P1.id)
5             FROM   Deans D1, Departments H1, Profs P1
6             WHERE  D1.id = H1.id AND H1.did = P1.did
7                 AND D1.id = D.id AND D1.id <> P1.id
8             GROUP BY D1.id) > 20
9       OR
10      (SELECT COUNT(DISTINCT major)
11      FROM   Deans D1, Departments H1, Profs P1
12      WHERE  D1.id = H1.id AND H1.did = P1.did
13          AND D1.id = D.lid
14      GROUP BY D1.id) > 5
15 );
```

2. (Universal Quantification) Consider the following schema:

```
1 CREATE TABLE Students (
2   matric VARCHAR(9) PRIMARY KEY,
3   sname  VARCHAR(50)
4 );
5 CREATE TABLE Projects (
6   pid    VARCHAR(9) PRIMARY KEY,
7   pname  VARCHAR(50)
8 );
9 CREATE TABLE Workings (
10  pid      VARCHAR(9) REFERENCES Projects(pid),
11  matric   VARCHAR(9) REFERENCES Students(matric),
12  since    DATE,
13  PRIMARY KEY(pid, matric)
14 );
15 CREATE TABLE Category (
16  pid      VARCHAR(9) REFERENCES Projects(pid),
17  cname    VARCHAR(9),
18  PRIMARY KEY(pid, cname)
19 );
```

Find all pair of distinct projects' pid (p1, p2) such that the two projects have exactly the same set of categories.

Suggested Guide:

Here we use the cardinality method. The two sets must be identical so the intersection and the union must have the same size.

```
1 SELECT P1.pid, P2.pid
2 FROM   Projects P1, Projects P2
3 WHERE  P1.pid <> P2.pid
4       AND ( SELECT COUNT(*)
5             FROM ( SELECT cname FROM Categories C0 WHERE P1.pid = C0.pid
6                   INTERSECT -- alternatively UNION
7                   SELECT cname FROM Categories C0 WHERE P2.pid = C0.pid ) AS T1 )
8       =
9       ( SELECT COUNT(*)
10      FROM ( SELECT cname FROM Categories C0 WHERE P1.pid = C0.pid
11            UNION      -- alternatively INTERSECT
12            SELECT cname FROM Categories C0 WHERE P2.pid = C0.pid ) AS T2 )
;
```